

# Parallel Local Graph Clustering

Kimion Fountoulakis, joint work with J. Shun, X. Cheng, F. Roosta-Khorasani, M. Mahoney, D. Gleich  
University of California Berkeley and Purdue University

Based on

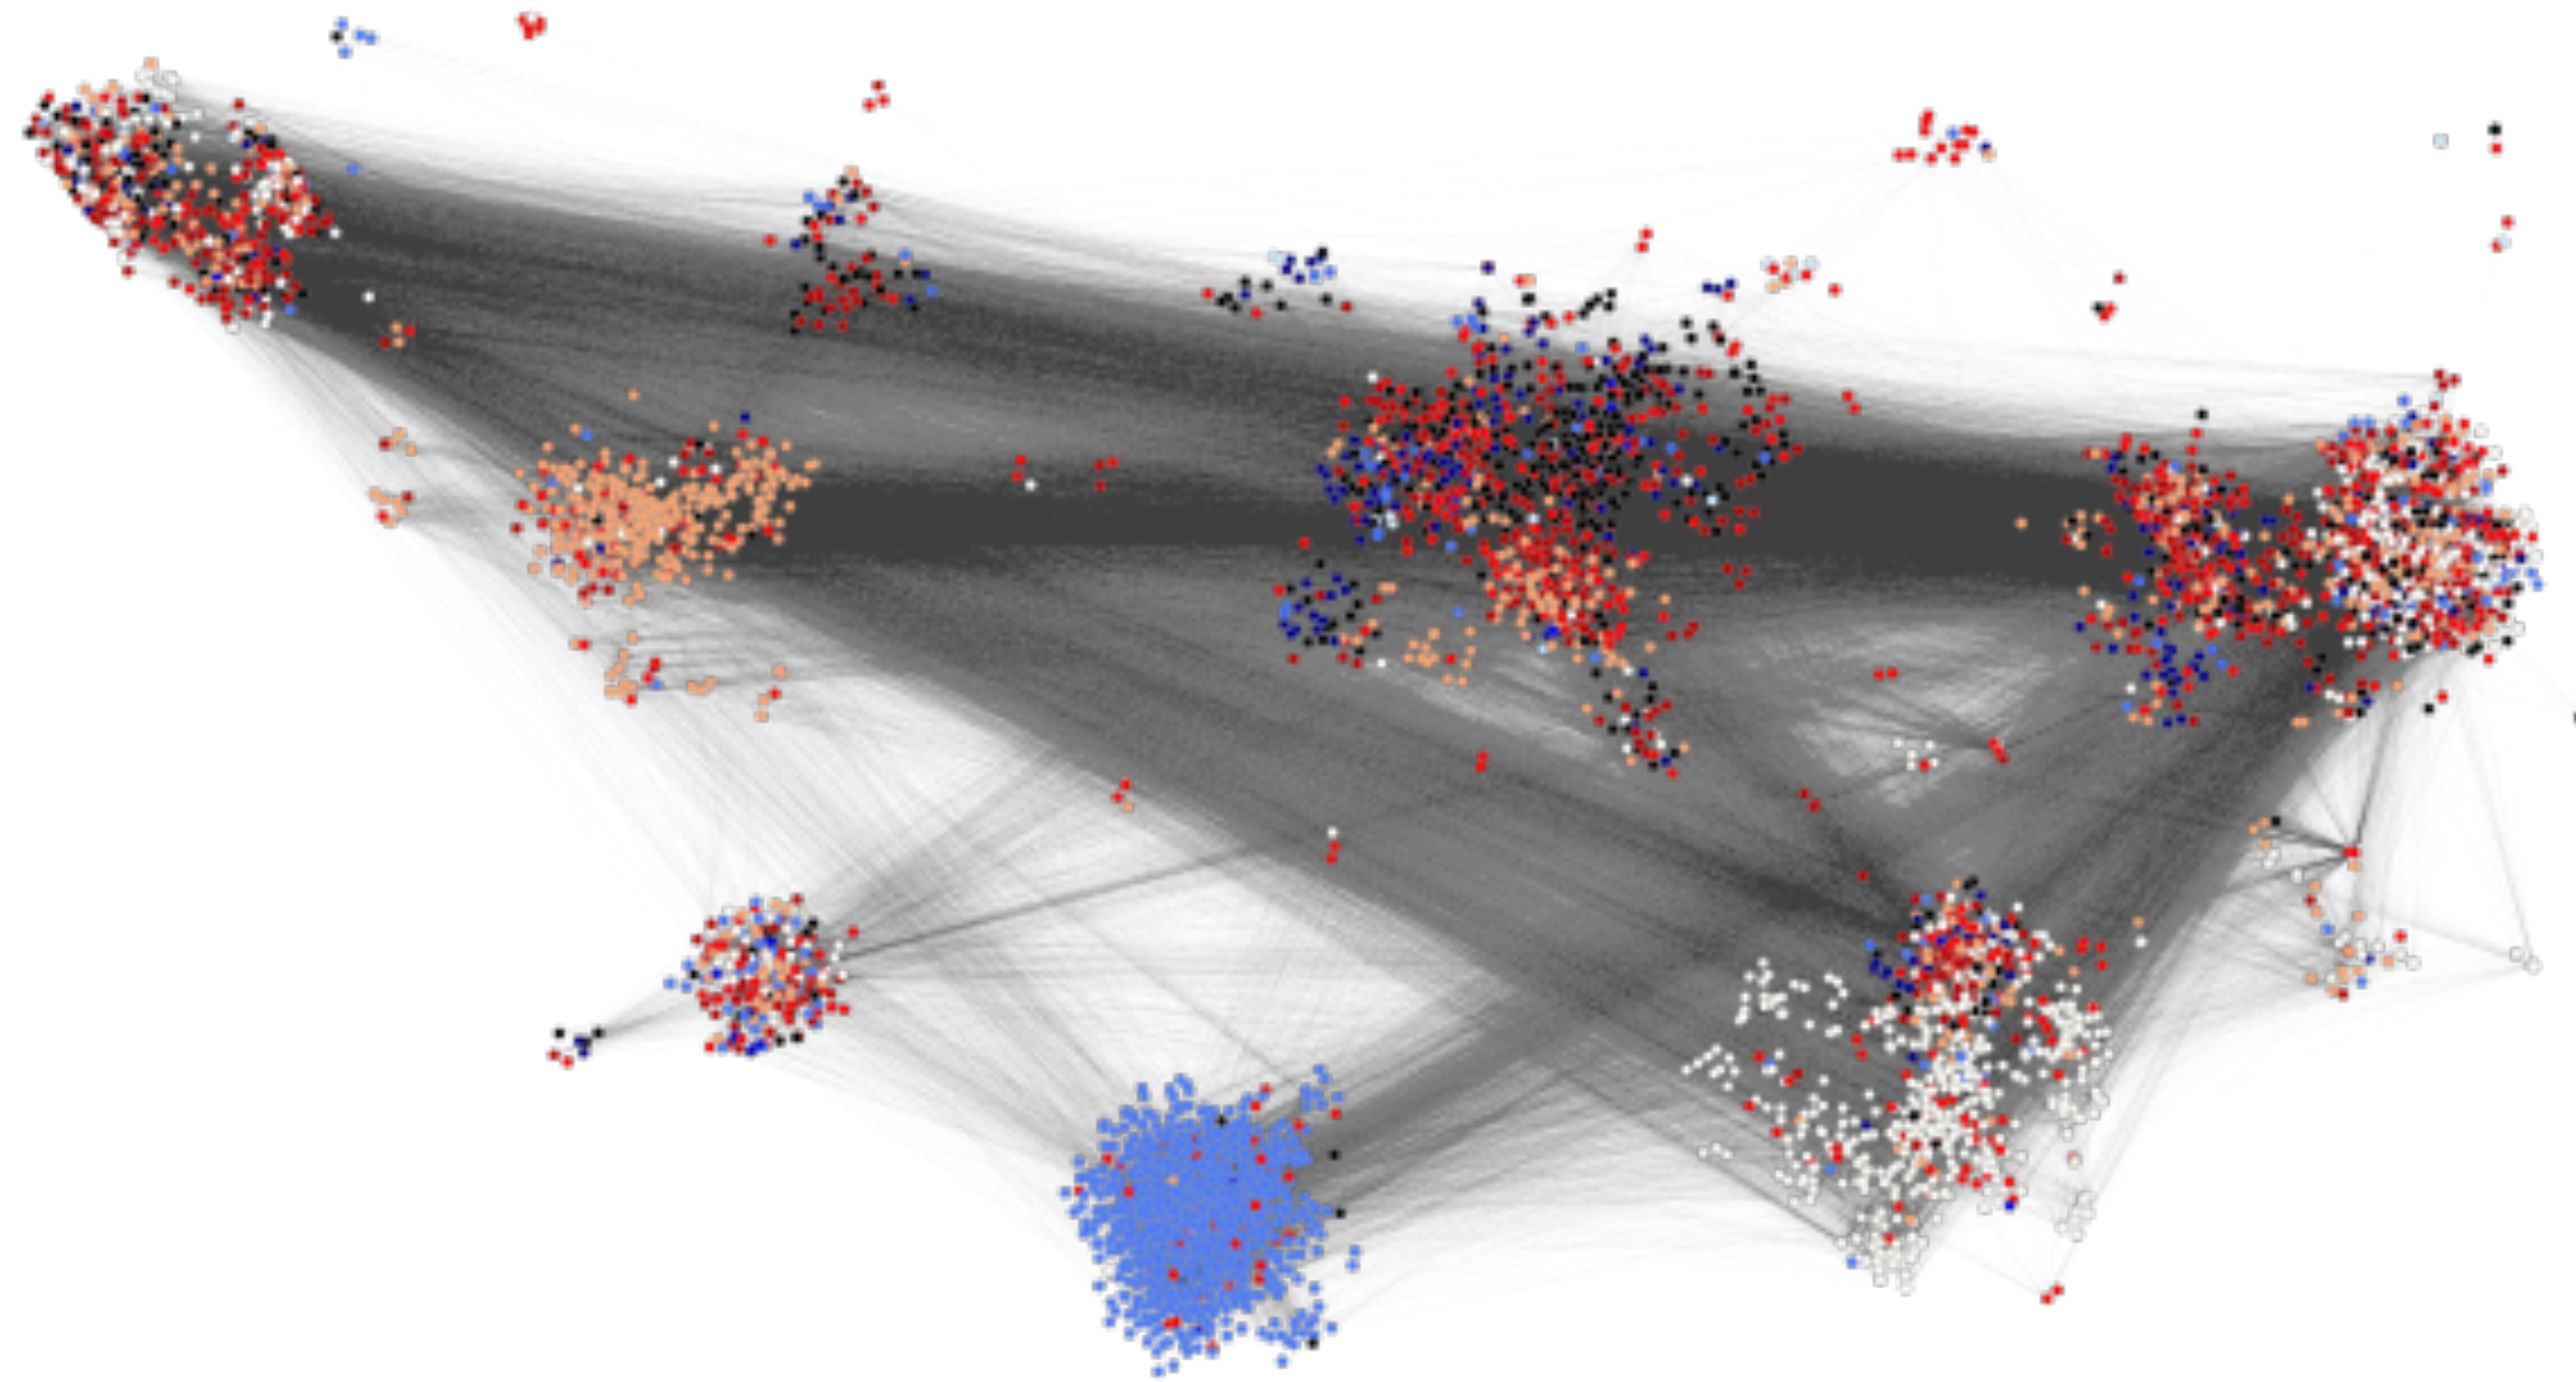
J. Shun, F. Roosta-Khorasani, KF, M. Mahoney. Parallel local graph clustering, VLDB 2016.

KF, X. Cheng, J. Shun, F. Roosta-Khorasani, M. Mahoney. Exploiting optimization for local graph clustering, arXiv:1602.01886v1.

KF, D. Gleich, M. Mahoney. An optimization approach to locally-biased graph algorithms, arXiv:1607.04940.

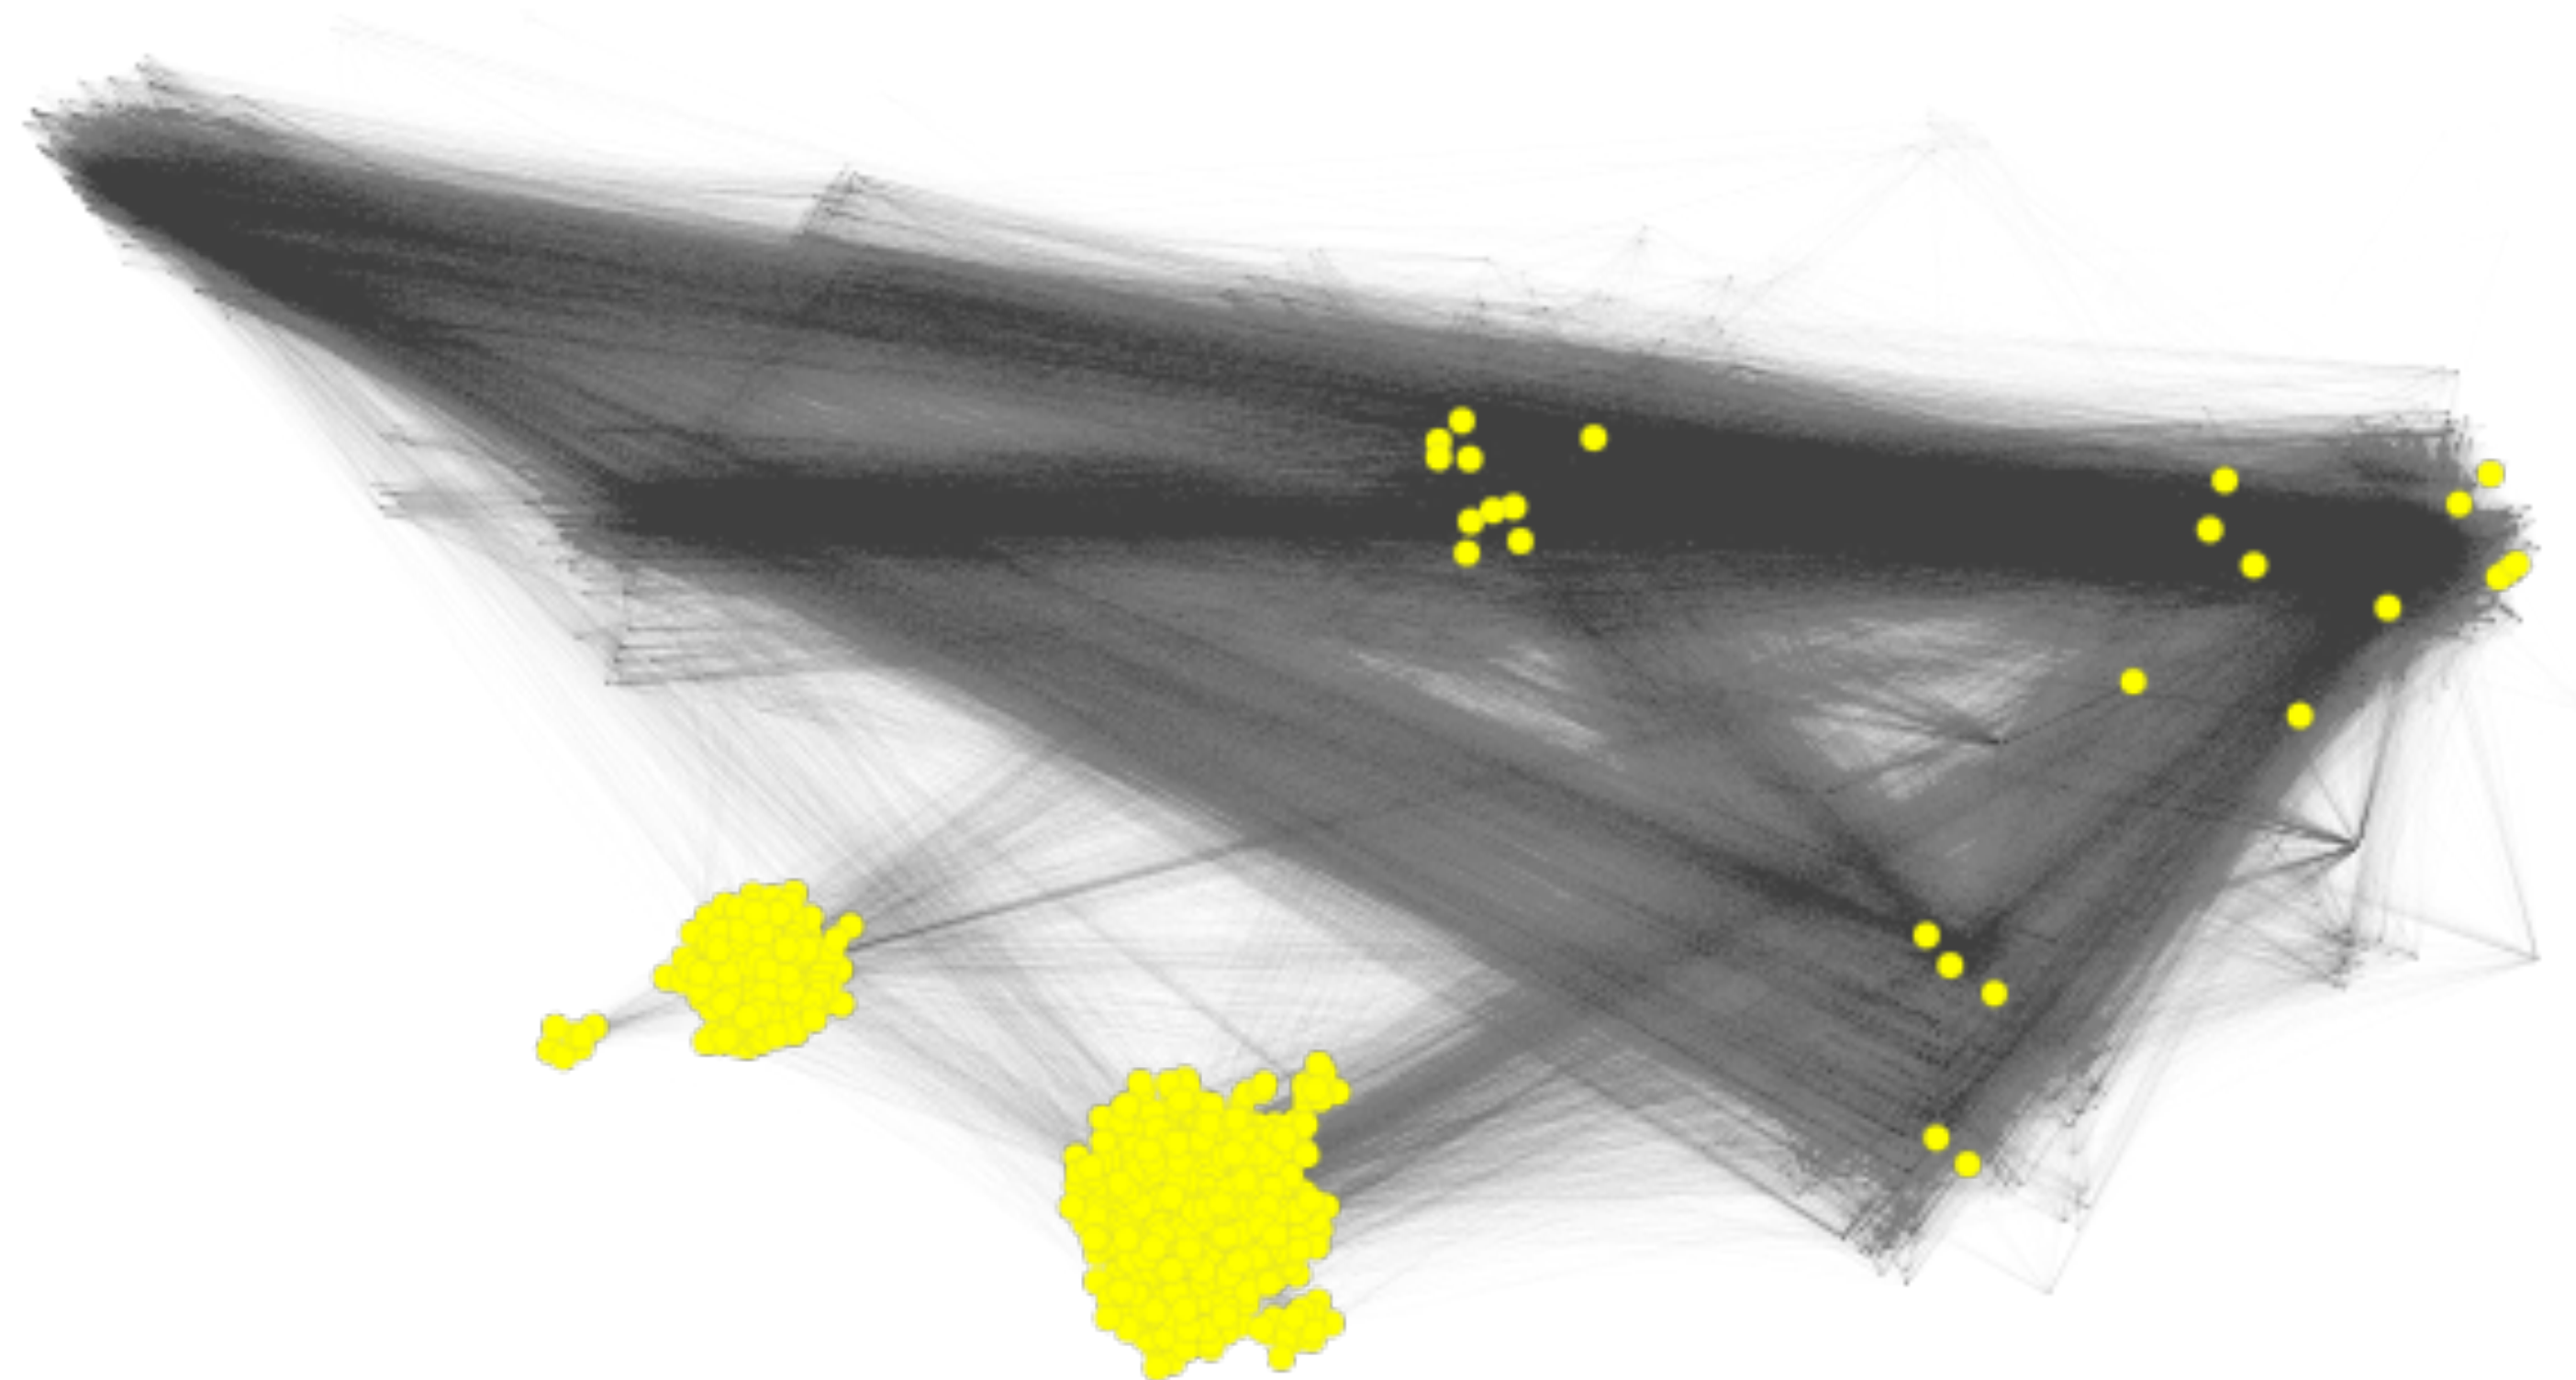
# Local graph clustering: motivation

# Facebook social network: colour denotes class year



Data: Facebook John Hopkins, A. L. Traud, P. J. Mucha and M. A. Porter, Physica A, 391(16), 2012

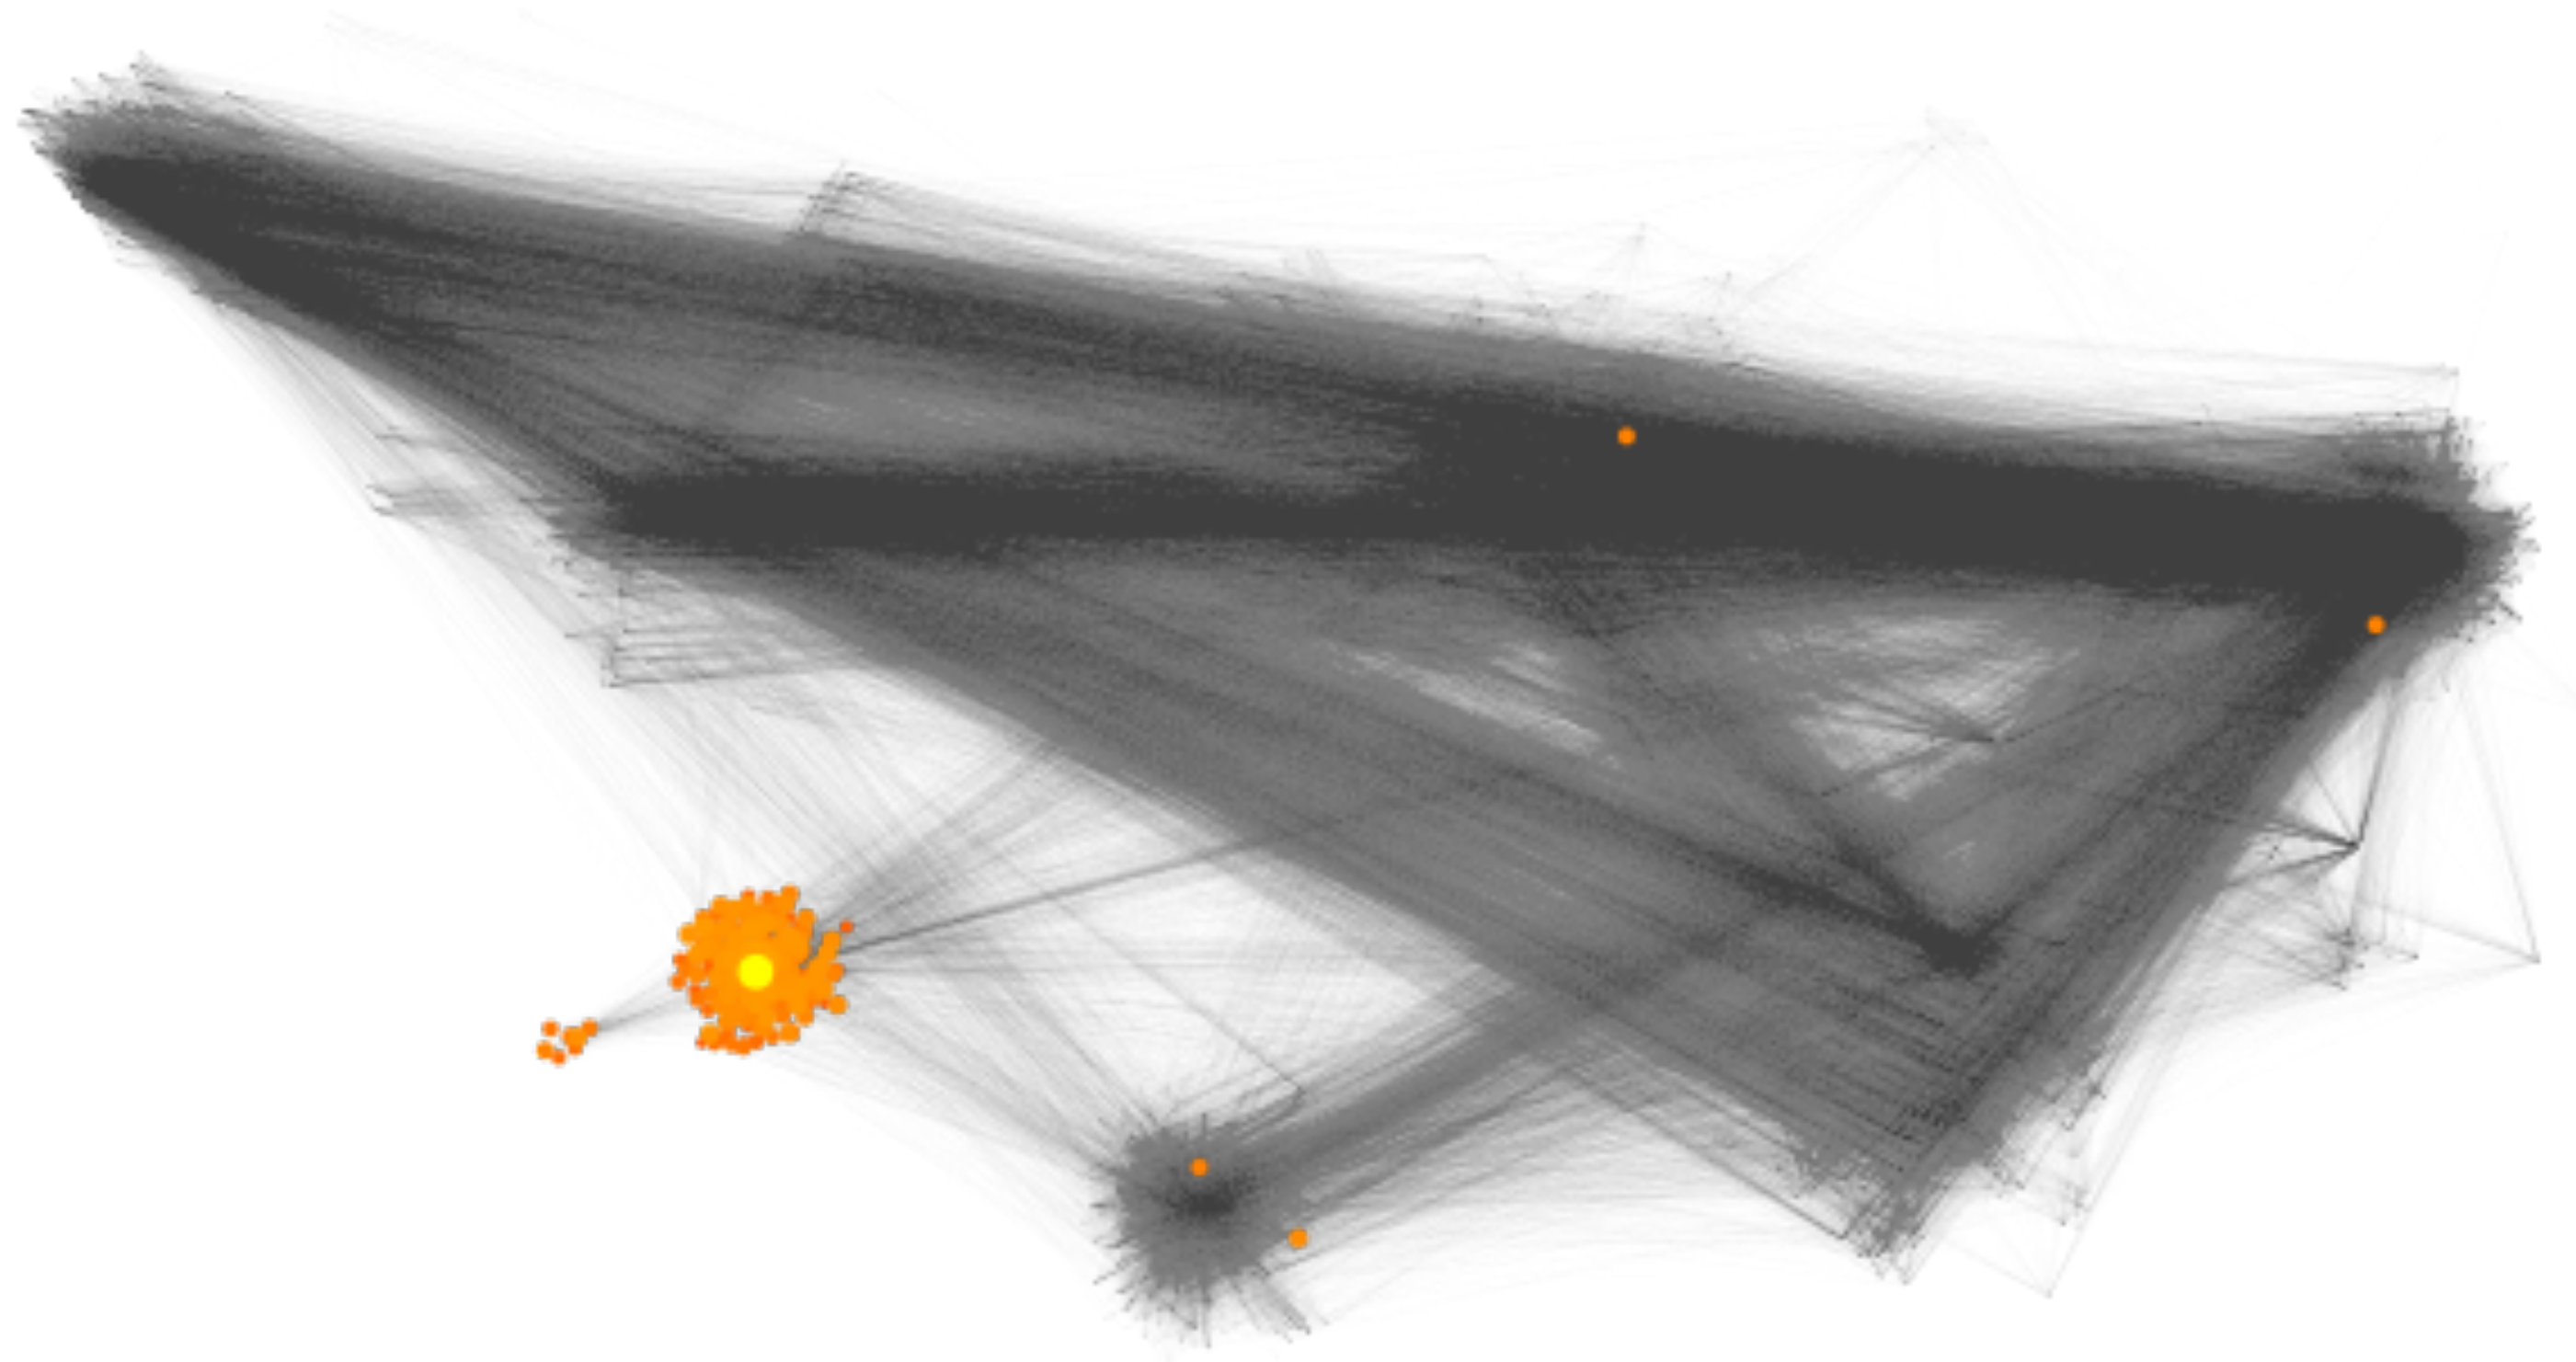
# Normalized cuts: finds 20% of the graph



Data: Facebook John Hopkins, A. L. Traud, P. J. Mucha and M. A. Porter, Physica A, 391(16), 2012

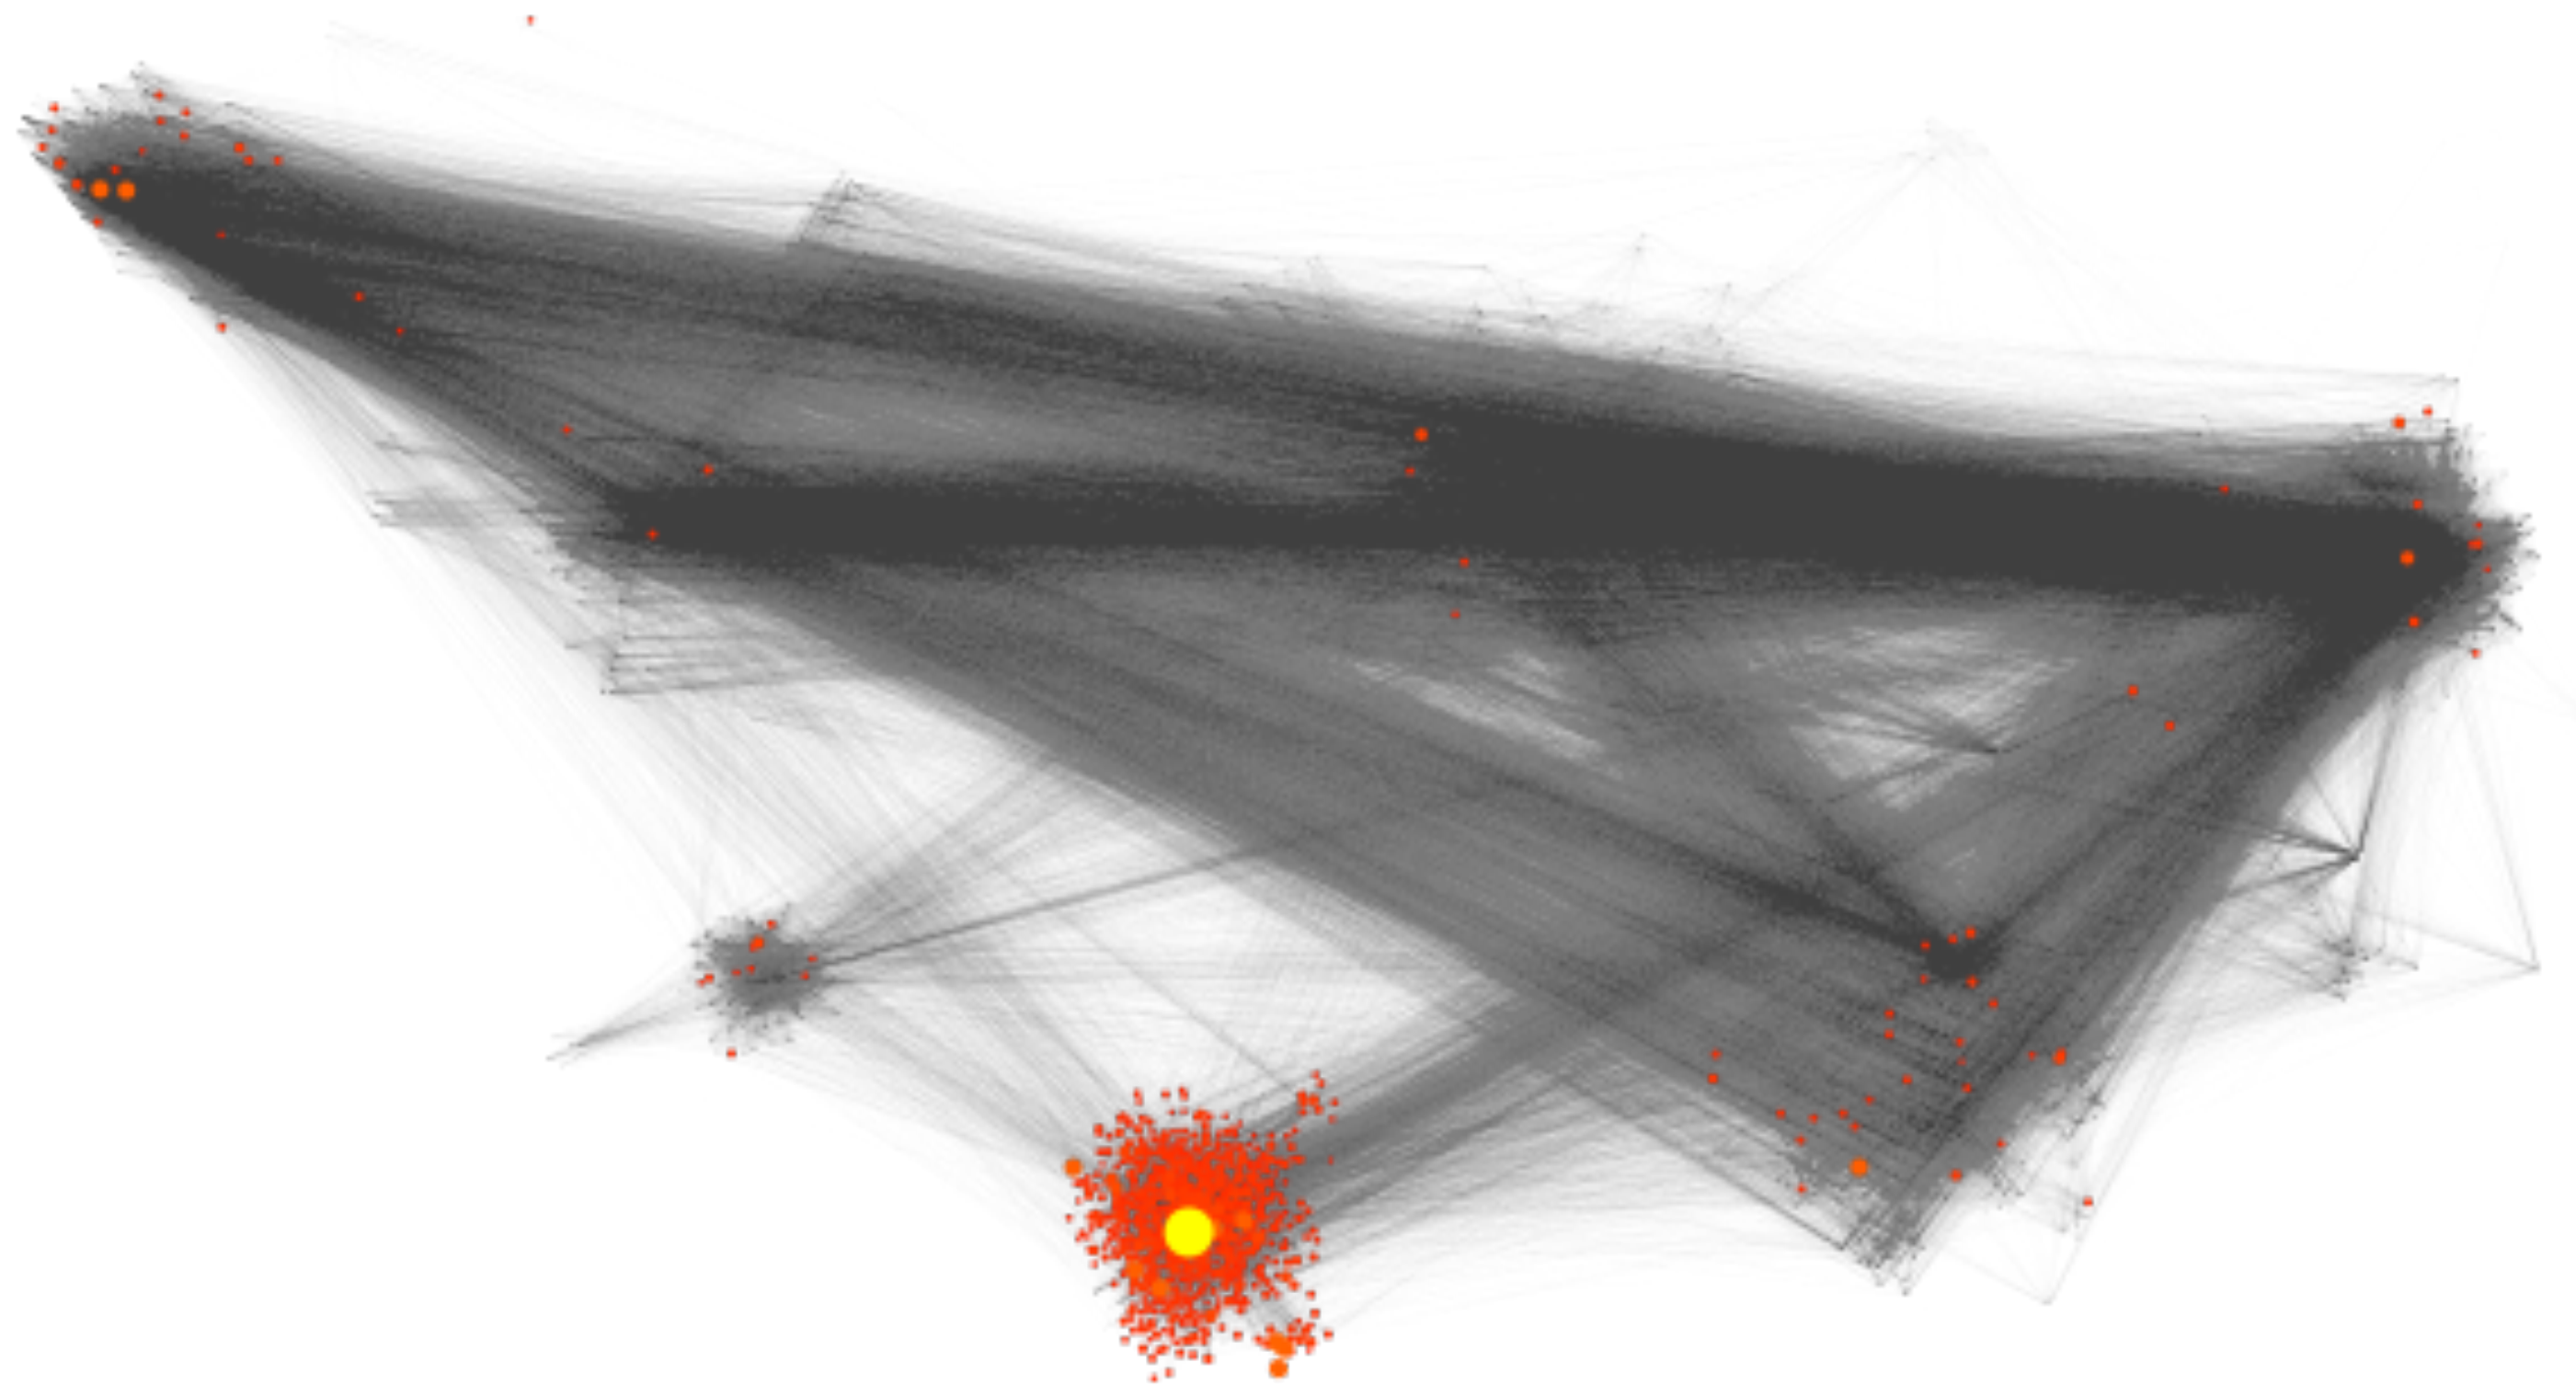


# Local graph clustering: finds 3% of the graph



Data: Facebook John Hopkins, A. L. Traud, P. J. Mucha and M. A. Porter, Physica A, 391(16), 2012

# Local graph clustering: finds 17% of the graph



Data: Facebook John Hopkins, A. L. Traud, P. J. Mucha and M. A. Porter, Physica A, 391(16), 2012

Current algorithms and running time

# Global, weakly and strongly local methods

## **Global methods:** $O(\text{volume of graph})$

- The workload depends on the size of the graph

## **Weakly local methods:** $O(\text{volume of graph})$

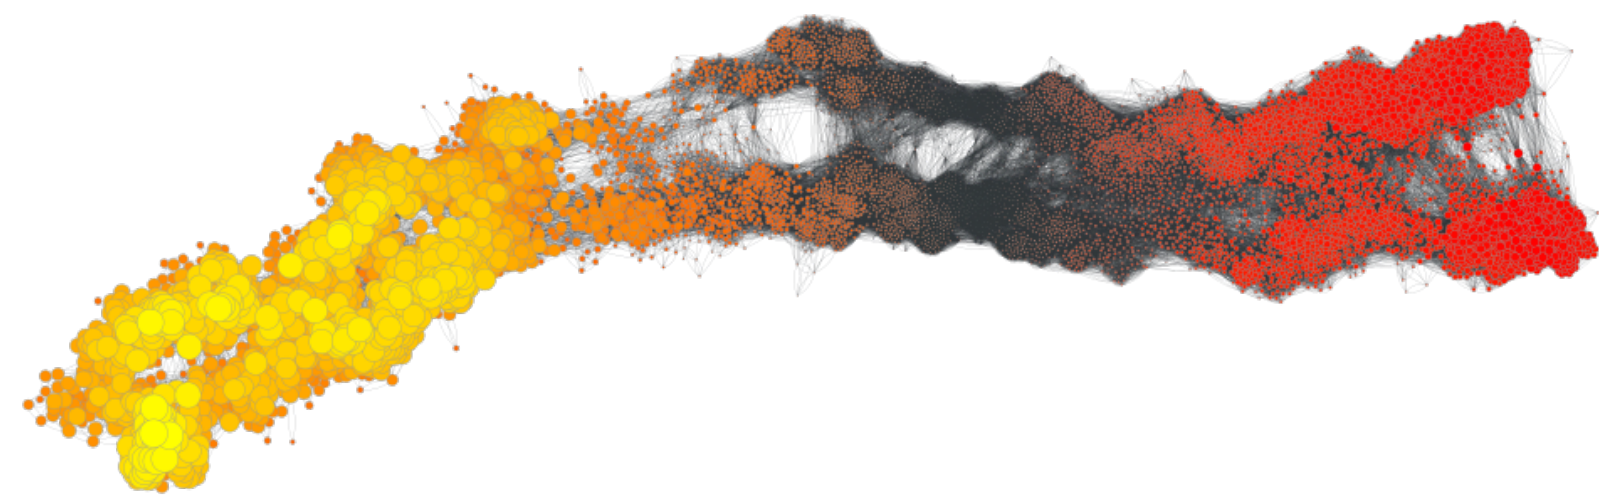
- A seed set of nodes is given
- The solution is locally biased to the input seed set
- The workload depends on the size of the graph

## **Strongly local methods:** $O(\text{volume of output cluster})$

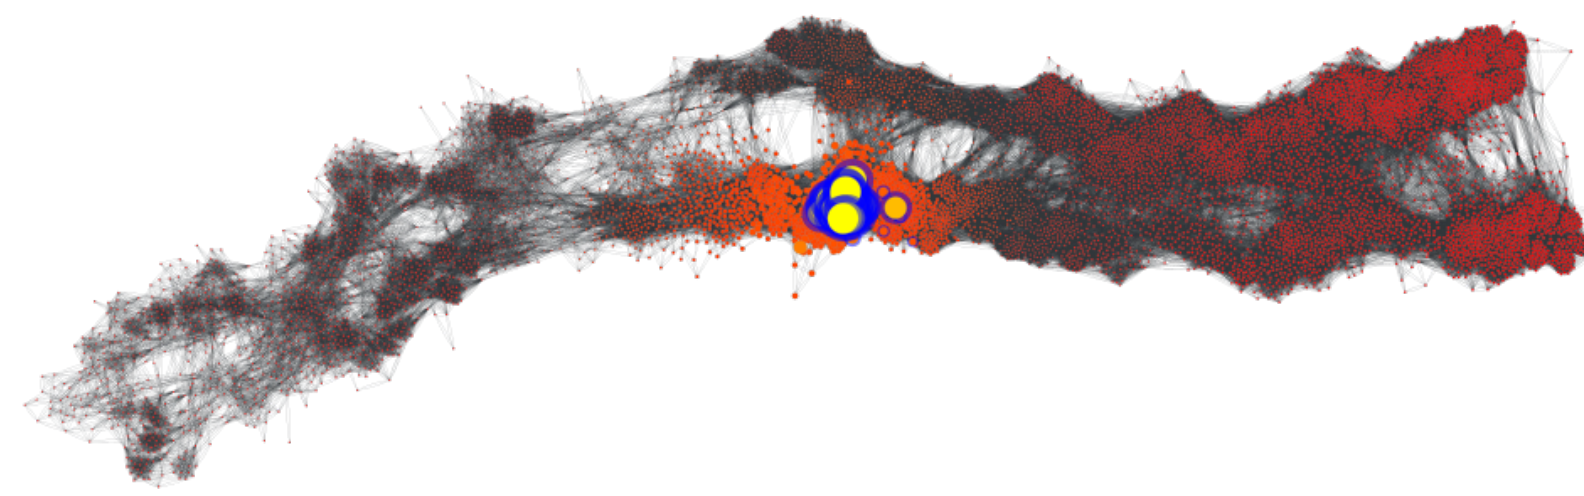
- A seed set of nodes is given
- The solution is locally biased to the input seed set



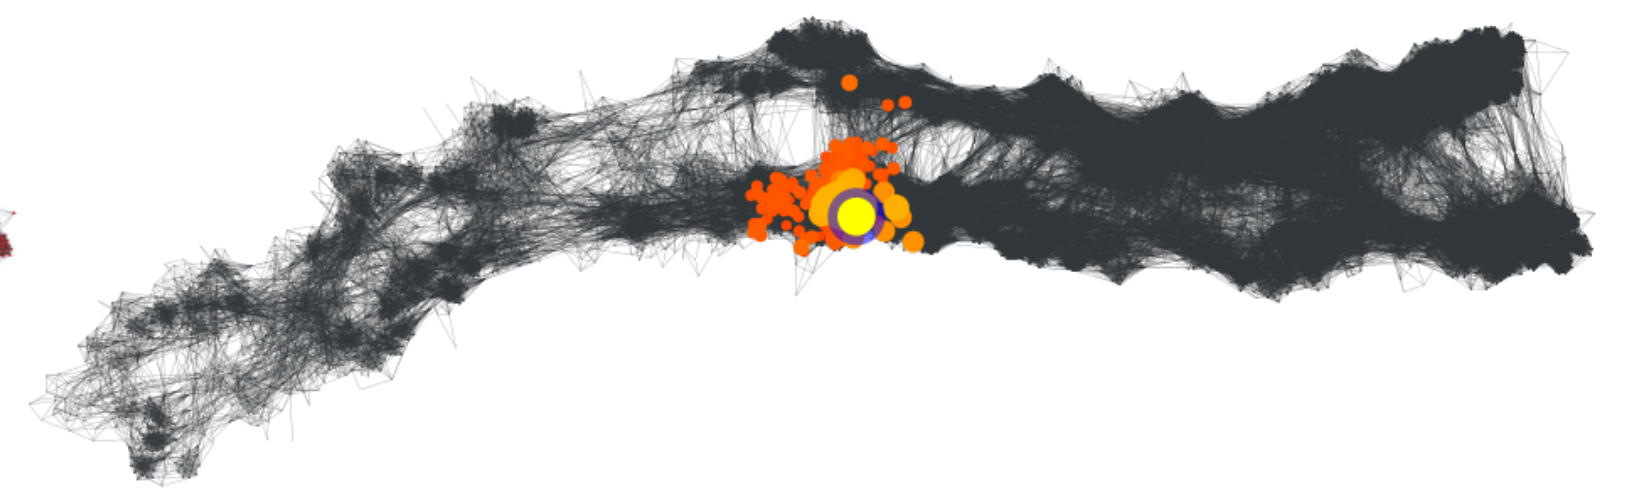
# Global, weakly and strongly local methods



Global



Weakly local

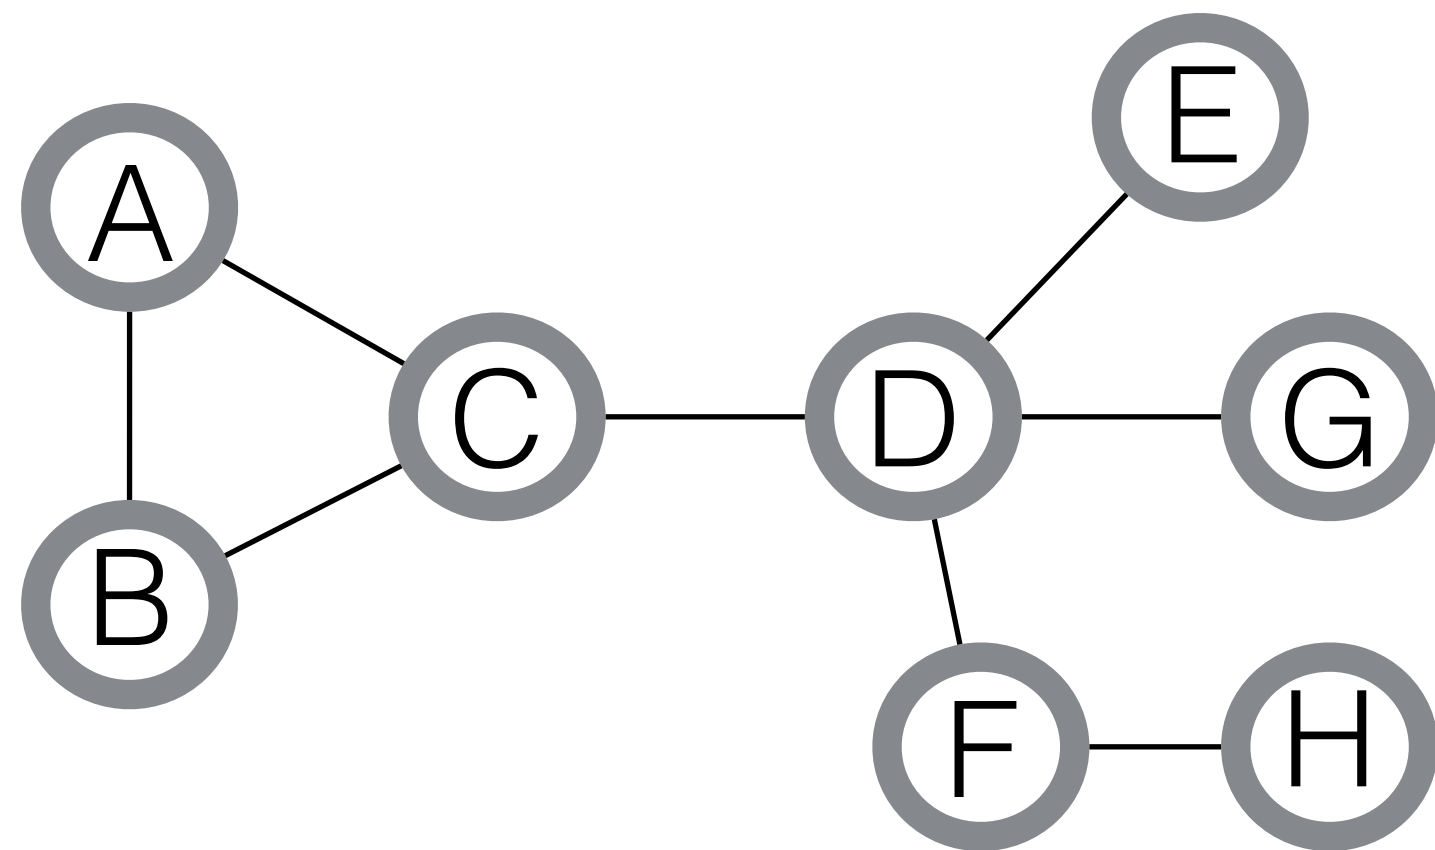


Strongly local

# Cluster quality

We measure cluster quality using

$$\text{Conductance} := \frac{\text{number of edges leaving cluster}}{\text{sum of degrees of vertices in cluster}}$$



$$\text{Conductance}(\{A,B\}) = 2/(2 + 2) = 1/2$$

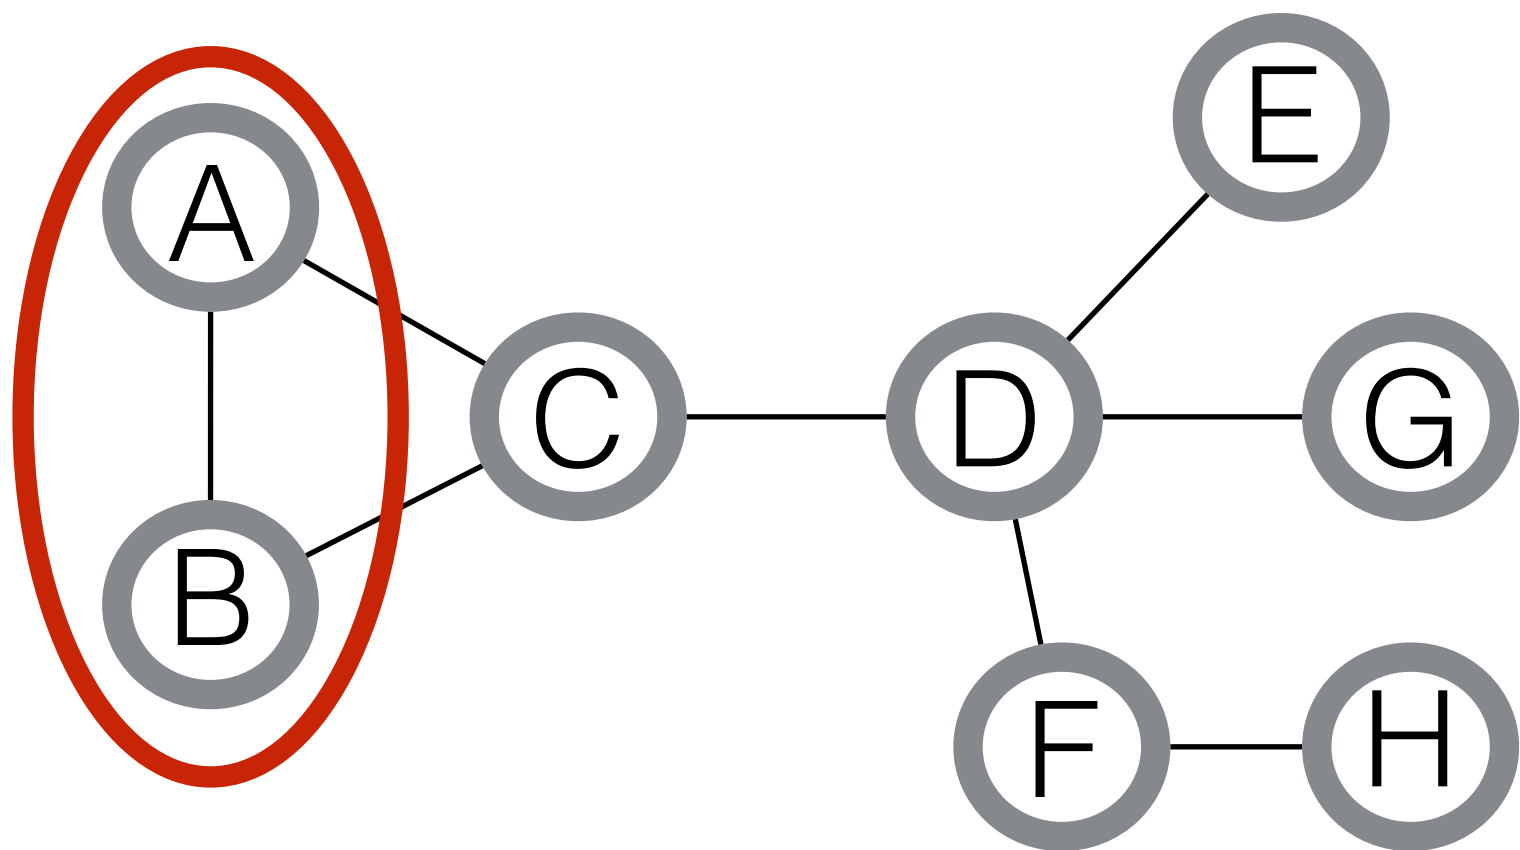
$$\text{Conductance}(\{A,B,C\}) = 1/(2 + 2 + 3) = 1/7$$

- The smaller the conductance value the better
- Minimizing conductance is NP-hard, we use approximation algorithms

# Cluster quality

We measure cluster quality using

$$\text{Conductance} := \frac{\text{number of edges leaving cluster}}{\text{sum of degrees of vertices in cluster}}$$



$$\text{Conductance}(\{\mathbf{A}, \mathbf{B}\}) = 2/(2 + 2) = 1/2$$

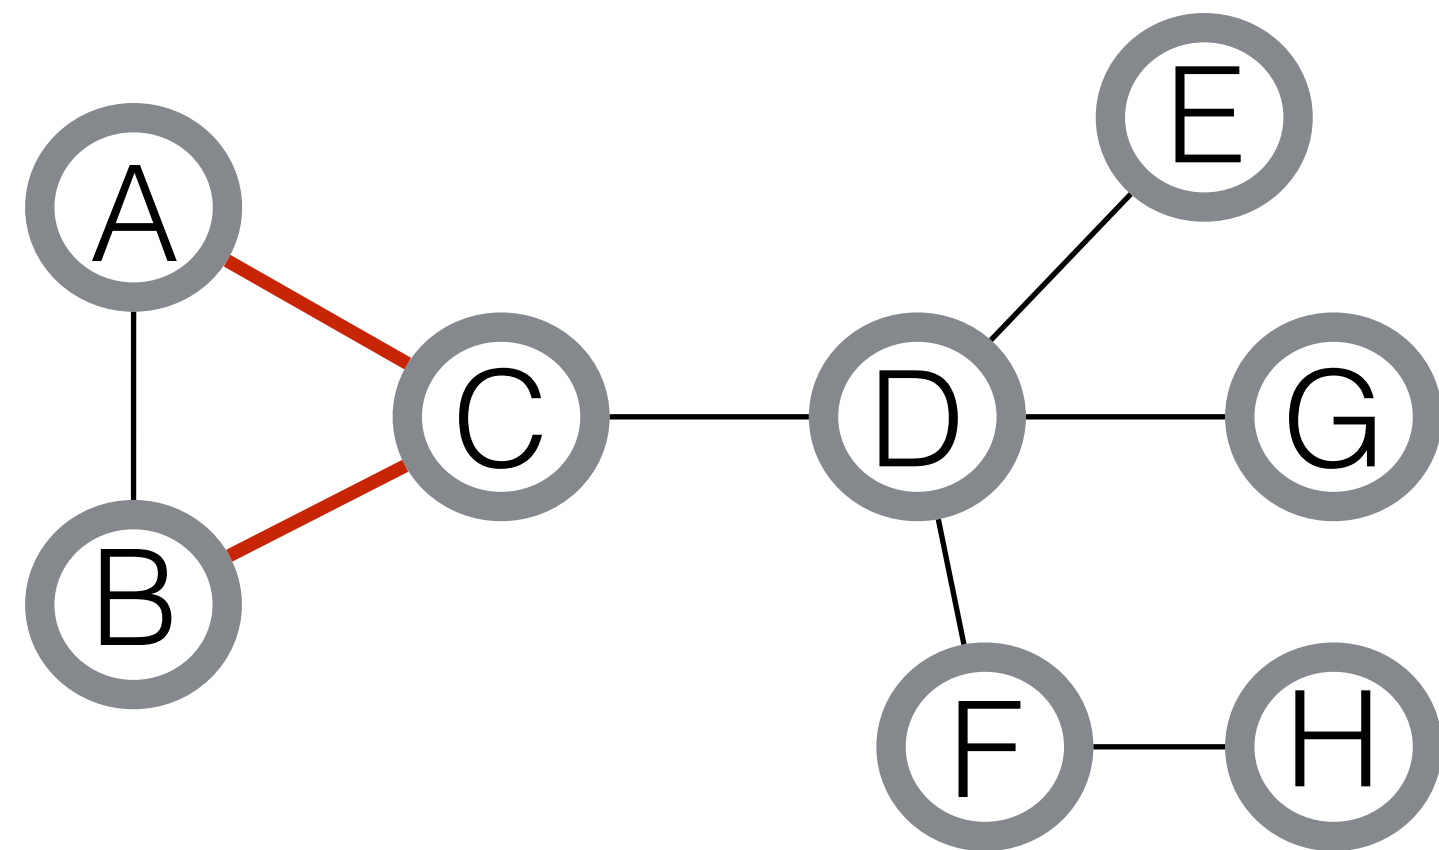
$$\text{Conductance}(\{A, B, C\}) = 1/(2 + 2 + 3) = 1/7$$

- The smaller the conductance value the better
- Minimizing conductance is NP-hard, we use approximation algorithms

# Cluster quality

We measure cluster quality using

$$\text{Conductance} := \frac{\text{number of edges leaving cluster}}{\text{sum of degrees of vertices in cluster}}$$



$$\text{Conductance}(\{A,B\}) = \mathbf{2}/(2 + 2) = 1/2$$

$$\text{Conductance}(\{A,B,C\}) = 1/(2 + 2 + 3) = 1/7$$

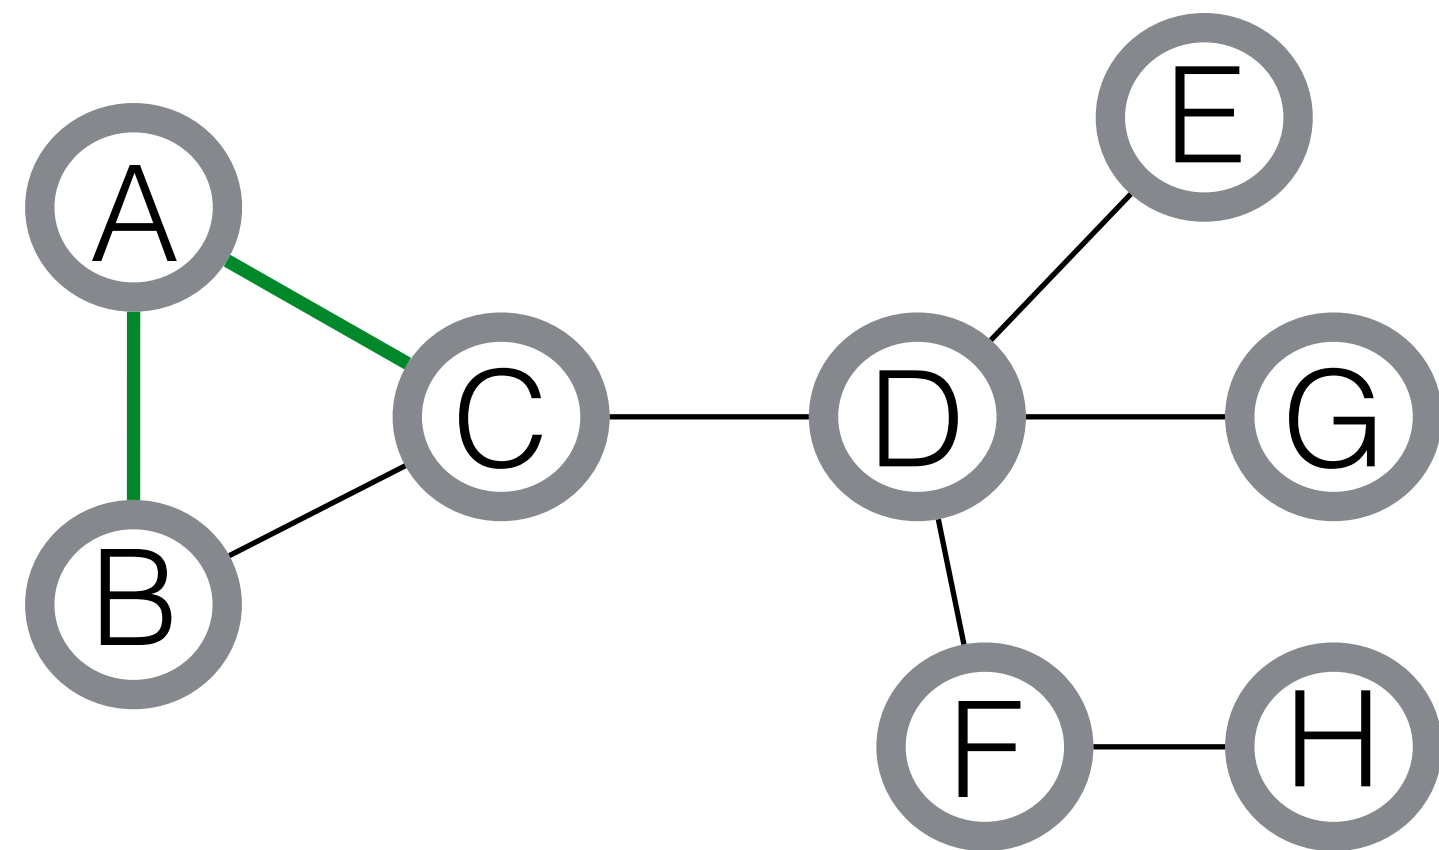
- The smaller the conductance value the better
- Minimizing conductance is NP-hard, we use approximation algorithms



# Cluster quality

We measure cluster quality using

$$\text{Conductance} := \frac{\text{number of edges leaving cluster}}{\text{sum of degrees of vertices in cluster}}$$



$$\text{Conductance}(\{A, B\}) = 2/(\mathbf{2} + 2) = 1/2$$

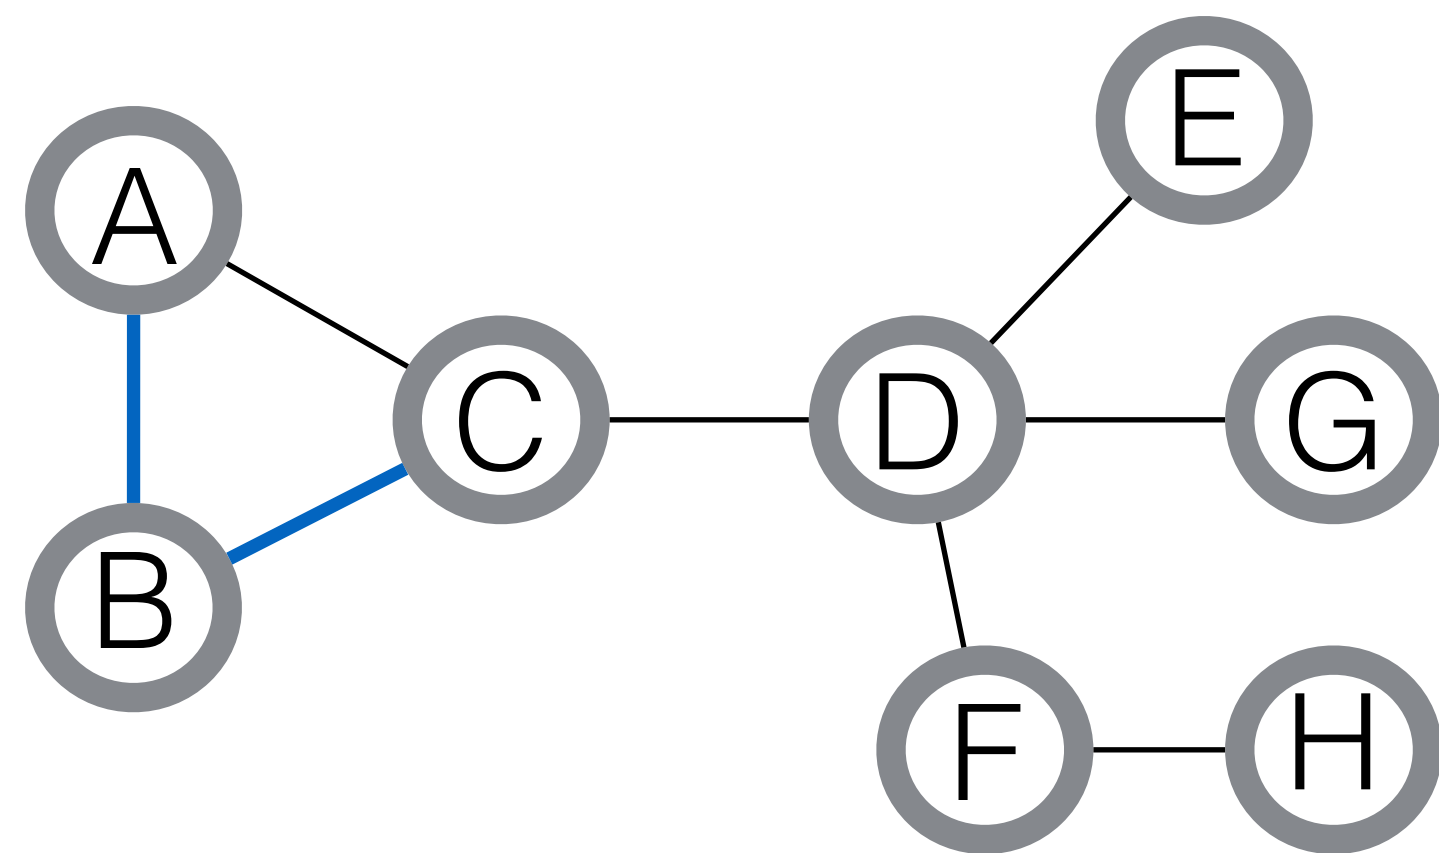
$$\text{Conductance}(\{A, B, C\}) = 1/(2 + 2 + 3) = 1/7$$

- The smaller the conductance value the better
- Minimizing conductance is NP-hard, we use approximation algorithms

# Cluster quality

We measure cluster quality using

$$\text{Conductance} := \frac{\text{number of edges leaving cluster}}{\text{sum of degrees of vertices in cluster}}$$



$$\text{Conductance}(\{A,B\}) = 2/(2 + \mathbf{2}) = 1/2$$

$$\text{Conductance}(\{A,B,C\}) = 1/(2 + 2 + 3) = 1/7$$

- The smaller the conductance value the better
- Minimizing conductance is NP-hard, we use approximation algorithms

# Local graph clustering methods

- **MQI (strongly local):** Lang and Rao, 2004
- **Approximate Page Rank (strongly local):** Andersen, Chung, Lang, 2006
- **spectral MQI (strongly local):** Chung, 2007
- **Flow-Improve (weakly local):** Andersen and Lang, 2008
- **MOV (weakly local):** Mahoney, Orecchia, Vishnoi, 2012
- **Nibble (strongly local):** Spielman and Teng, 2013
- **Local Flow-Improve (strongly local):** Orecchia, Zhu, 2014
- **Deterministic HeatKernel PR (strongly local):** Kloster, Gleich, 2014
- **Randomized HeatKernel PR (strongly local):** Chung, Simpson, 2015
- **Sweep cut rounding algorithm**

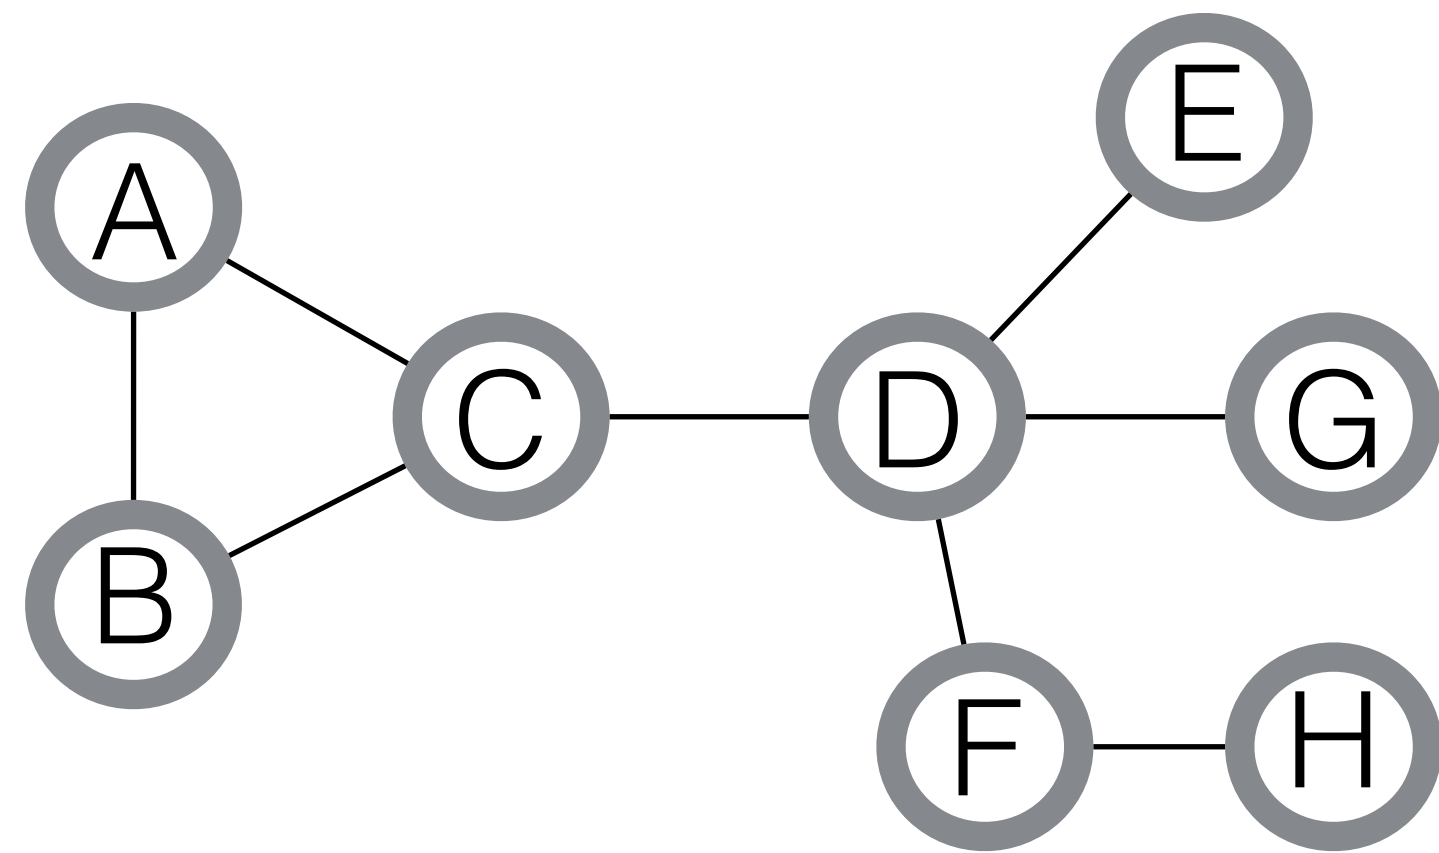
# Shared memory parallel methods

- **We parallelize 4 strongly local spectral methods + rounding**
  1. Approximate Page Rank ←--- this talk
  2. Nibble
  3. Deterministic HeatKernel Approximate Page-Rank
  4. Randomized HeatKernel Approximate Page-Rank
  5. Sweep cut rounding algorithm ←--- this talk
- **All local methods take various parameters**
  - Parallel method 1: try different parameters independently in parallel
  - Parallel method 2: **parallelize algorithm for individual run**
    - Useful for **interactive** setting where tweaking of parameters is needed



# Approximate Page-Rank

# Personalized Page-Rank vector



Adjacency matrix A

	A	B	C	D	E	F	G	H
A		1	1					
B	1		1					
C	1	1		1				
D			1		1	1	1	
E				1				
F				1				1
G				1				
H						1		

Degree matrix D

	A	B	C	D	E	F	G	H
A	2							
B		2						
C			3					
D				4				
E					1			
F						2		
G							1	
H								1

Pick a vertex  $u$  of interest and define a vector:

$$s[u] = 1, \quad s[v] = 0 \quad \forall v \neq u$$

a teleportation parameter  $0 \leq \alpha \leq 1$  and  $W = AD^{-1}$  then the PPR vector is given by solving:

$$((1 - \alpha)W + \alpha se^T)p = p \quad \Leftrightarrow \quad (I - (1 - \alpha)W)p = \alpha s$$

# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006

Algorithm idea: iteratively spread probability mass from vector  $s$  around the graph.

- $r$  is the residual vector,  $p$  is the solution vector
- $\rho > 0$  is tolerance parameter

Run a coordinate descent solver for PPR until: any vertex  $u$  satisfies  $r[u] \geq -\alpha d[u]$

**Initialize:  $p = 0$ ,  $r = -\alpha s$**

**While termination criterion is not met do**

**1. Choose any vertex  $u$  where  $r[u] < -\alpha d[u]$**

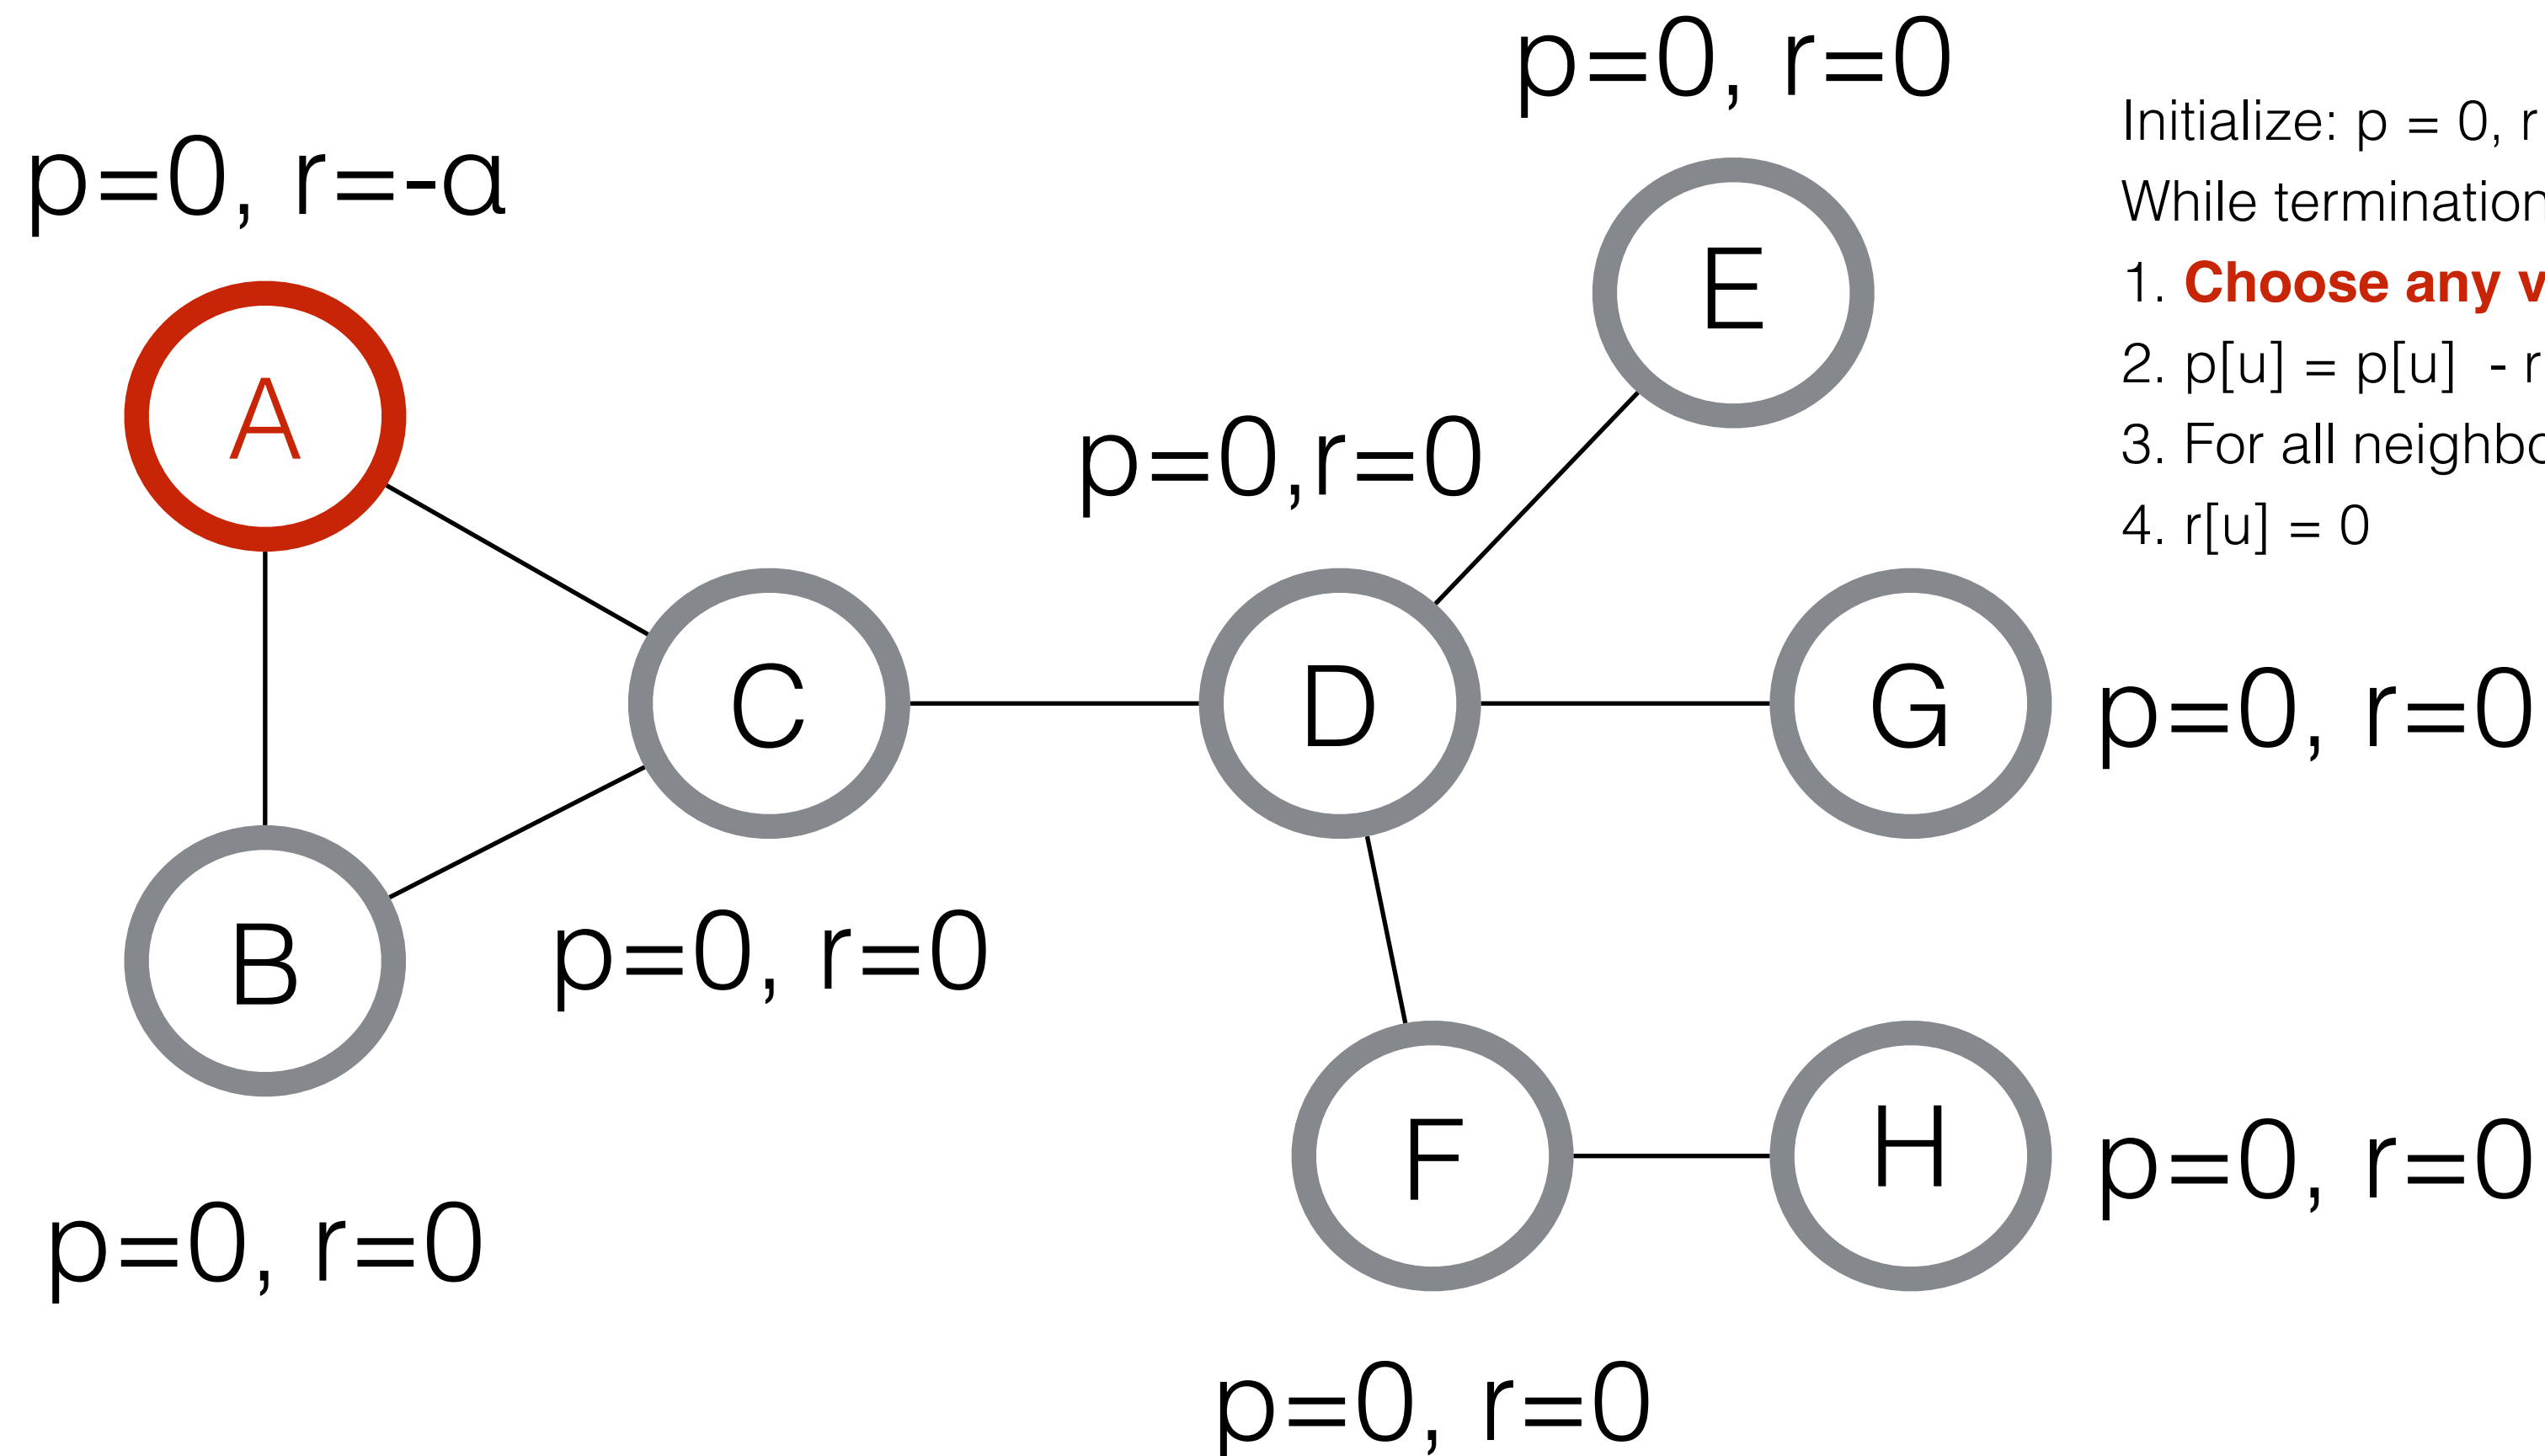
**2.  $p[u] = p[u] - r[u]$**

residual update { **3. For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)r[u]A[u,v]/d[u]$**   
**4.  $r[u] = 0$**

Final step: round the solution  $p$  using sweep cut.

# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006



Initialize:  $p = 0, r = -\alpha s$

While termination criterion is not met do

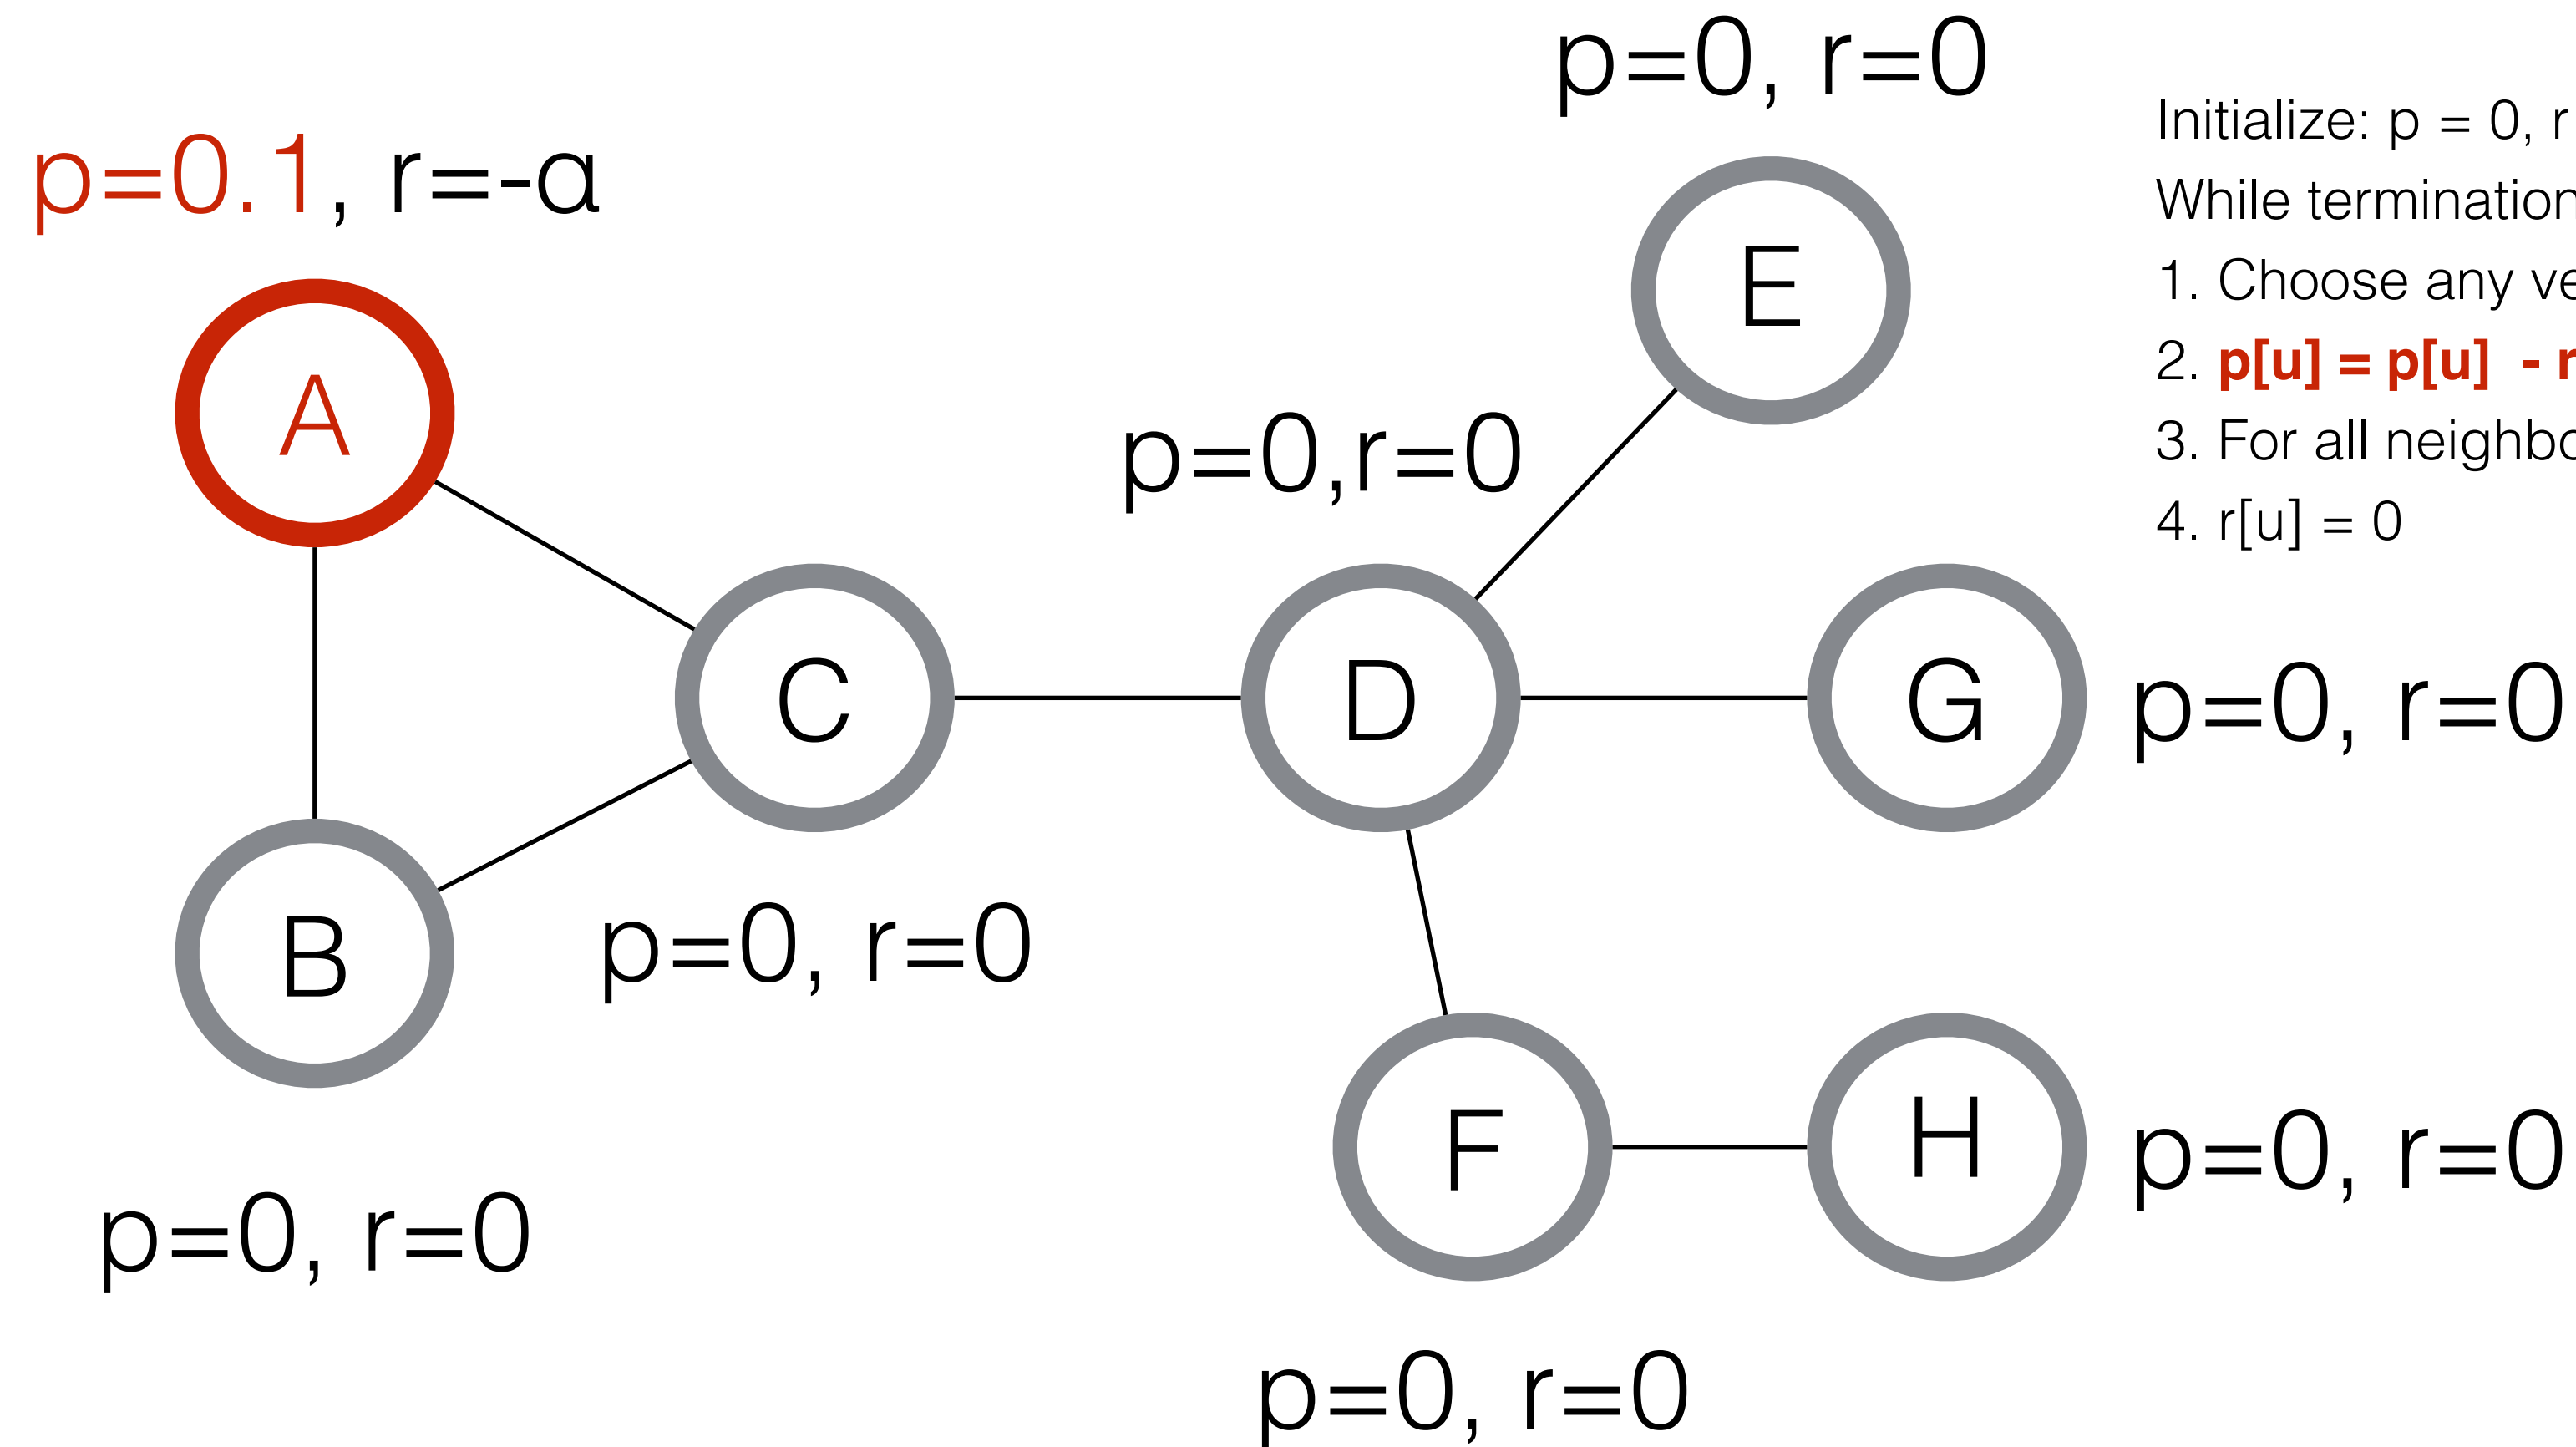
1. **Choose any vertex  $u$  where  $r[u] < -\alpha p d[u]$**
2.  $p[u] = p[u] - r[u]$
3. For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)r[u]A[u,v]/d[u]$
4.  $r[u] = 0$

$$\frac{\|r\|_1}{\alpha} = 1, \quad \|p\|_1 = 0, \quad \frac{\|r\|_1}{\alpha} + \|p\|_1 = 1$$



# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006



Initialize:  $p = 0, r = -\alpha s$

While termination criterion is not met do

1. Choose any vertex  $u$  where  $r[u] < -\alpha p d[u]$

2.  $p[u] = p[u] - r[u]$

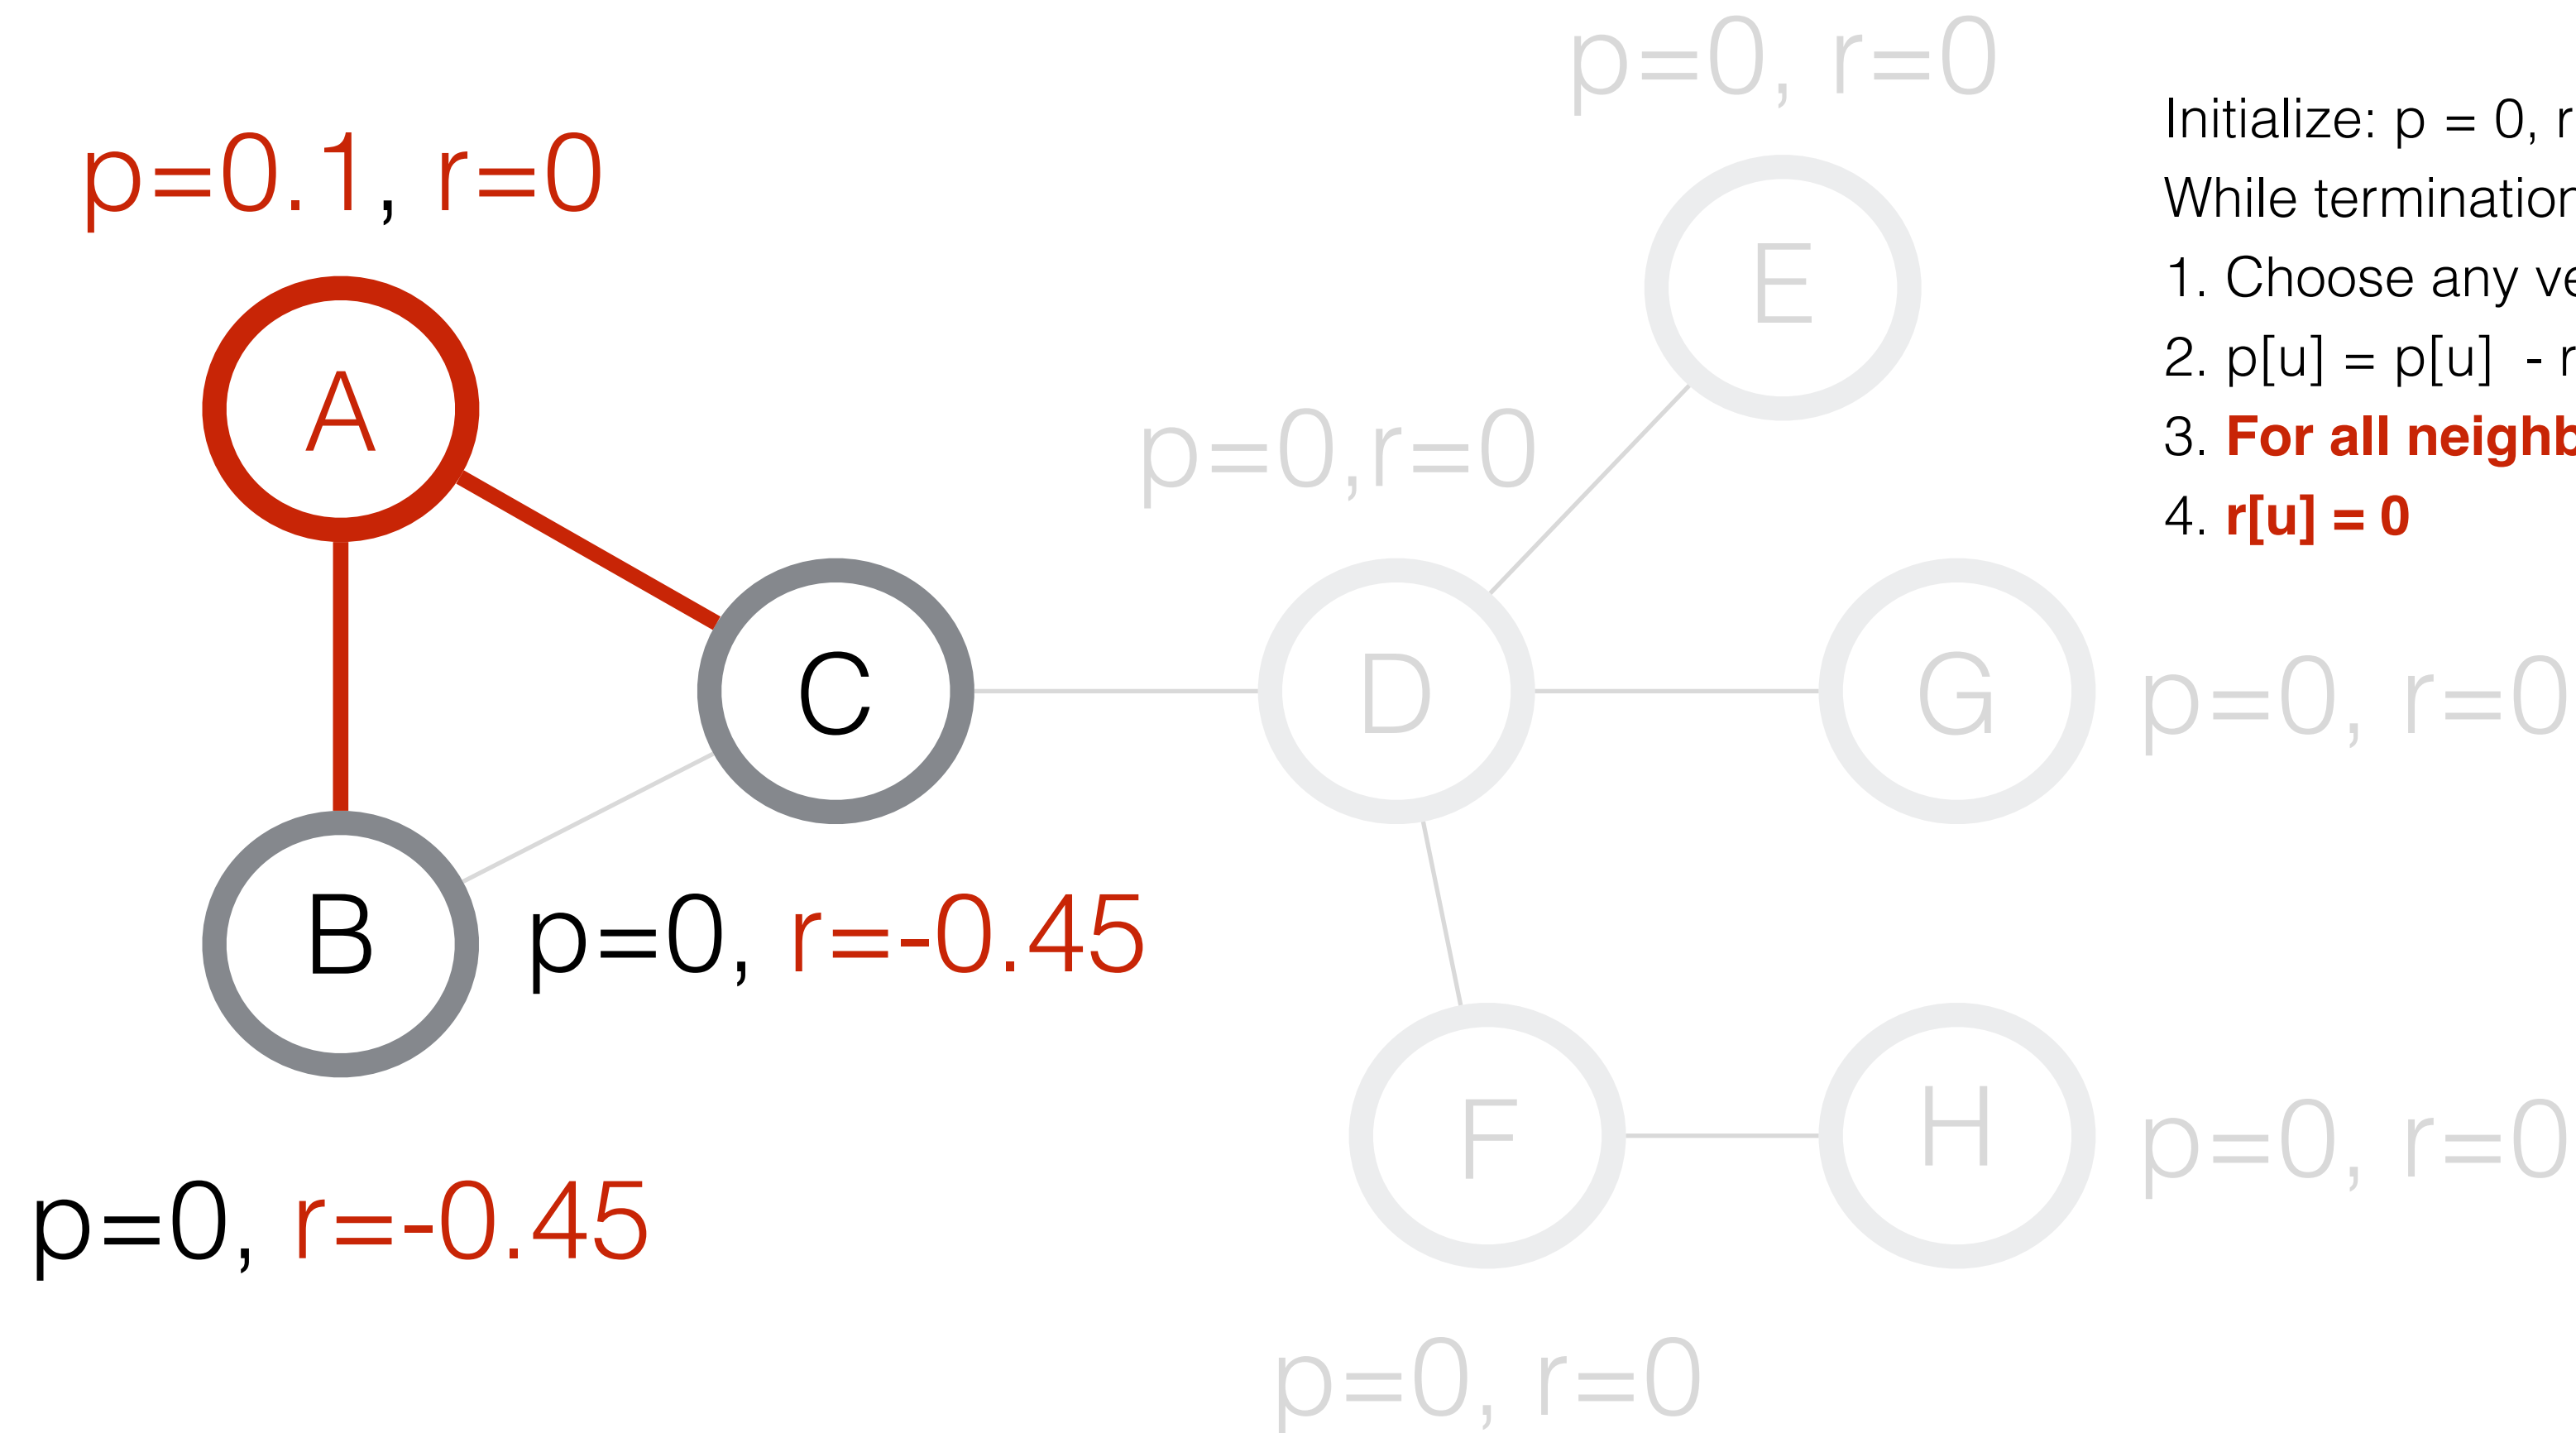
3. For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)r[u]A[u,v]/d[u]$

4.  $r[u] = 0$

$$\frac{\|r\|_1}{\alpha} = 1, \quad \|p\|_1 = 0.1, \quad \frac{\|r\|_1}{\alpha} + \|p\|_1 = 1.1$$

# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006



Initialize:  $p = 0, r = -\alpha s$ , where  $s$  is a probability vector

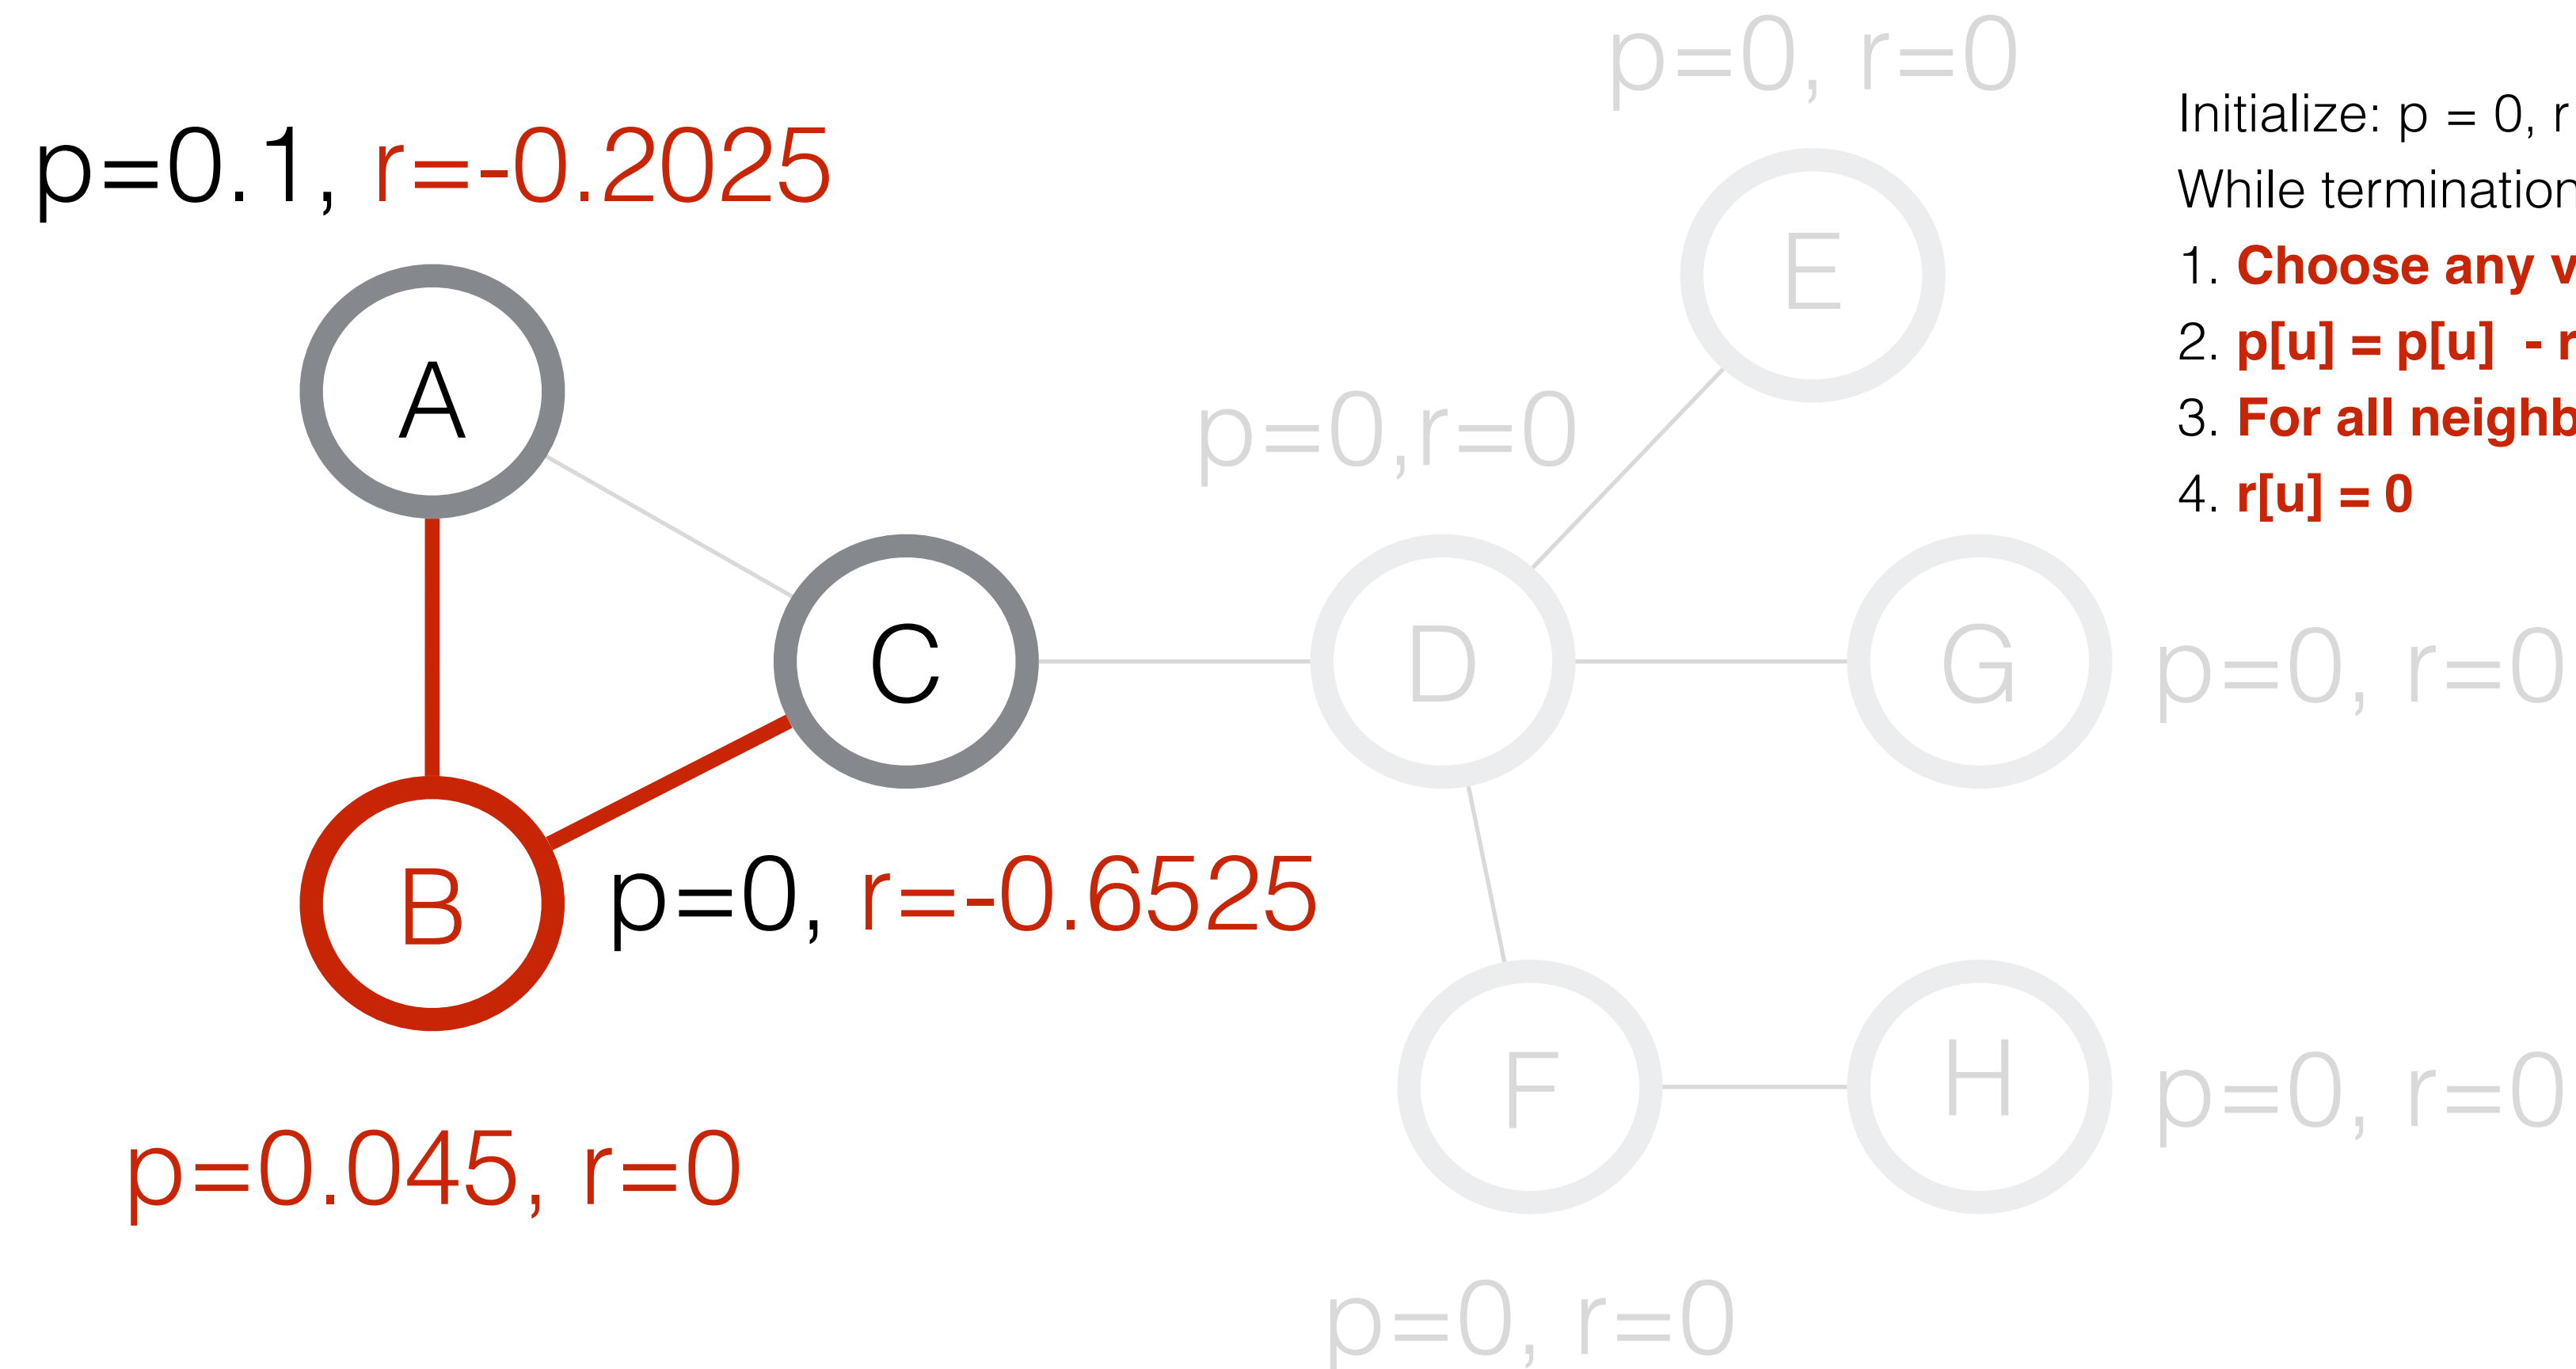
While termination criterion is not met do

1. Choose any vertex  $u$  where  $r[u] < -\alpha p[u]$
2.  $p[u] = p[u] - r[u]$
3. **For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)r[u]A[u,v]/d[u]$**
4.  **$r[u] = 0$**

$$\frac{\|r\|_1}{\alpha} = 0.9, \quad \|p\|_1 = 0.1, \quad \frac{\|r\|_1}{\alpha} + \|p\|_1 = 1.0$$

# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006



Initialize:  $p = 0, r = -\alpha s$

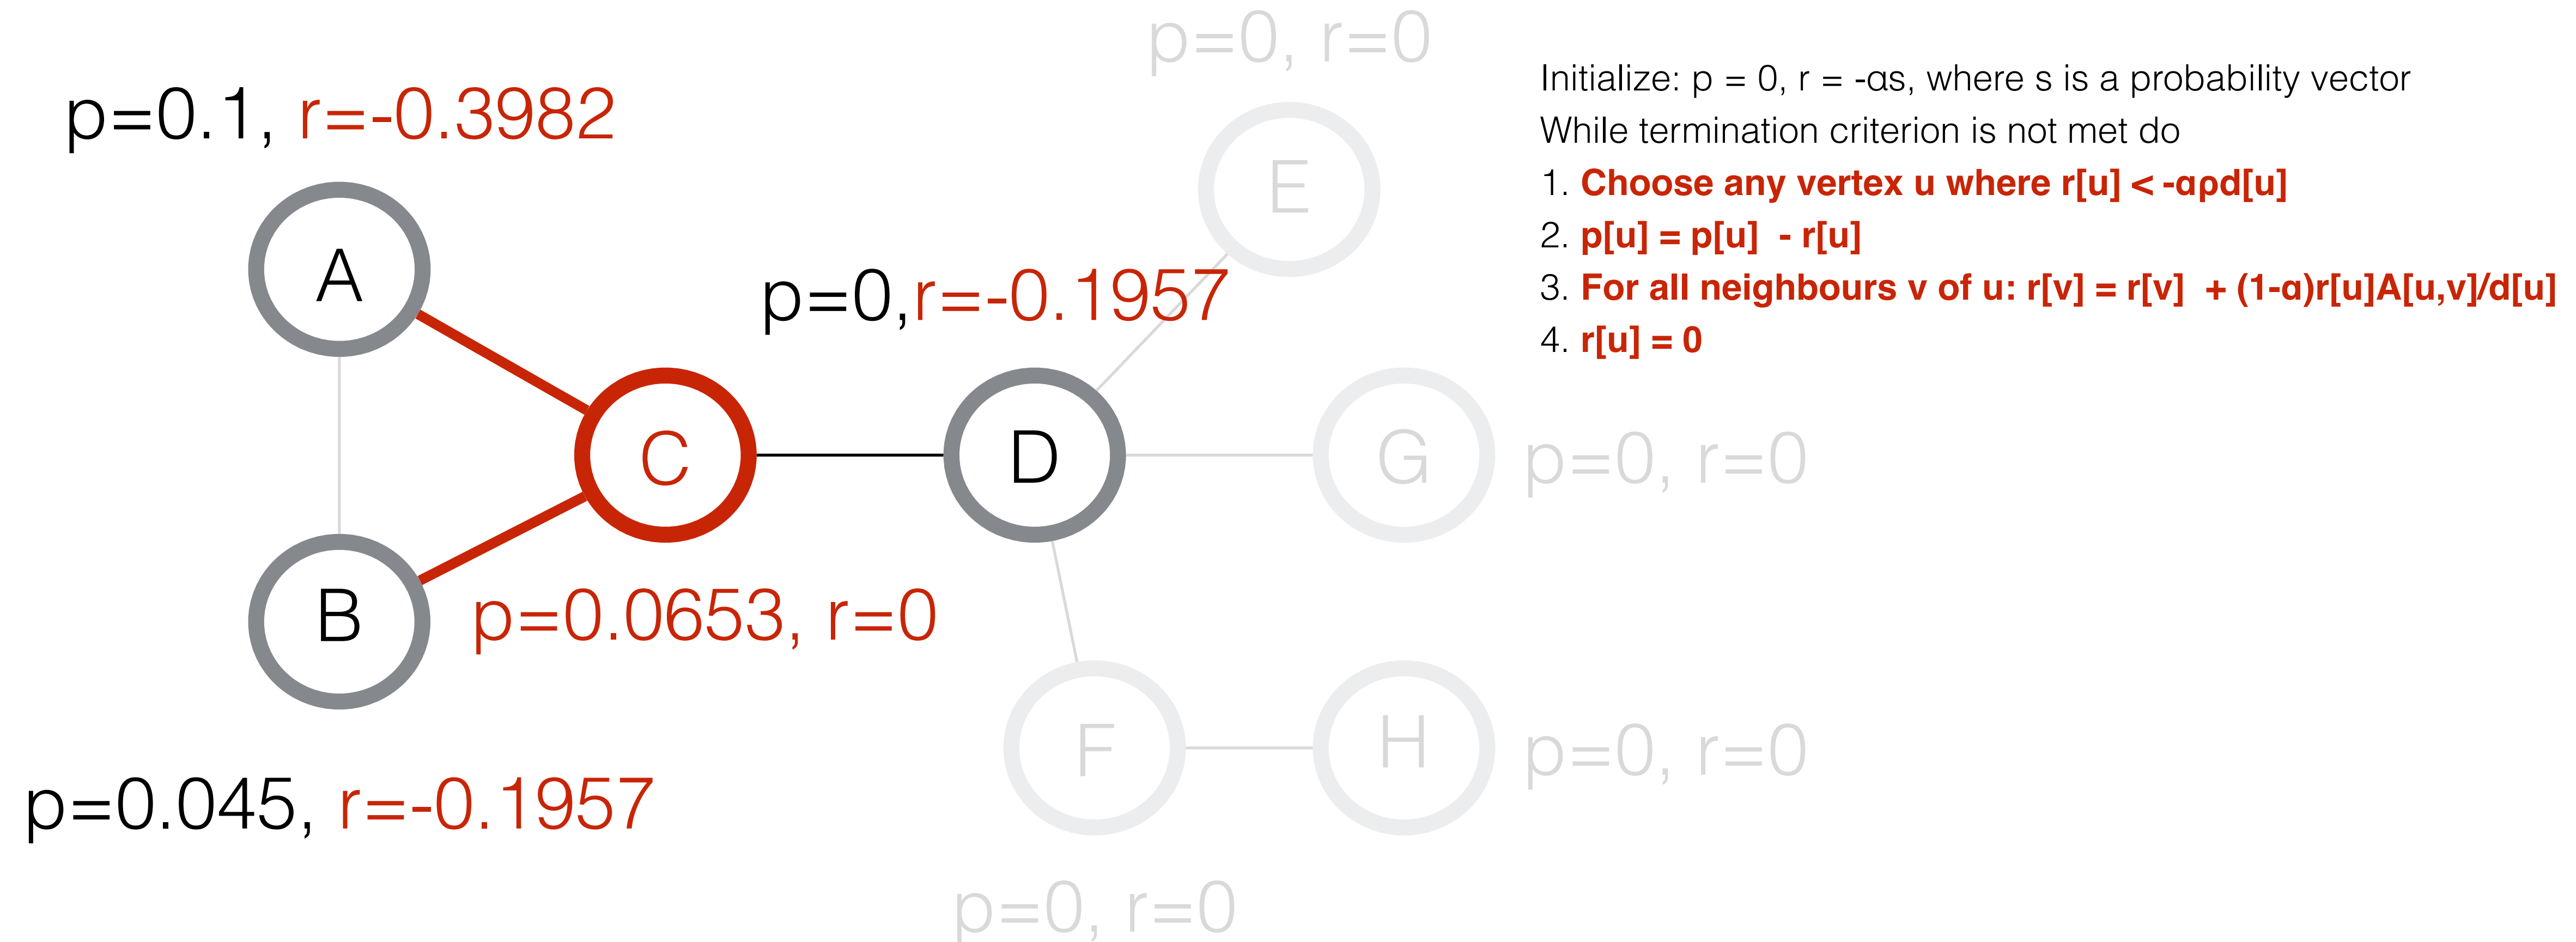
While termination criterion is not met do

1. **Choose any vertex  $u$  where  $r[u] < -\alpha p d[u]$**
2.  **$p[u] = p[u] - r[u]$**
3. **For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)r[u]A[u,v]/d[u]$**
4.  **$r[u] = 0$**

$$\frac{\|r\|_1}{\alpha} = 0.855, \quad \|p\|_1 = 0.145, \quad \frac{\|r\|_1}{\alpha} + \|p\|_1 = 1.0$$

# Approximate Personalized Page-Rank

R. Andersen, F. Chung and K. Lang. Local graph partitioning using Page-Rank, FOCS, 2006



$$\frac{\|r\|_1}{\alpha} = 0.7897, \quad \|p\|_1 = 0.2103, \quad \frac{\|r\|_1}{\alpha} + \|p\|_1 = 1.0$$



# Running time APPR

- At each iteration APPR touches a single node and its neighbours
  - Let  $\text{supp}(p)$  be the support of vector  $p$  at termination which satisfies  **$\text{vol}(\text{supp}(p)) \leq 1/(\alpha p)$**
  - **Overall until termination the work is:  $O(1/(\alpha p))$**  [Andersen, Chung, Lang, FOCS, 2006]
- We store vectors  $p$  and  $r$  using sparse sets
  - **We can only afford to do work proportional to nodes and edges currently touched**
  - We used *unordered\_map* data structure in STL (Standard Template Library)
  - Guarantees  $O(1/(\alpha p))$  work

# L1-regularized Page-Rank

APPR is an approximation algorithm but what is it minimizing?

$$\text{minimize } \frac{1-\alpha}{2} \|Bp\|_2^2 + \alpha \|H(\mathbf{1} - p)\|_2^2 + \alpha \|Zp\|_2^2 + \rho\alpha \|Dp\|_1$$

where

- B: is the incidence matrix
- Z, H: are diagonal scaling matrices

Incidence matrix B								
	A	B	C	D	E	F	G	H
A-B	1	-1						
A-C	1		-1					
B-C		1	-1					
C-D			1	-1				
D-E				1	-1			
D-F				1		-1		
D-G				1			-1	
F-H						1		-1

KF, X. Cheng, J. Shun, F. Roosta-Khorasani, M. Mahoney. Exploiting optimization for local graph clustering, arXiv:1602.01886v1.

KF, D. Gleich, M. Mahoney. An optimization approach to locally-biased graph algorithms, arXiv:1607.04940.

Shared memory

# Running time: work depth model

Work depth model: J. Jaja. Introduction to parallel algorithms. Addison-Wesley Profesional, 1992

Note that our results are not model dependent.

## Model

- Work: number of operations required
- Depth: longest chain of sequential dependencies

Let  $P$  be the number of cores available.

By Brent's theorem [1] **an algorithm with work  $W$  and depth  $D$  has overall running time:  $W/P + D$ .**

In practice  $W/P$  dominates. Thus parallel efficient algorithms require the same work as its sequential version.

# Parallel Approximate Personalized Page-Rank

While termination criterion is not met do

1. Choose **ALL (instead of any)** vertex  $u$  where  $r[u] < \alpha \deg[u]$
2.  $p[u] = p[u] - r[u]$
3. For all neighbours  $v$  of  $u$ :  $r[v] = r[v] + (1-\alpha)/(2\deg[u])r[u]$
4.  $r[u] = (1-\alpha)r[u]/2$

- **Asymptotic work remains the same:**  $O(1/(\alpha p))$ .
- **Parallel randomized implementation: work  $O(1/(\alpha p))$  and depth  $O(\log(1/(\alpha p)))$ .**
  - Keep track of two **sparse** copies of  $p$  and  $r$
  - Concurrent hash table for sparse sets  $\leftarrow$  important for  $O(1/(\alpha p))$  work
  - Use atomic increment to deal with conflicts
  - Use of Ligra (Shun and Blelloch 2013) to process only “active” vertices and their edges
- **Same theoretical graph clustering guarantees**, Fountoulakis et al. 2016.

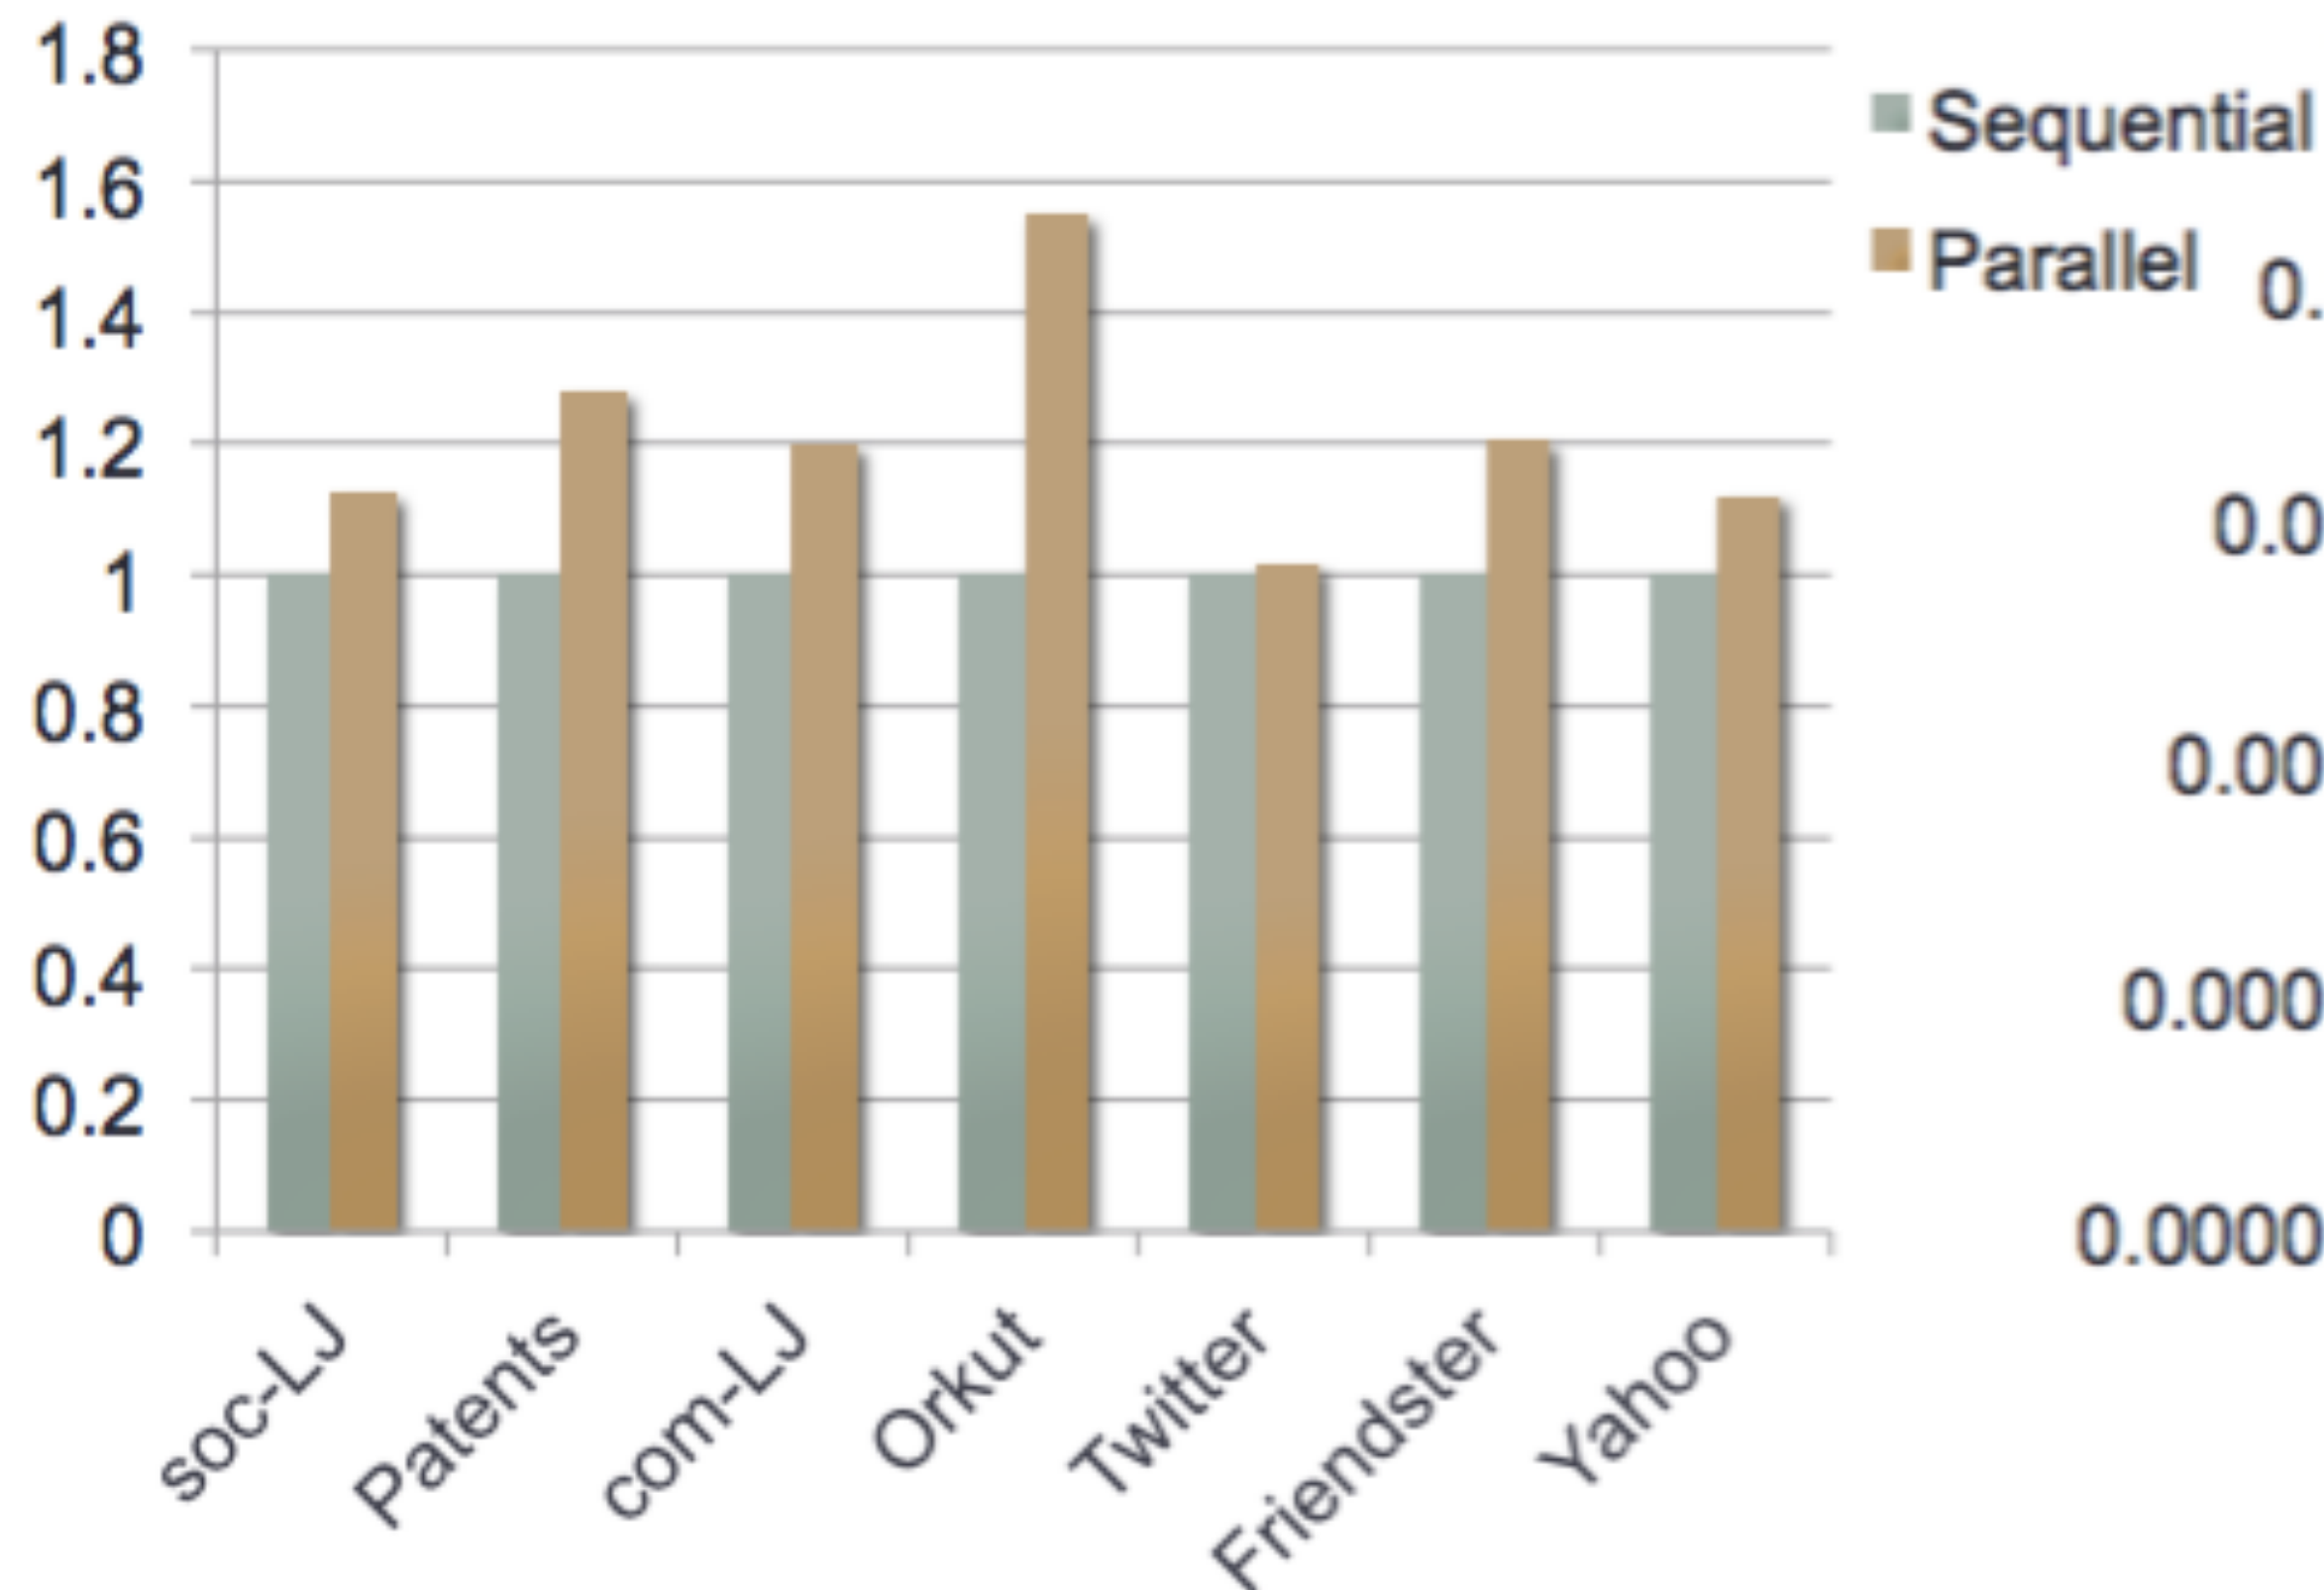
# Data

Input graph	Num. vertices	Num. edges
<b>soc-JL</b>	4,847,571	42,851,237
<b>cit-Patents</b>	6,009,555	16,518,947
<b>com-LJ</b>	4,036,538	34,681,189
<b>com-Orkut</b>	3,072,627	117,185,083
<b>Twitter</b>	41,652,231	1,202,513,046
<b>Friendster</b>	124,836,180	1,806,607,135
<b>Yahoo</b>	1,413,511,391	6,434,561,035

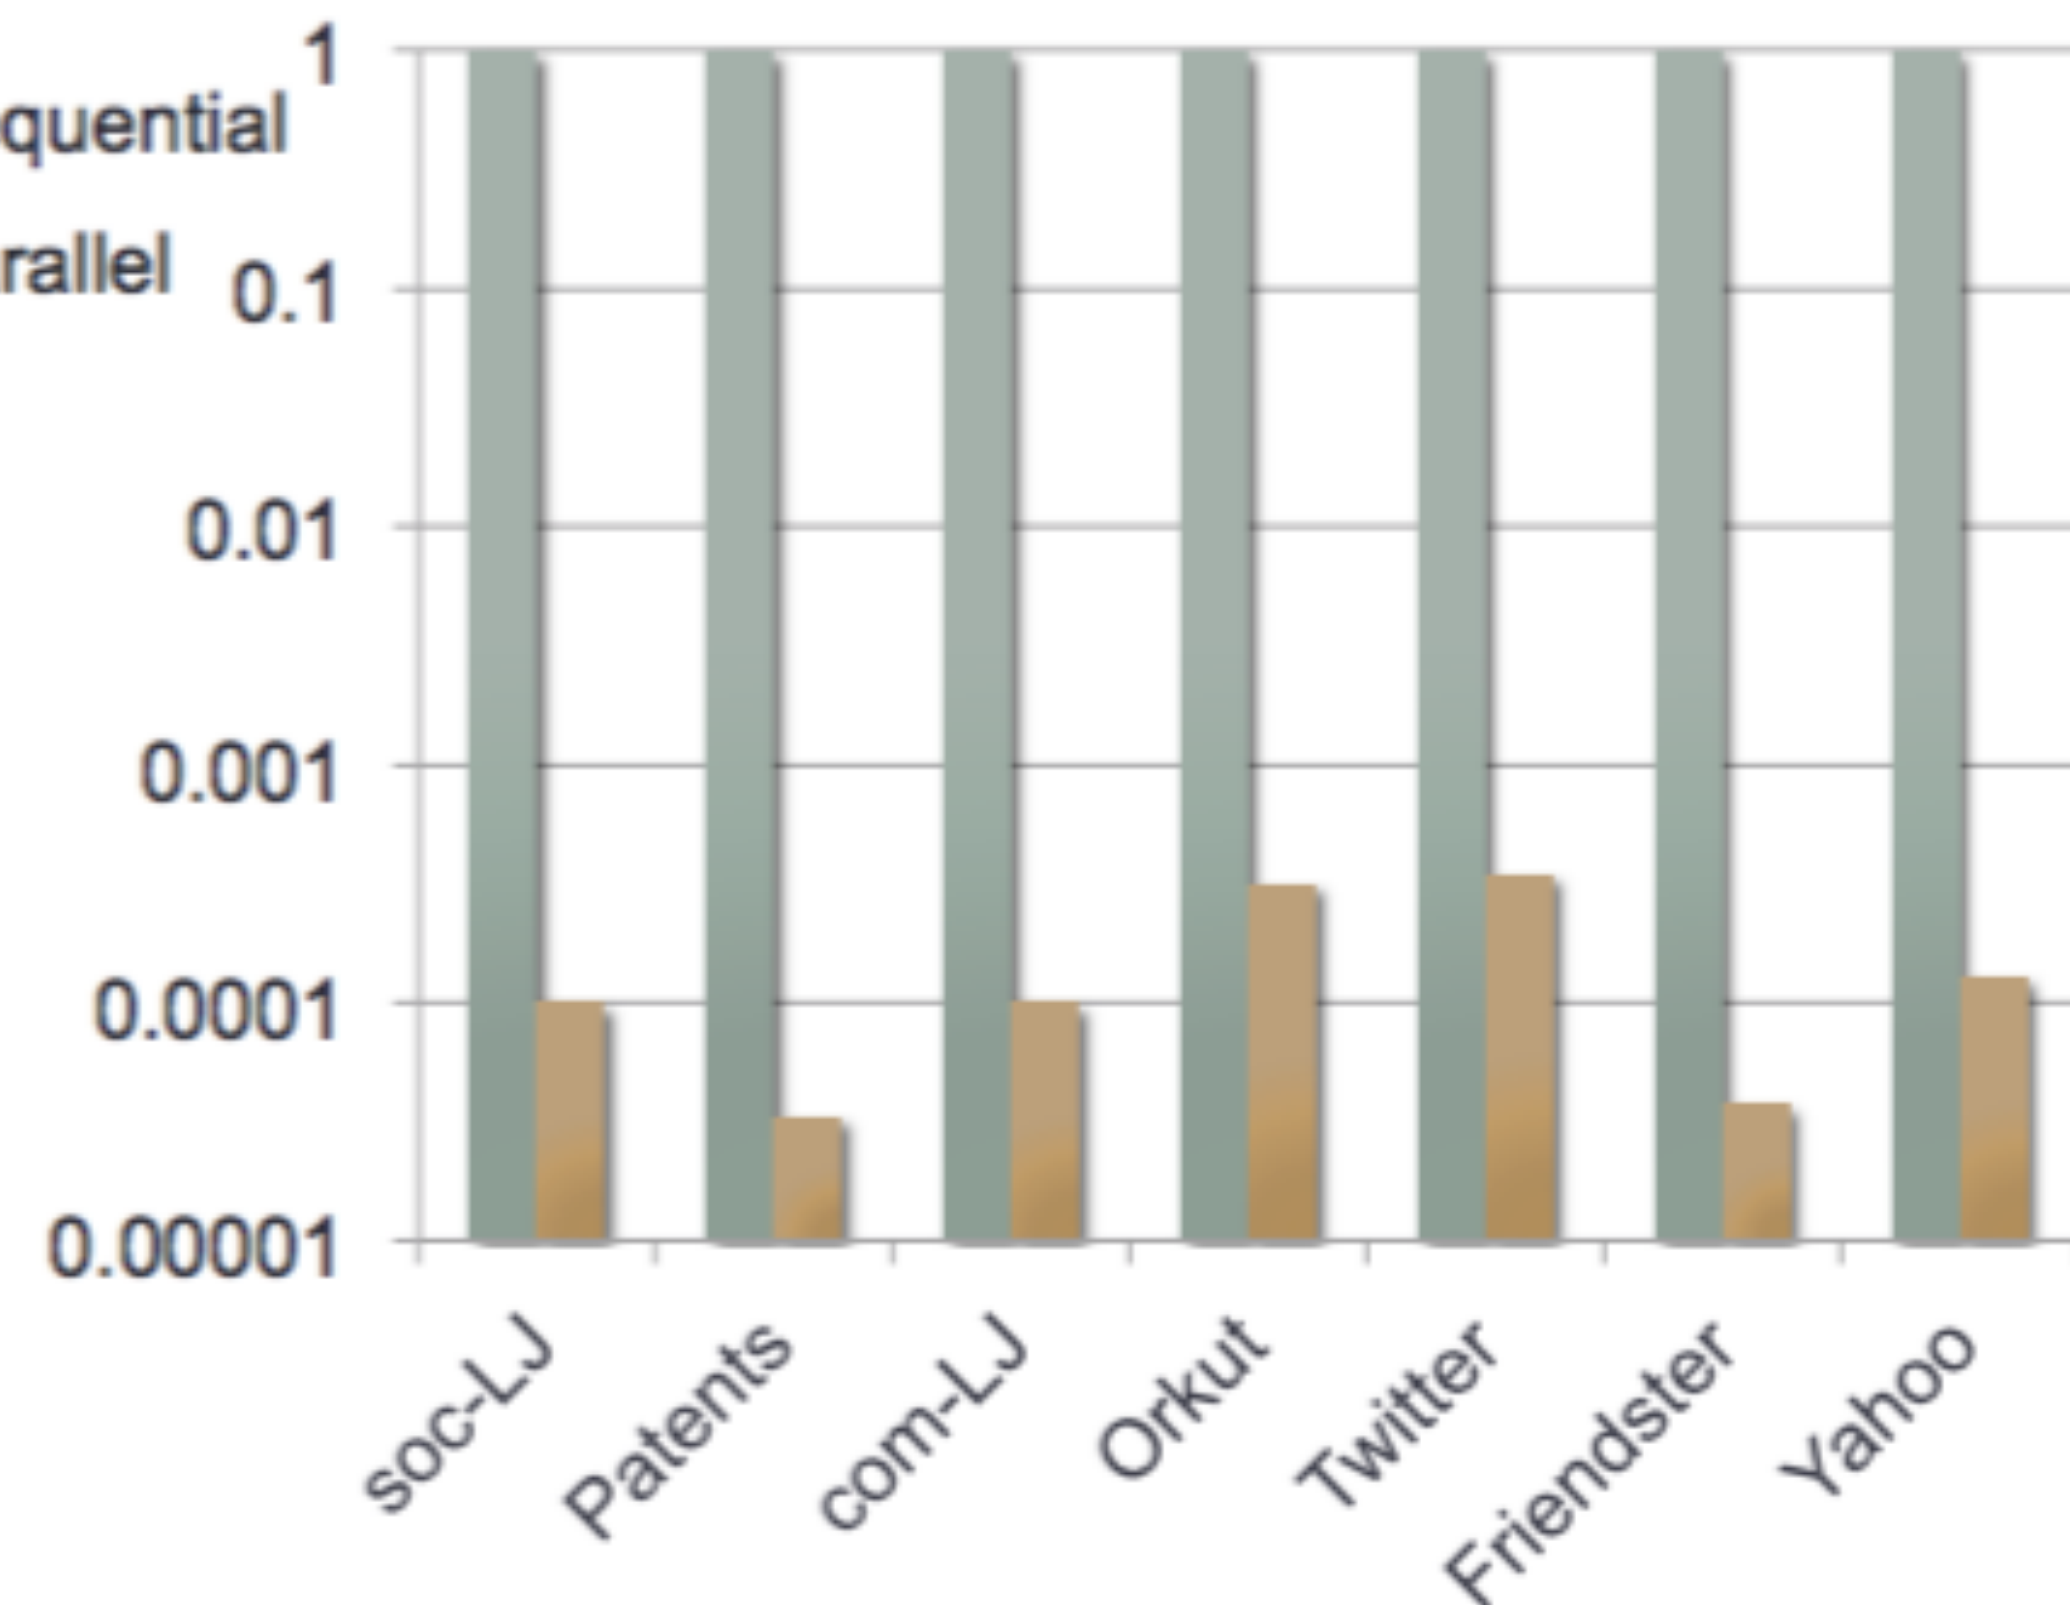


# Performance

Num. vertices processed (normalized)

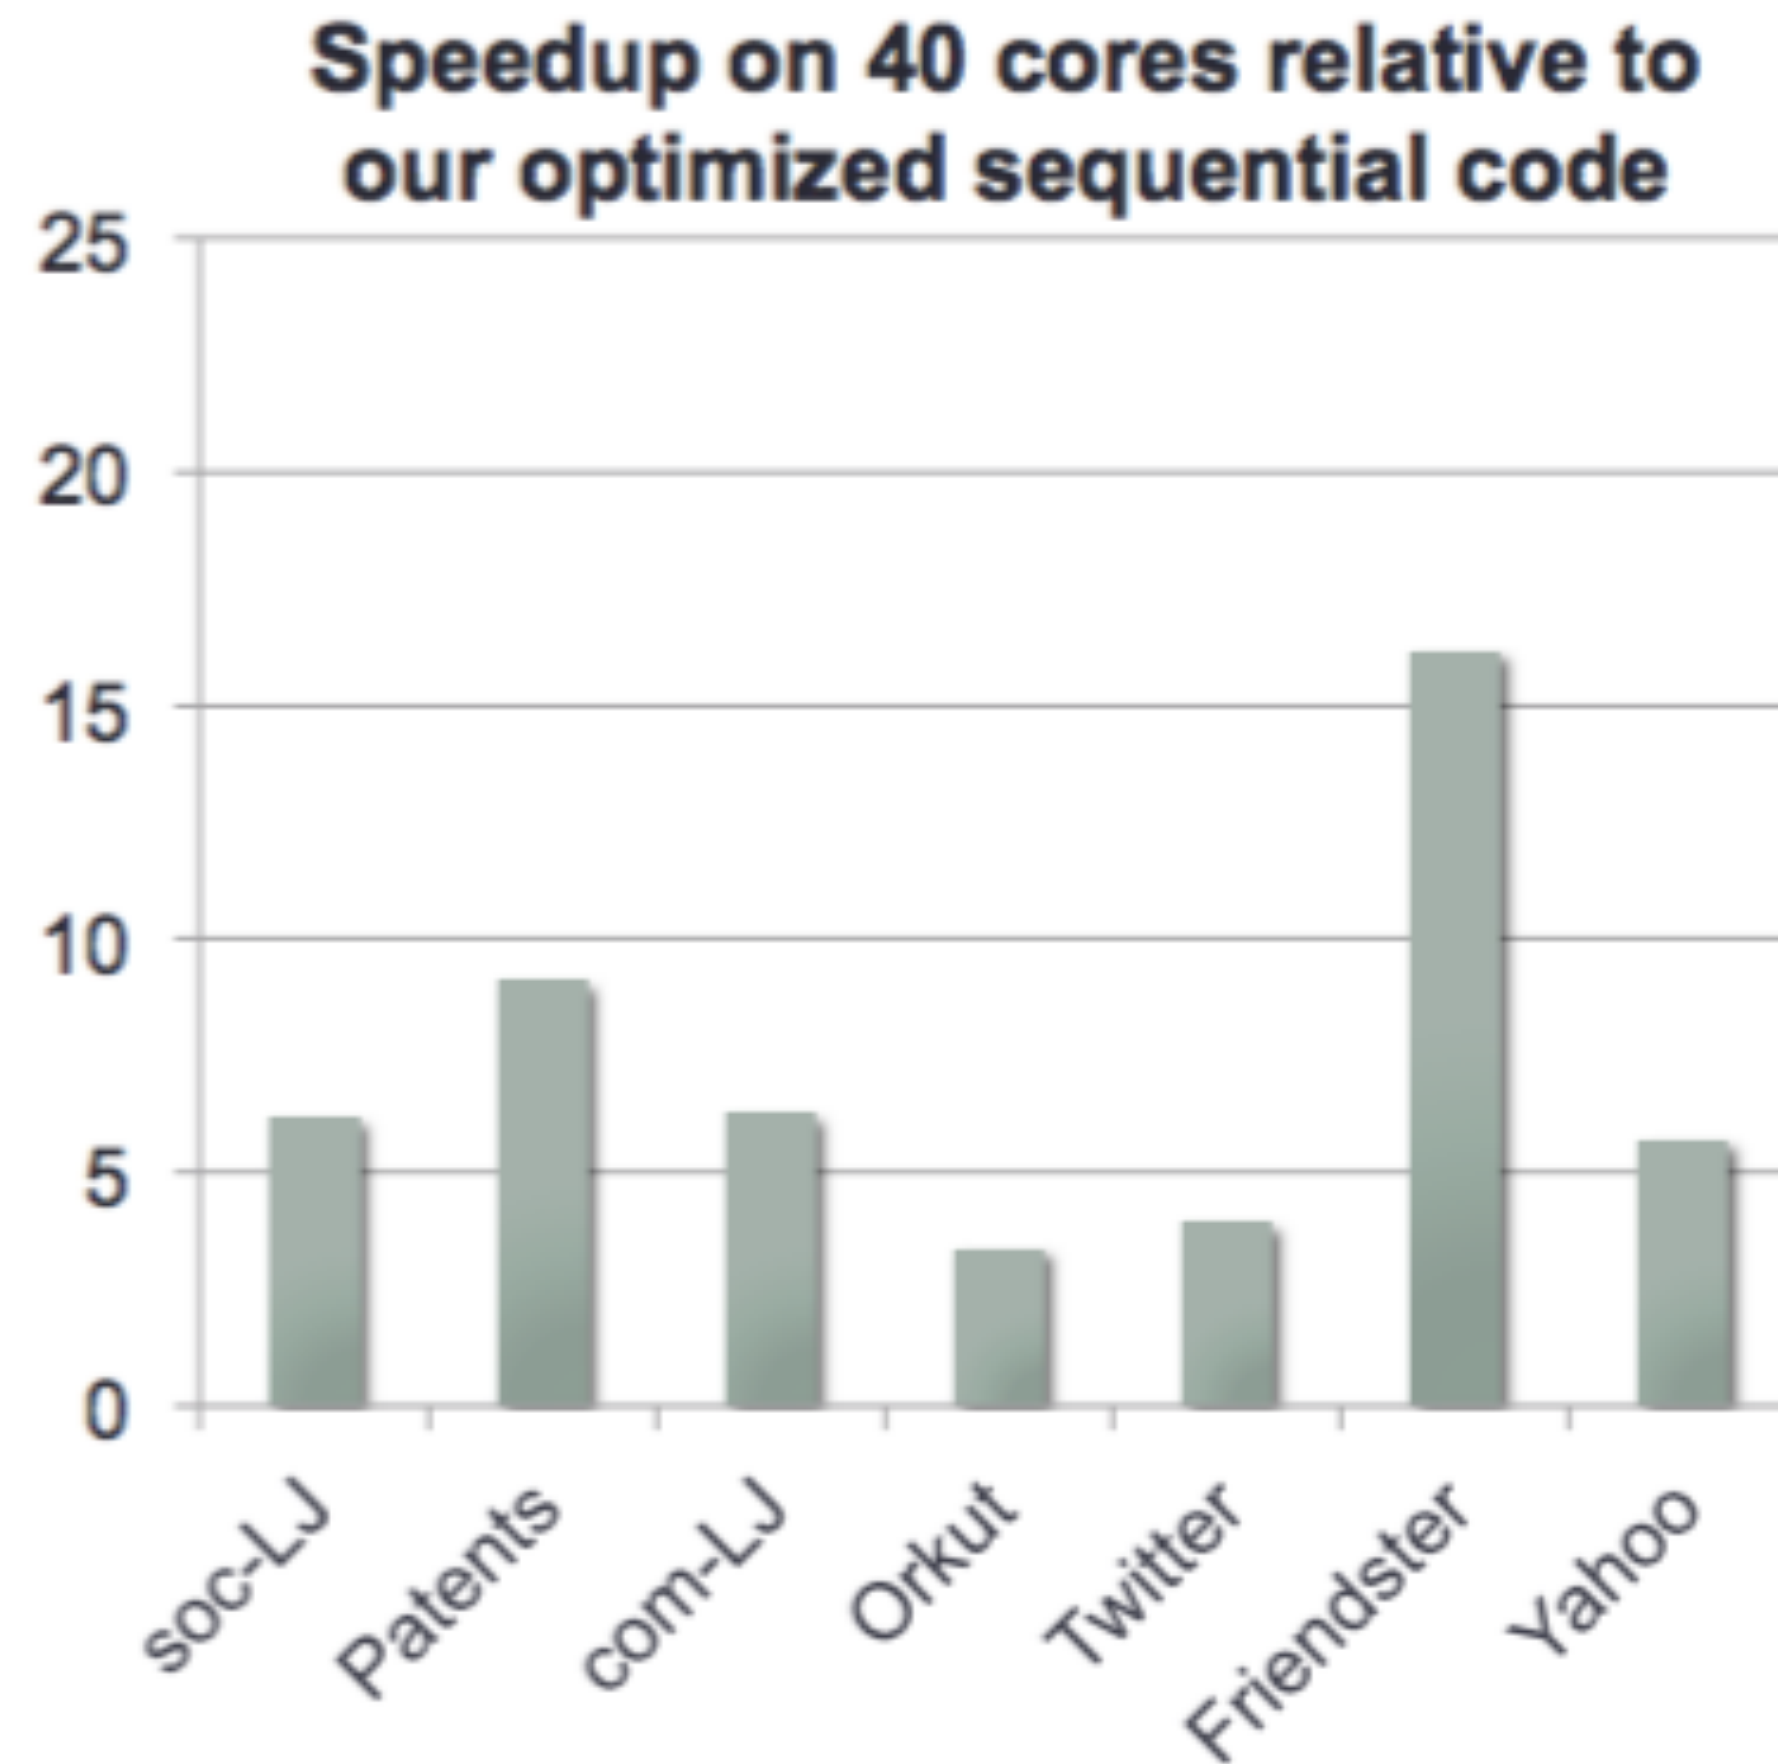
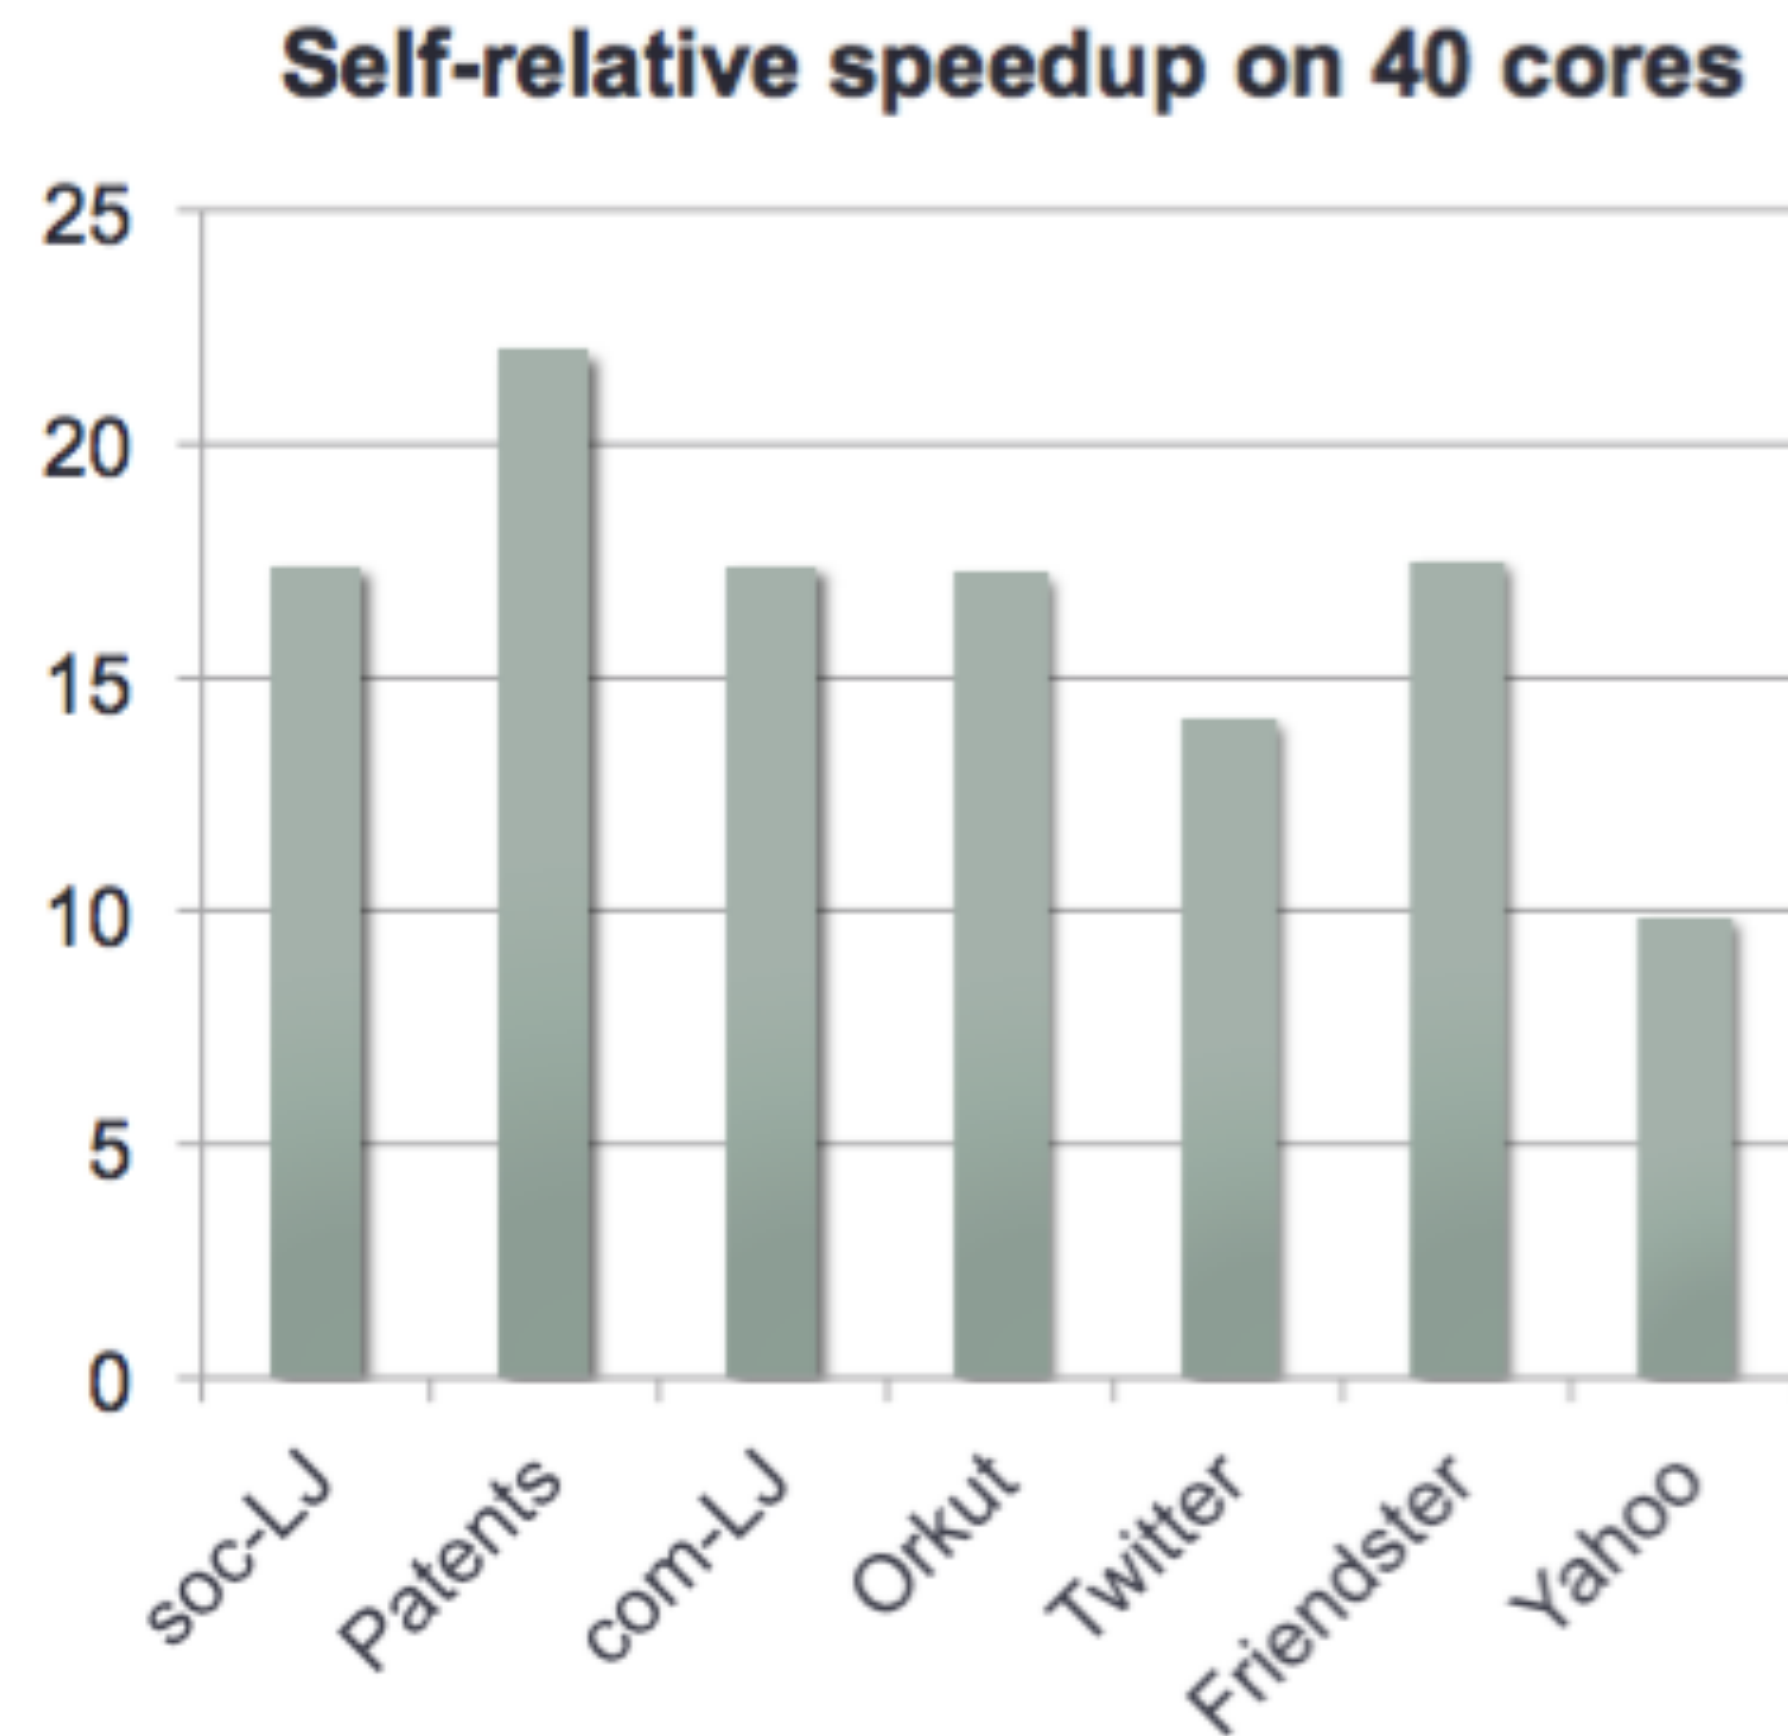


Num. iterations (normalized)



- Slightly more work for the parallel version
- Number of iterations is significantly less

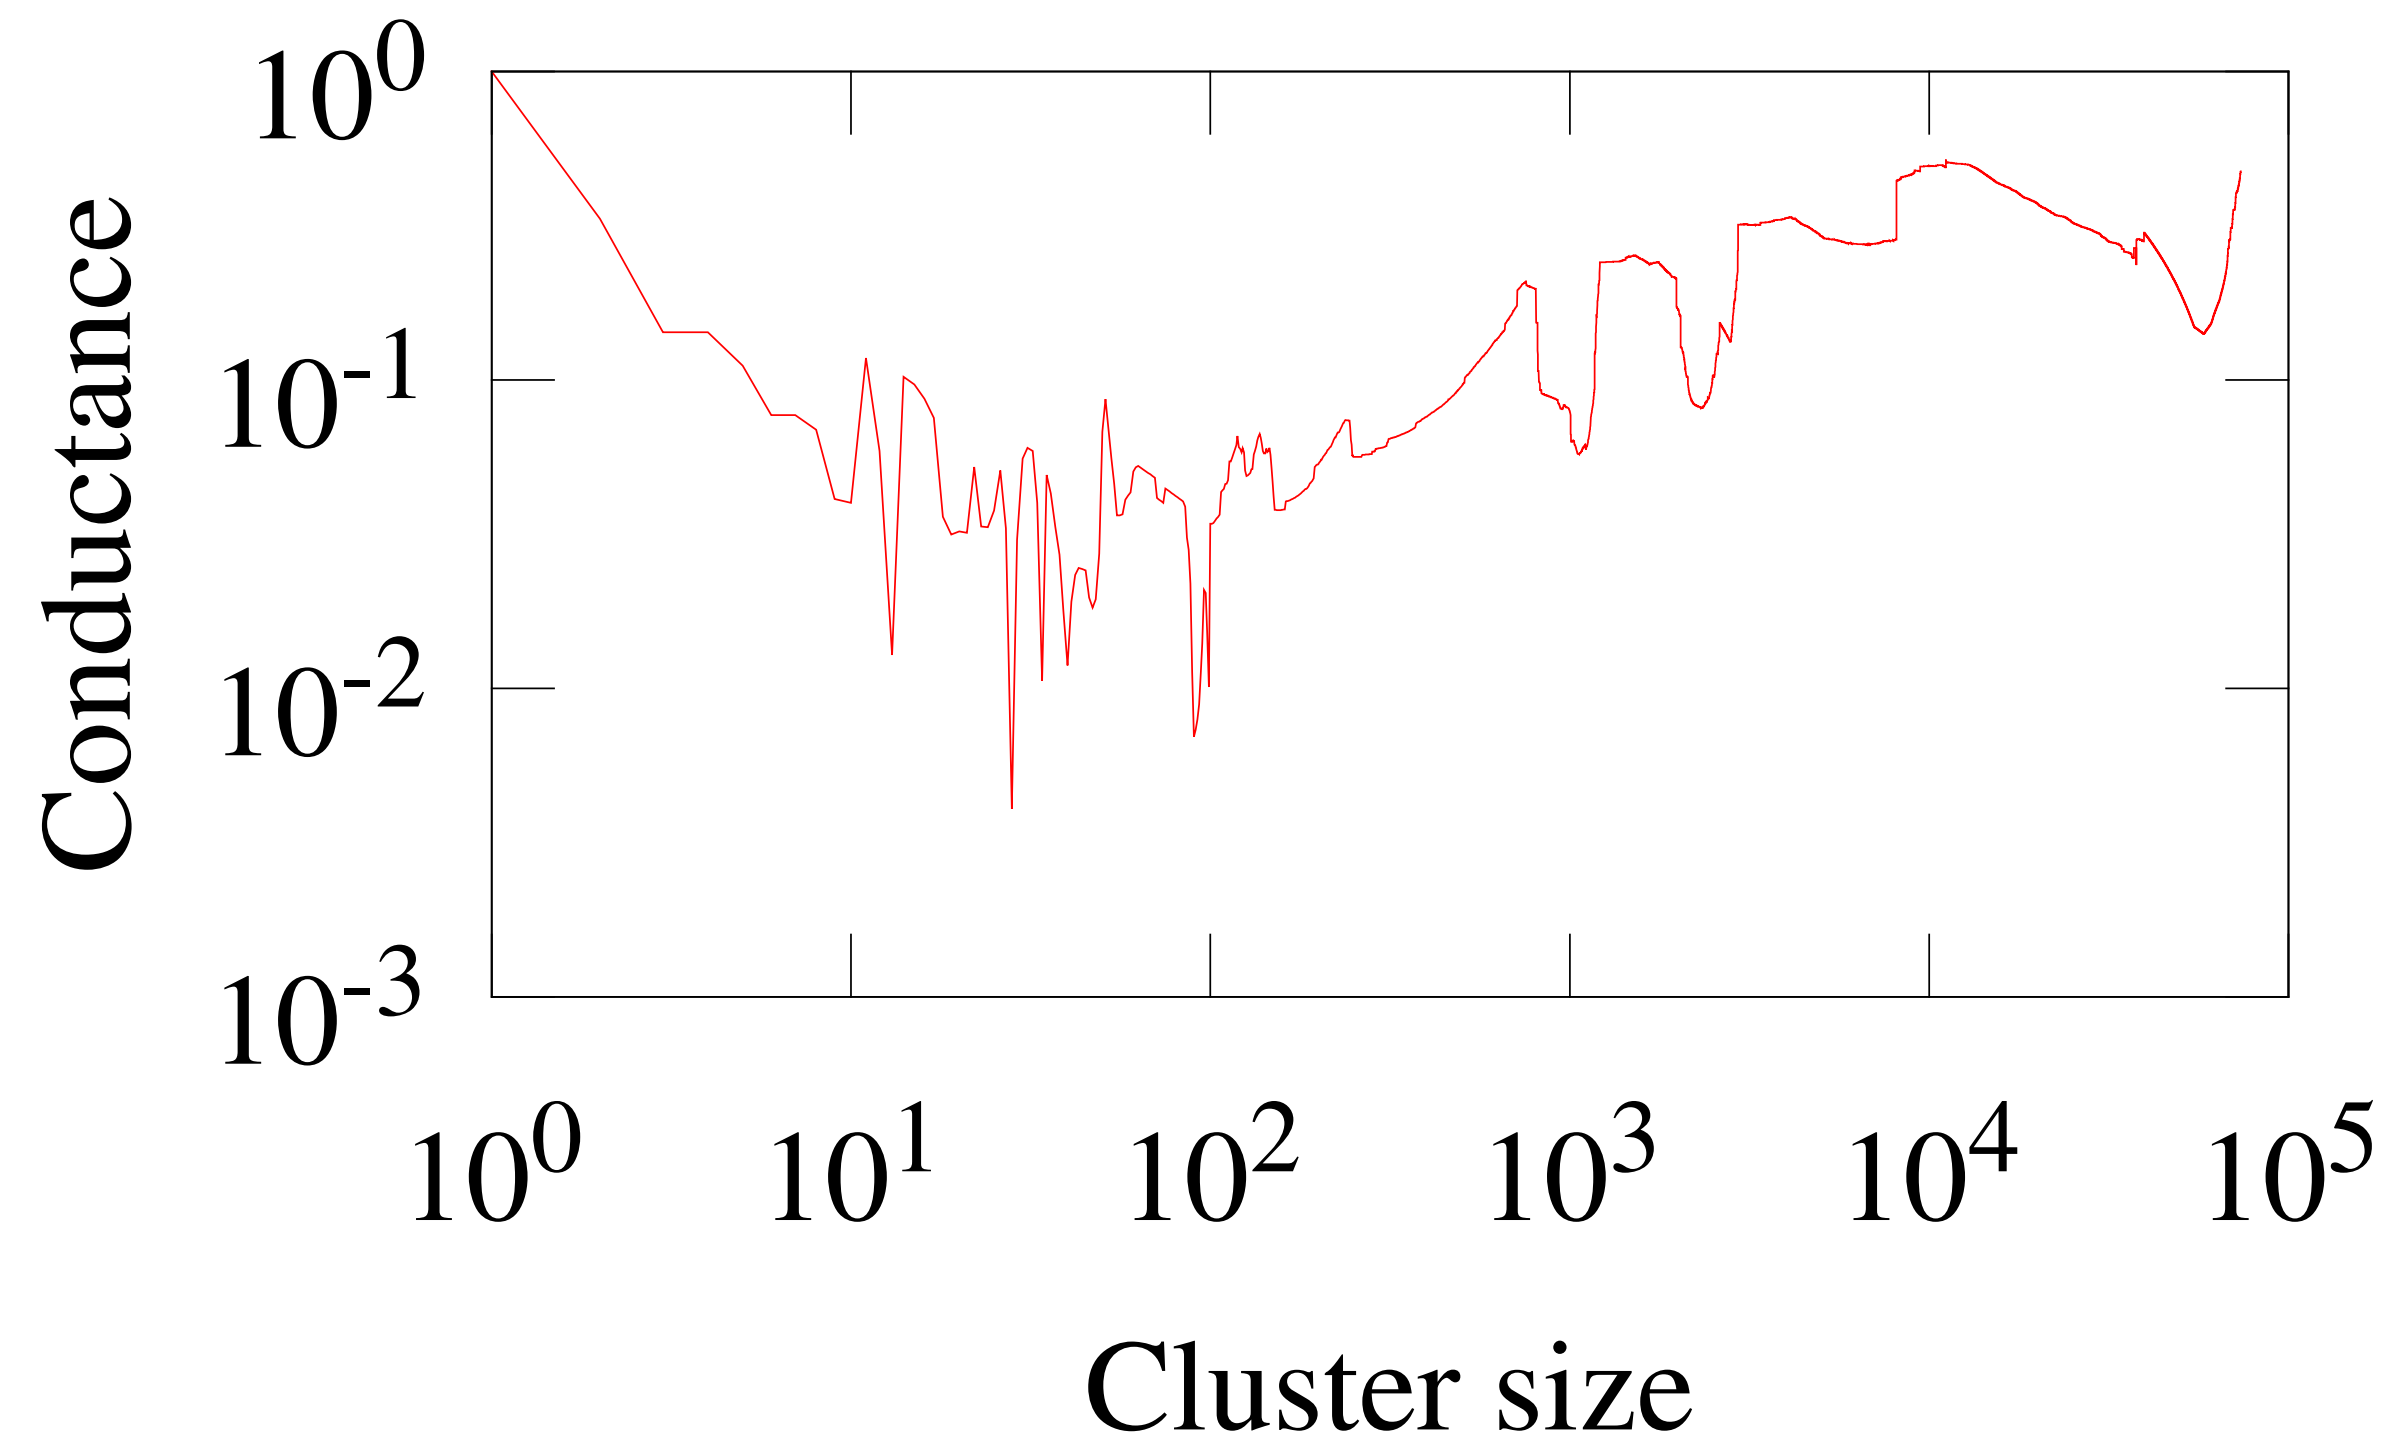
# Performance



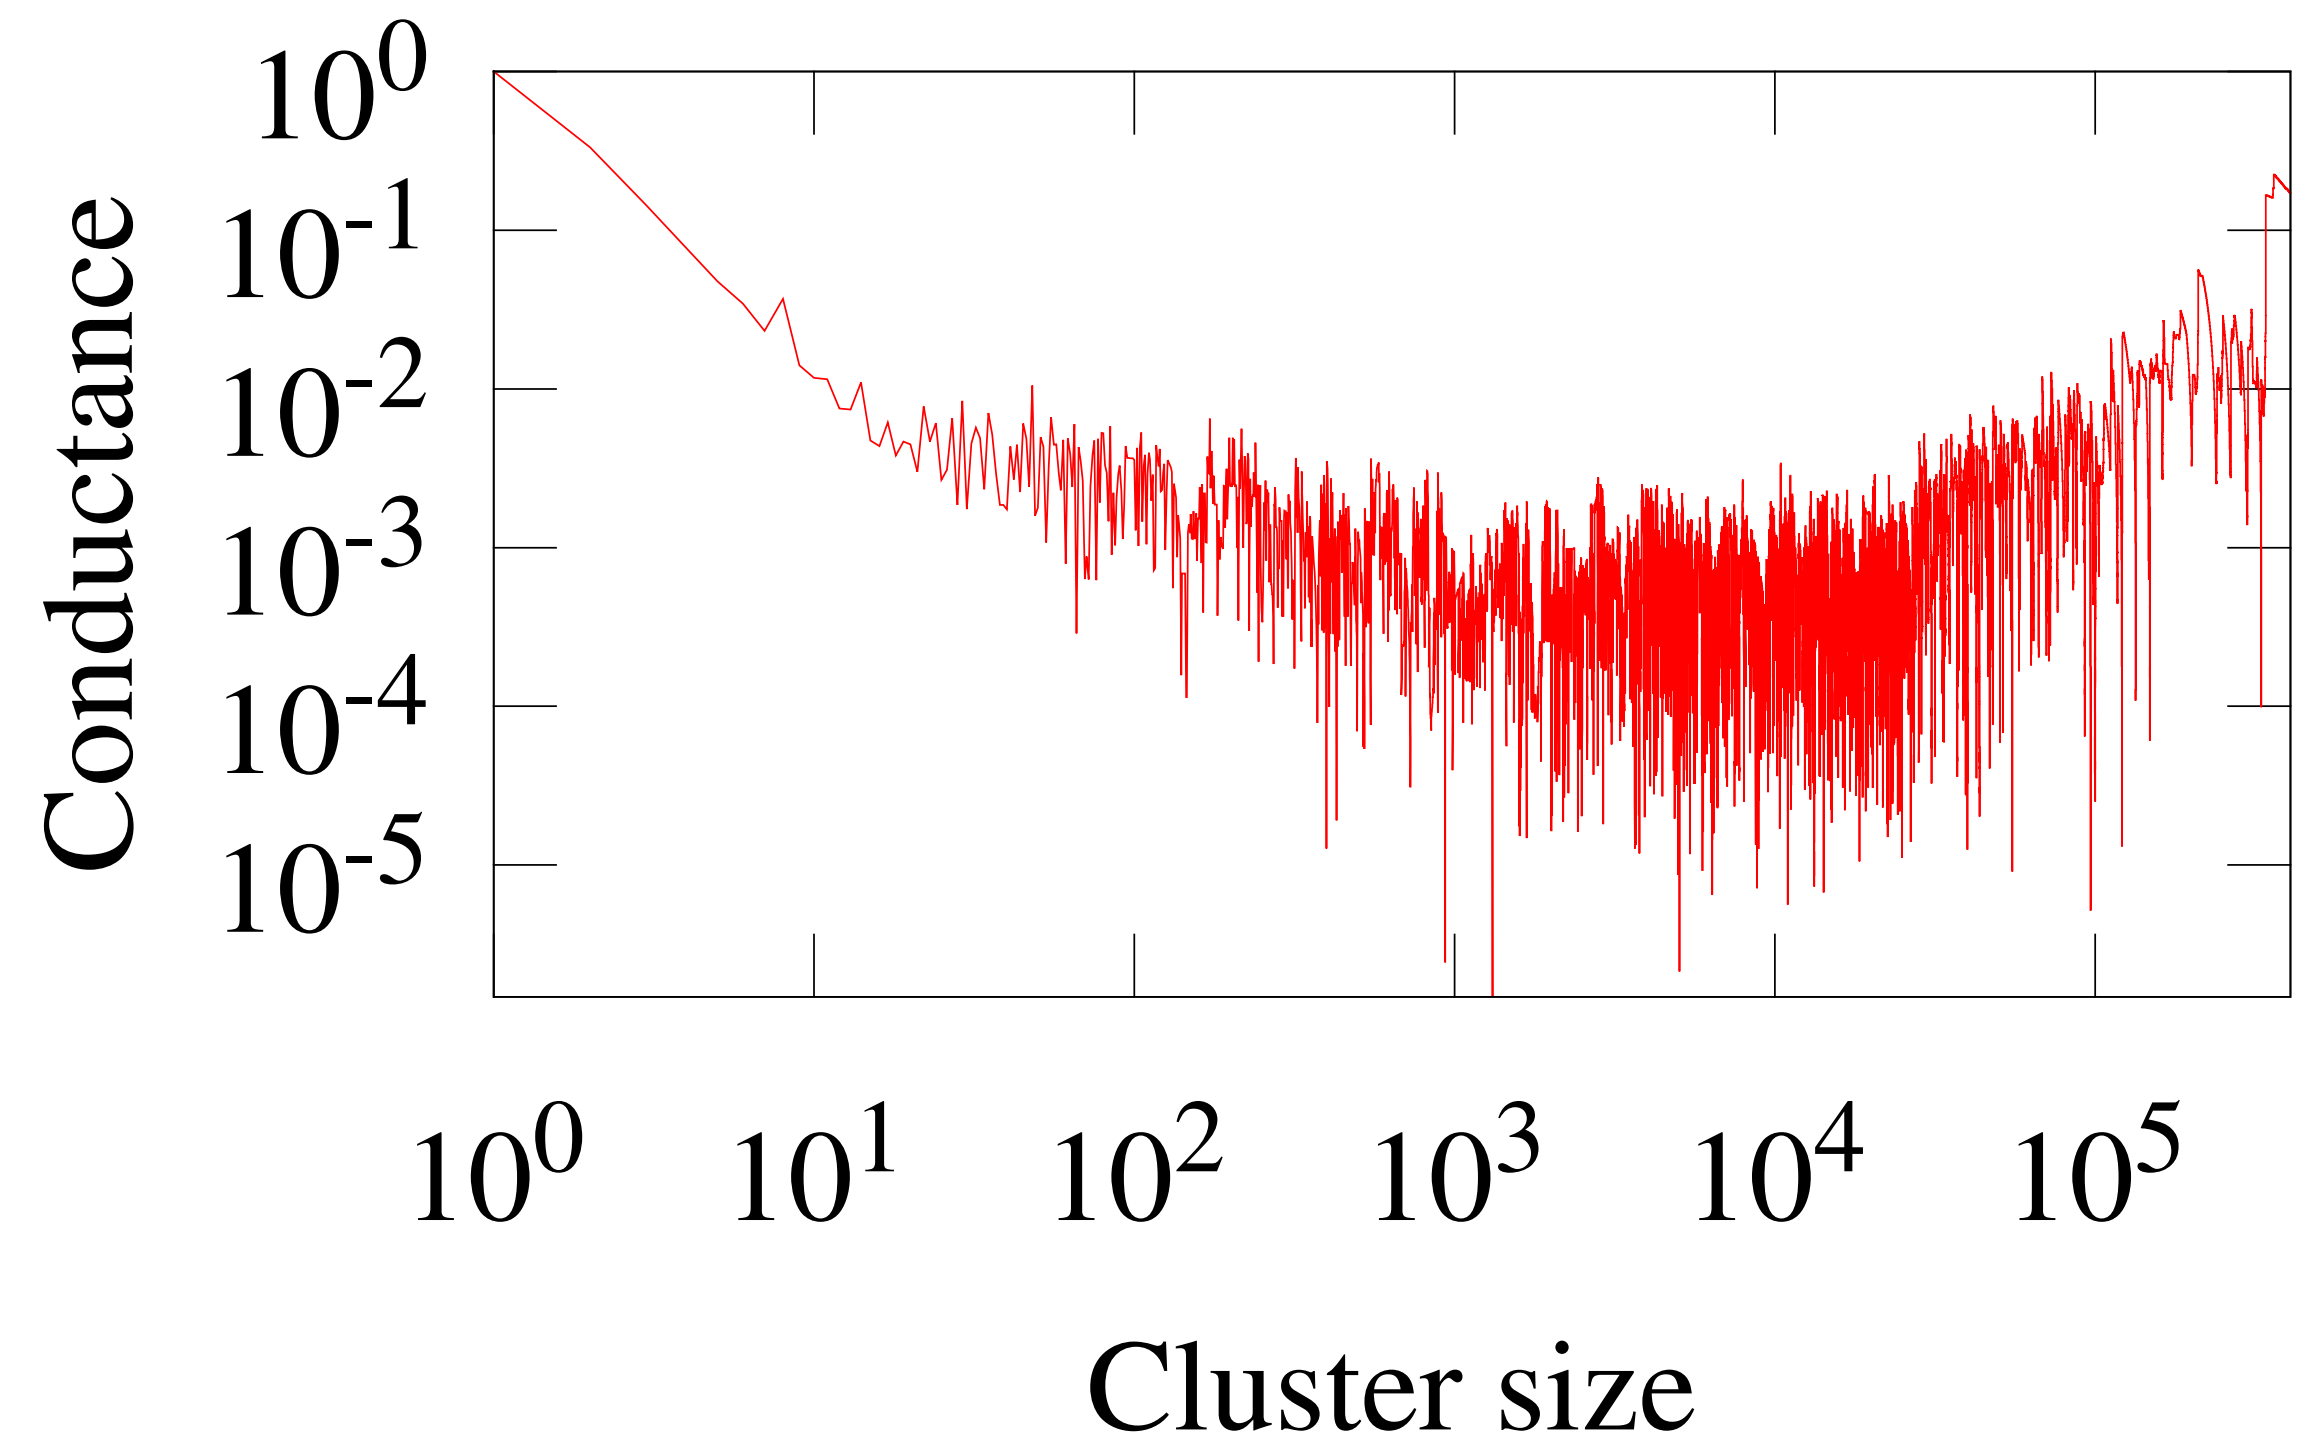
- 3-16x speed up
- Speedup is limited by small active set in some iterations and memory effects

# Network community profile plots

Friendster, 124M nodes, 1.8B edges



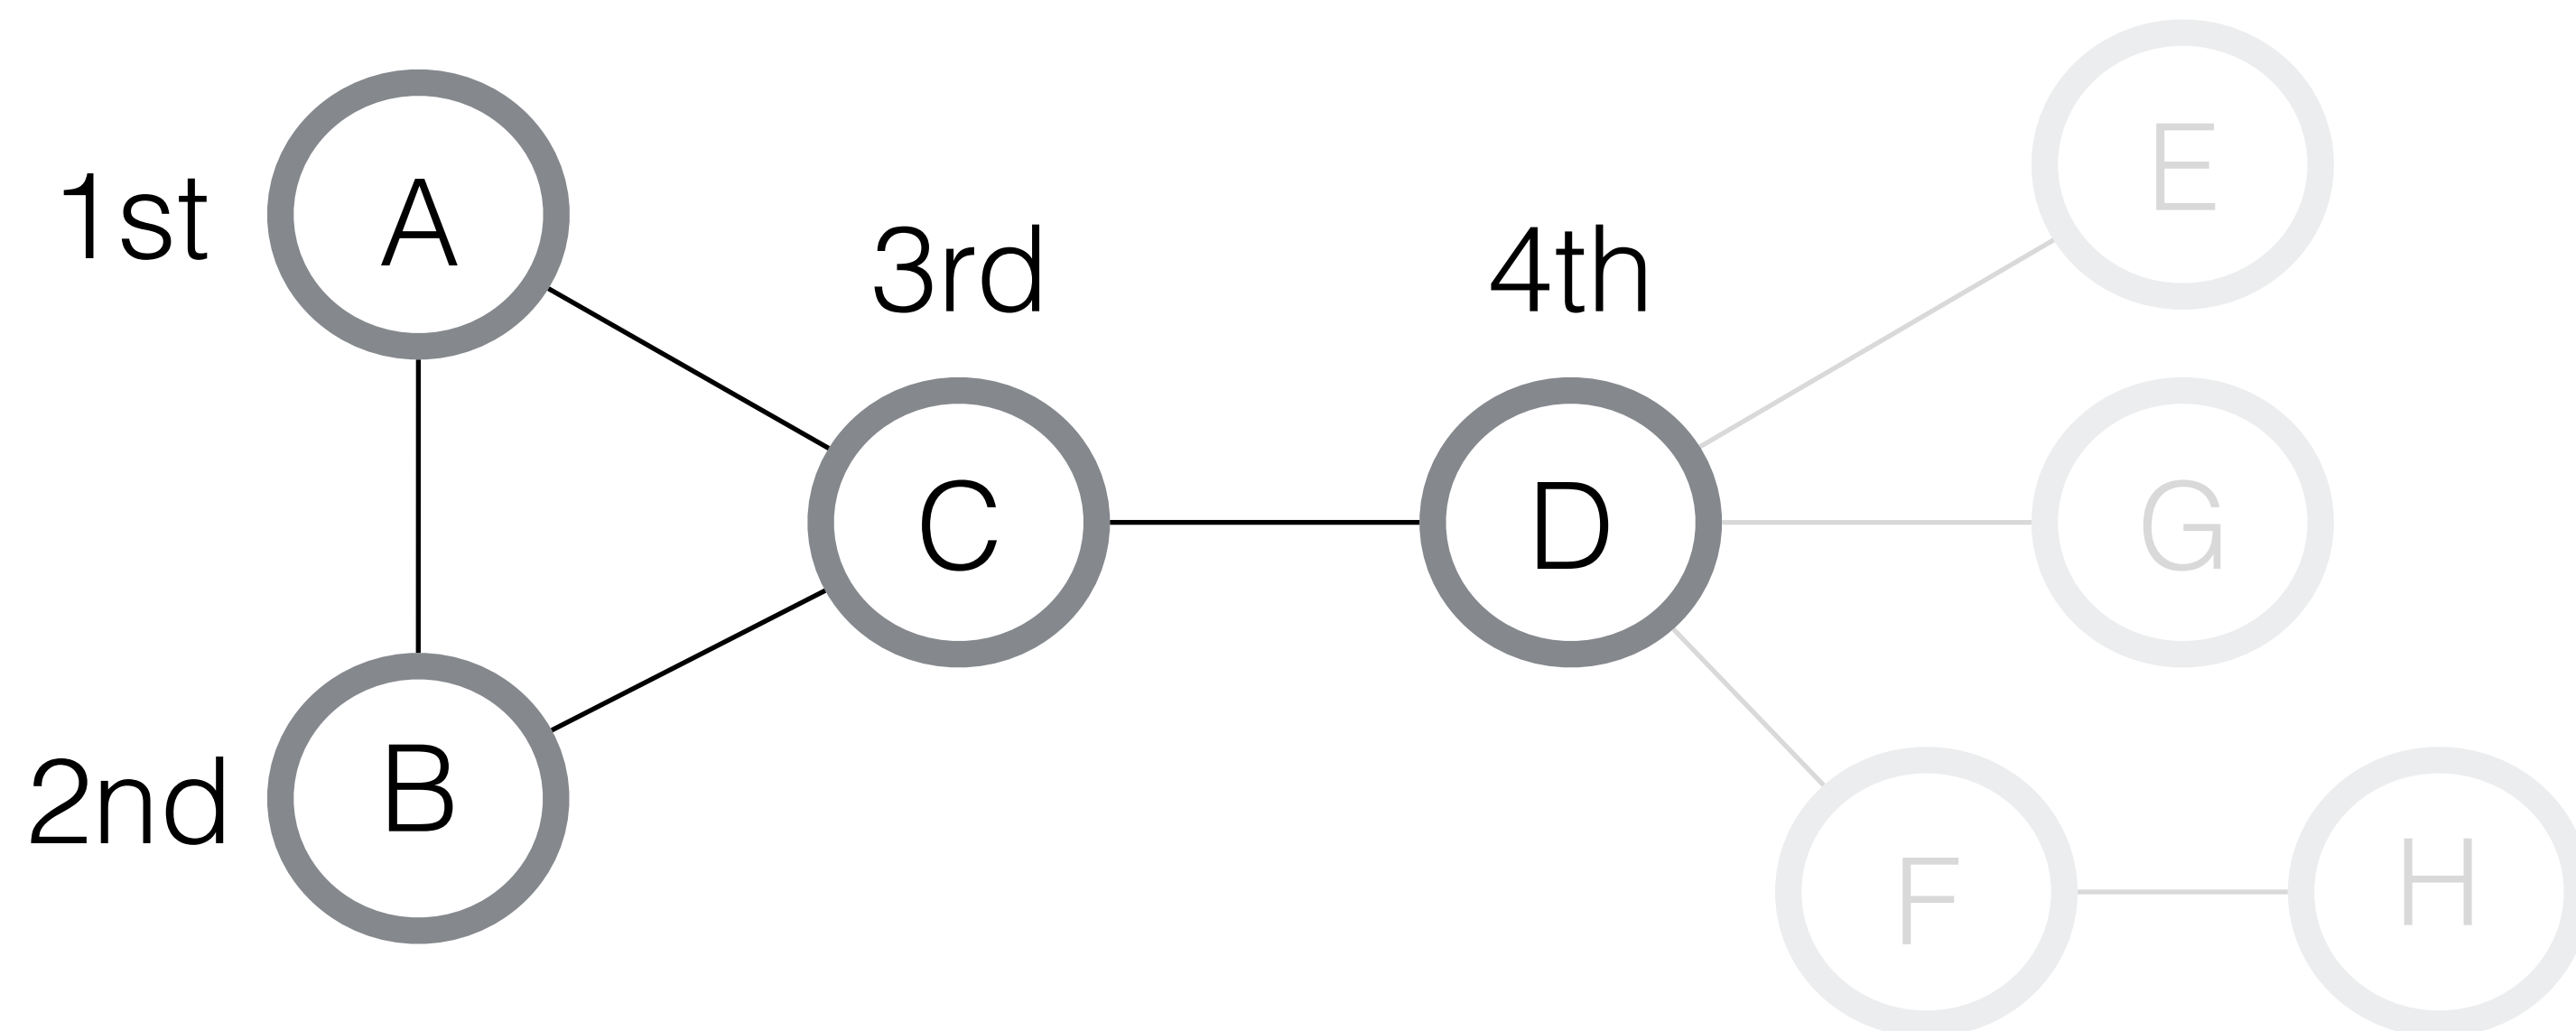
Yahoo, 1.4B nodes, 6.4B edges



- $O(10^5)$  approximate PPR problems were solved in parallel for each plot,
- Agrees with conclusions of [Leskovec et al. 2008], i.e., good clusters tend to be small.

# Rounding: sweep cut

- Round returned vector  $p$  of approximate PPR
  - **1st step ( $O(1/(ap) \log(1/(ap)))$  work):** Sort vertices by non-increasing value of non-zero  $p[u]/d[u]$
  - **2nd step ( $O(1/(ap))$  work):** Look at all prefixes of sorted order and return the cluster with minimum conductance,



Sorted vertices: {A,B,C,D}

Cluster	Conductance
{A}	1
{A,B}	1/2
<b>{A,B,C}</b>	<b>1/7</b>
{A,B,C,D}	3/11

# Parallel sweep cut

- **1st step:** Sort vertices by non-increasing value of non-zero  $p[u]/d[u]$ .
  - Use parallel sorting algorithm,  $O(1/(\alpha p) \log(1/(\alpha p)))$  work and  $O(\log(1/(\alpha p)))$  depth.
- **2nd step:** Look at all prefixes of sorted order and return the cluster with minimum conductance.
  - Naive implementation: for each sorted prefix compute conductance,  $O((1/(\alpha p))^2)$ .
  - **We design a parallel algorithm based on integer sorting and prefix sums that takes  $O(1/(\alpha p))$  time.**
  - **The algorithm computes the conductance of ALL sets with a single pass over the nodes and the edges.**



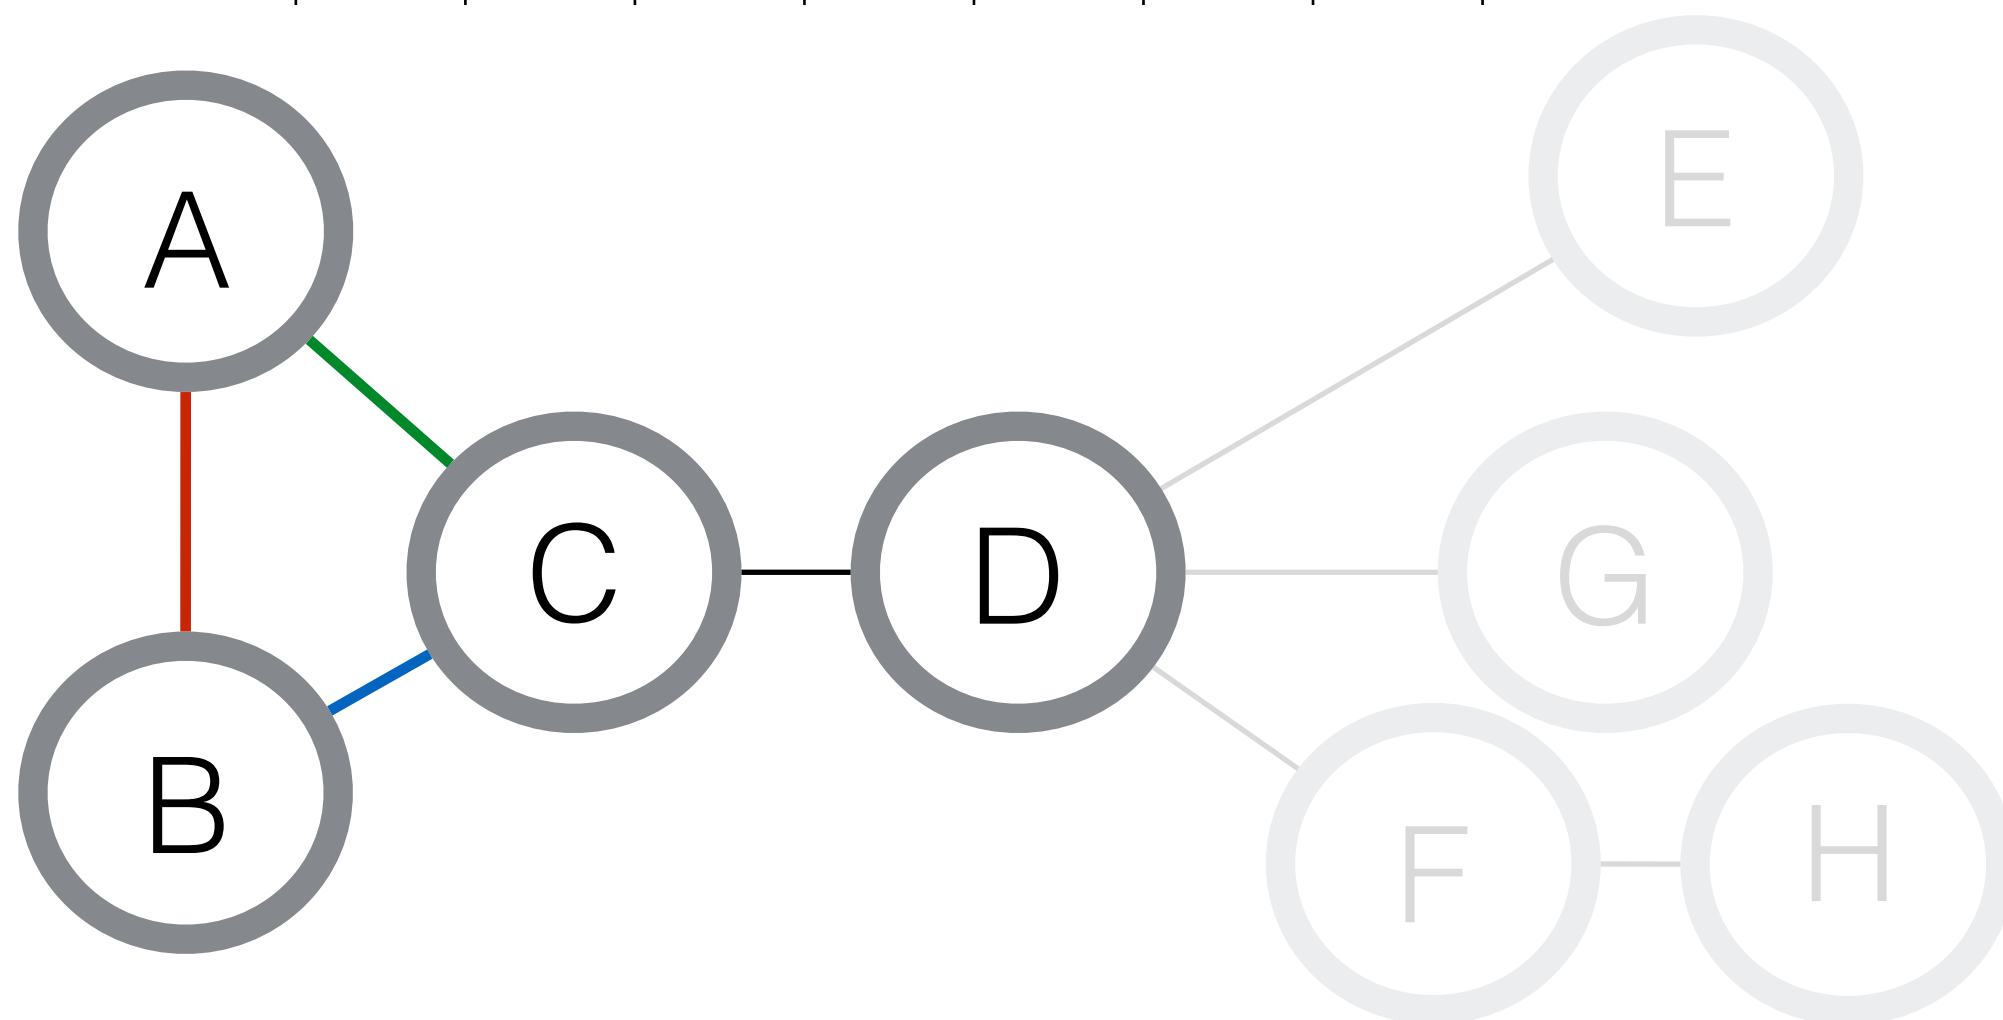
# Parallel sweep cut: 2nd step

Incidence matrix B

	A	B	C	D	E	F	G	H
A-B	1	-1						
A-C	1		-1					
B-C		1	-1					
C-D			1	-1				
D-E				1	-1			
D-F				1		-1		
D-G				1			-1	
F-H						1		-1

Sorted vertices: {A,B,C,D}

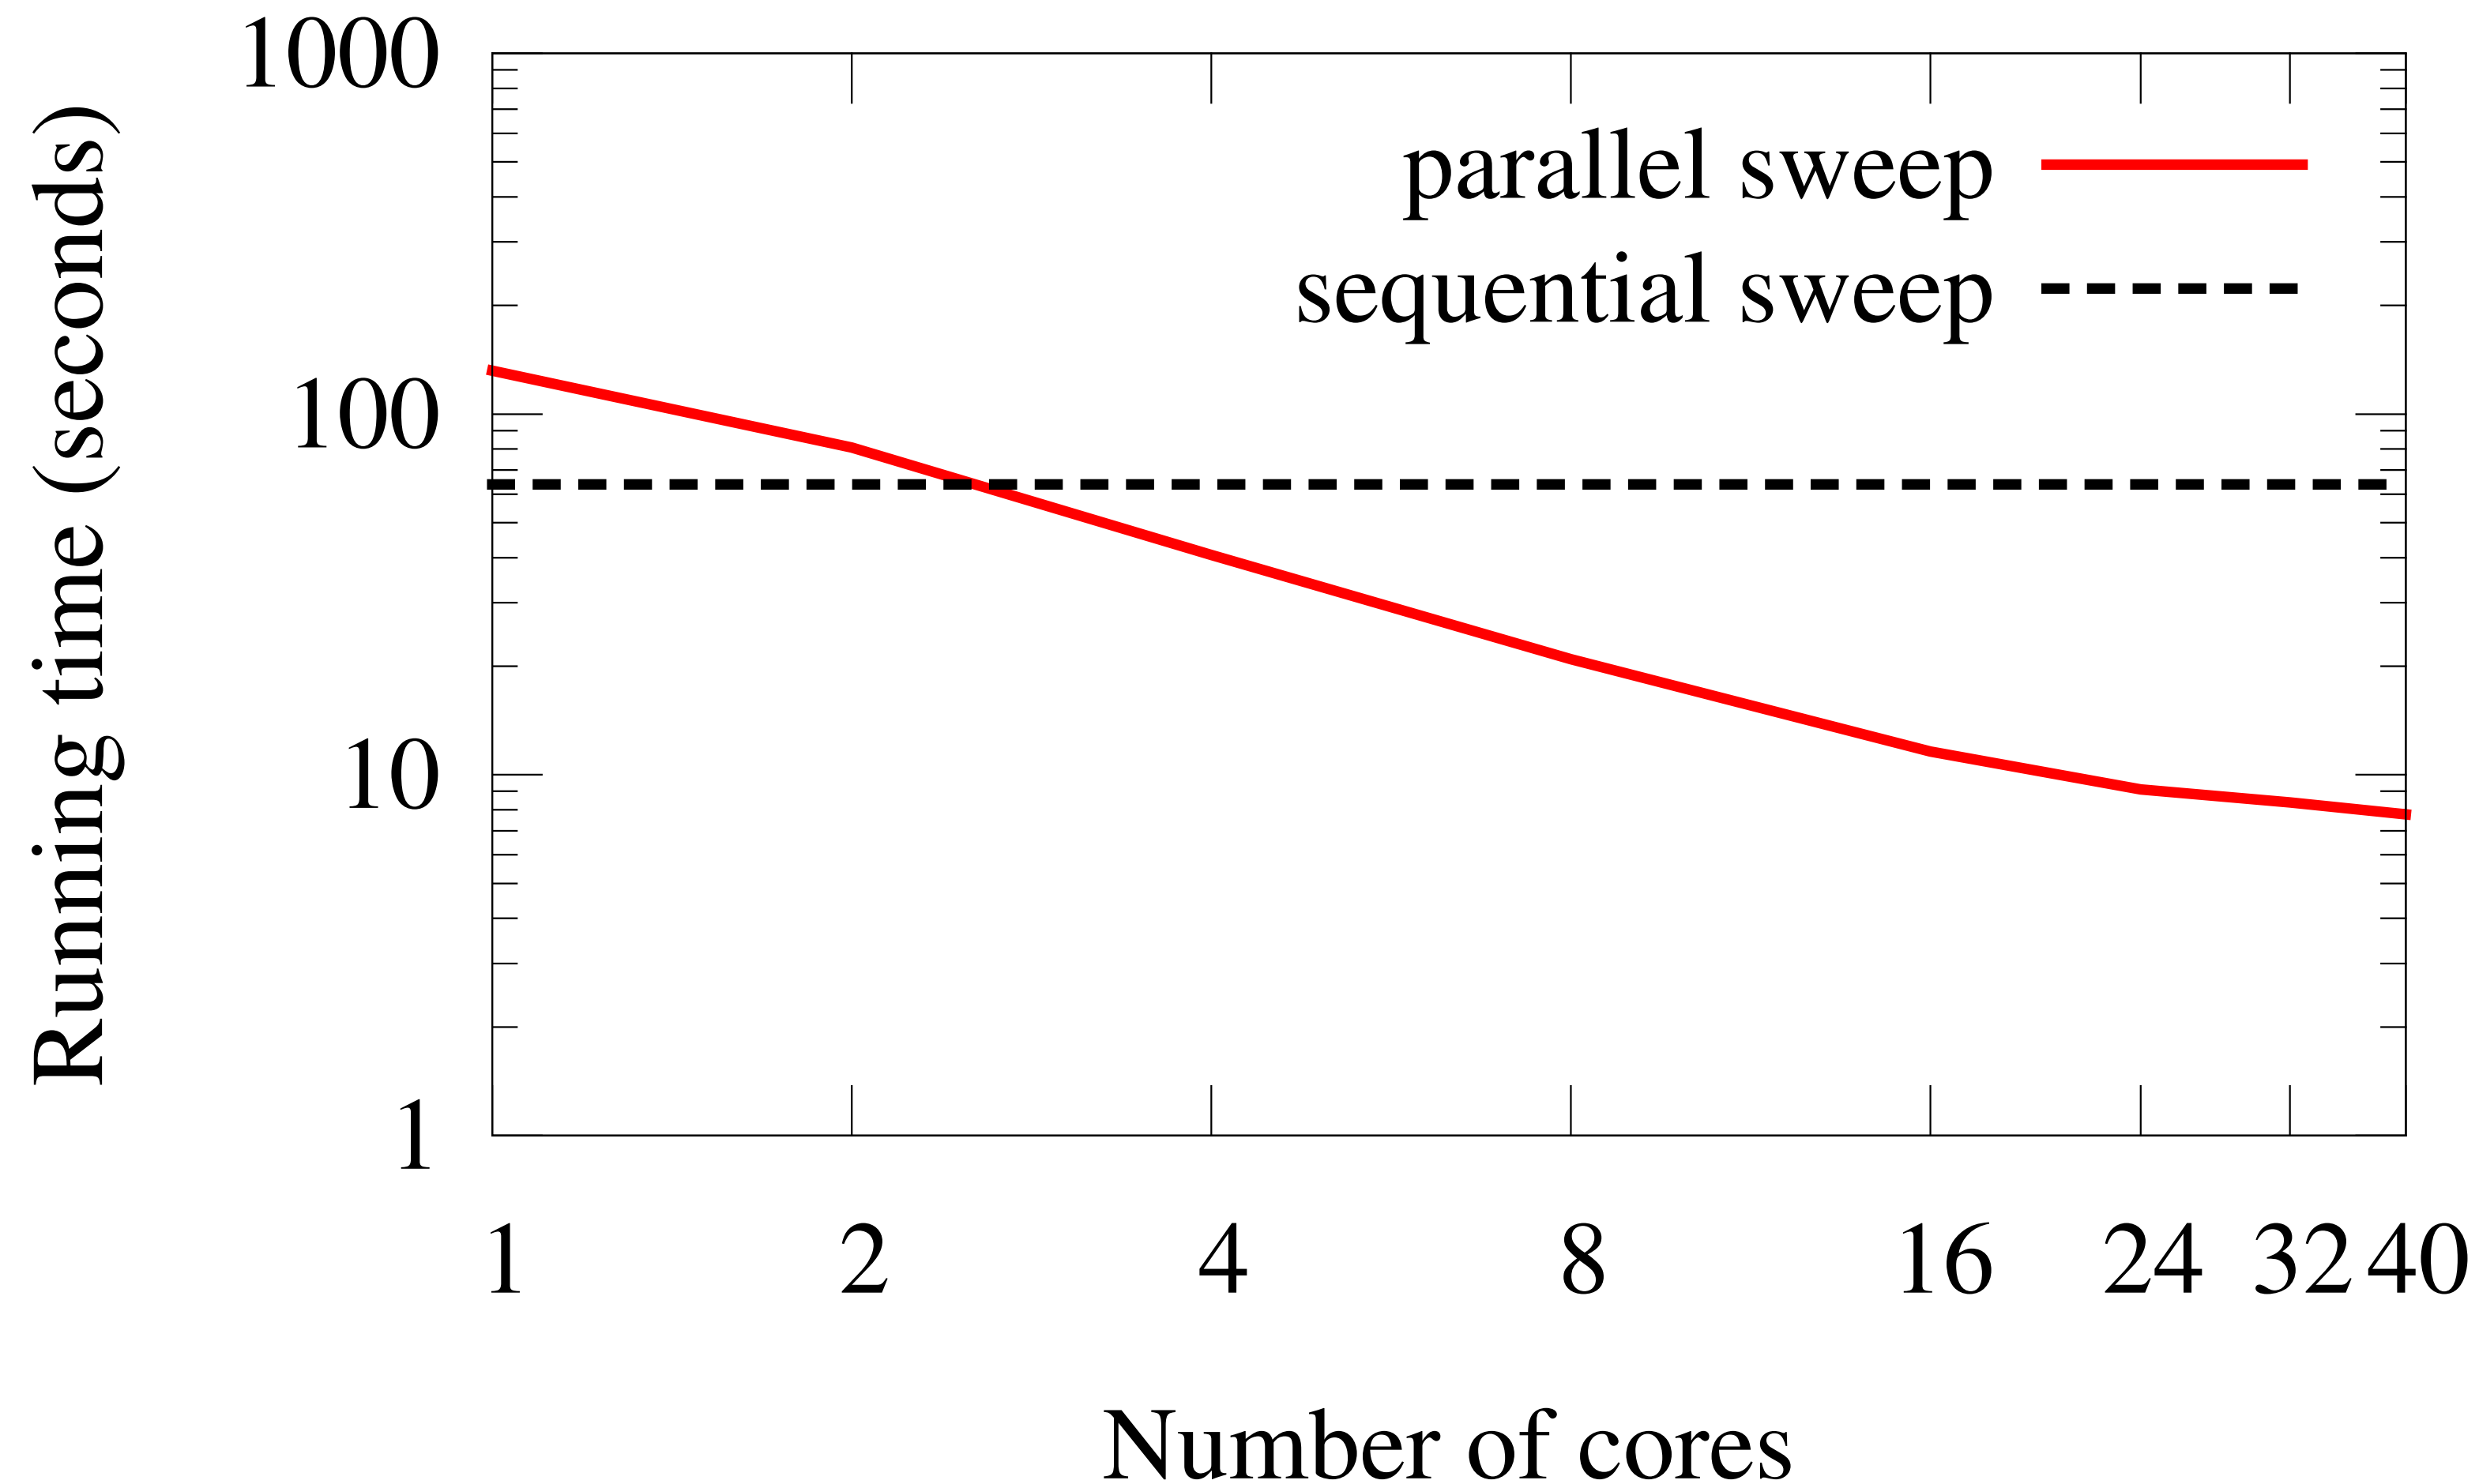
Cluster	Sum cols B	Volume	Conductance
{A}	<b>2</b>	<b>2</b>	$2/2=1$
{A,B}	2	4	$2/4=1/2$
{A,B,C}	1	7	$1/7$
{A,B,C,D}	3	11	$3/11$



- Sort vertices
  - work:  $O(1/(ap) \log(1/(ap)))$ , depth:  $O(\log(1/(ap)))$
- Represent matrix B with a sparse set using vertex identifiers and the order of vertices
  - work:  $O(1/(ap))$ , depth:  $O(\log(1/(ap)))$
- Use prefix sums to sum elements of the columns
  - work:  $O(1/(ap))$ , depth:  $O(\log(1/(ap)))$



# Parallel sweep cut: performance



# Summary

- We parallelise 4 spectral algorithms for local graph clustering.
- The proposed algorithms are work efficient, i.e., same worst-case work.
- We parallelise the rounding procedure to obtain the clusters.
- Useful in interactive setting where one has to experiment with parameters.
- 3-15x faster than sequential version
- Parallelisation allowed us to solve problems of billions of nodes and edges.

# Further work: distributed block coordinate descent

- Generalization to  $l_2$ -regularized least-squares and kernel learning:

A. Devarakonda, KF, J. Demmel, M. Mahoney: Avoiding communication in primal and dual block coordinate descent methods (work in progress  $\leq$  month)

- Given a positive integer  $k$ 
  - We reduce latency for BCD by a factor of  $k$
  - at the expense of a factor of  $k$  more work and number of words.

Thank you!